

# Artificial Intelligence

## Prolog

# What is Prolog?

- Can be downloaded at <http://www.swi-prolog.org/>
- Invented early 70s by Alain Colmerauer in France and Robert Kowalski in Britain.
- Programmation en Logique (Programming in Logic).
- differs from most common programming languages
- is a declarative language
  - programmer specifies a goal to be achieved
  - Prolog system works out how to achieve it



# What is Prolog?

- traditional programming languages are said to be **procedural**
- procedural programmer must specify in detail how to solve a problem:
  - mix ingredients;
  - beat until smooth;
  - bake for 20 minutes in a moderate oven;
  - remove tin from oven;
  - put on bench;
  - close oven;
  - turn off oven;
- in purely **declarative** languages, the programmer only states what the problem is and leaves the rest to the language system

# Applications of Prolog

- intelligent data base retrieval
- natural language understanding
- expert systems
- specification language
- machine learning
- robot planning
- automated reasoning
- problem solving
- ...



# Relations

- Prolog programs specify relationships among objects and properties of objects.
- When we say, "John owns the book", we are declaring the ownership relationship between two objects: John and the book.
- When we ask, "Does John own the book?" we are trying to find out about a relationship.
- Relationships can also be specified as rules such as:  
*Two people are sisters **if**  
they are both female **and** they have the same parents.*
- A rule allows us to find out about a relationship even if the relationship isn't explicitly stated as a fact.

# Programming in Prolog

- declare facts describing ***explicit relationships*** between objects and properties objects might have (e.g. Mary likes pizza, grass has\_colour green)
- define rules defining ***implicit relationships*** between objects (e.g. the sister rule above) and/or rules defining implicit object properties (e.g. X is a parent if there is a Y such that Y is a child of X).

One then uses the system by:

- asking questions about relationships between objects, and/or about object properties (e.g. does Mary like pizza? is Joe a parent?)



# Facts

- Properties of objects, *or* relationships between objects;
- "Dr Turing lectures in course 9020", is written in Prolog as:  
*lectures(turing, 9020).*
- *Notice:*
  - names of properties/relationships begin with lower case letters.
  - the relationship name appears as the first term
  - objects appear as comma-separated arguments within parentheses.
  - A period "." must end a fact.
  - objects also begin with lower case letters. They also can begin with digits (like 9020), and can be strings of characters enclosed in quotes (as in *reads(fred, "War and Peace")*).
- *lectures(turing, 9020).* is also called a *predicate*

# Facts

- Facts about a hypothetical computer science department:

% lectures(X, Y): person X lectures in course Y

lectures(turing, 9020).

lectures(codd, 9311).

lectures(backus, 9021).

lectures(ritchie, 9201).

lectures(minsky, 9414).

lectures(codd, 9314).

% studies(X, Y): person X studies in course Y

studies(fred, 9020).

studies(jack, 9311).

studies(jill, 9314).

studies(jill, 9414).

studies(henry, 9414).

studies(henry, 9314).

%year(X, Y): person X is in year Y

year(fred, 1).

year(jack, 2).

year(jill, 2).

year(henry, 4).



# Queries

- Once we have a database of facts (and, soon, rules) we can ask questions about the stored information.
- Suppose we want to know if Turing lectures in course 9020. We can ask:

|   |  |
|---|--|
| <pre>% <i>prolog -s facts03.pro</i> (multi-line welcome message) ?- <i>lectures(turing, 9020).</i> Yes/True ?- <i>&lt;control-D&gt;</i> %</pre> | <pre>facts03 loaded into Prolog "?-" is Prolog's prompt output from Prolog hold down control &amp; press D to leave Prolog</pre> |
|---|--|

- **Notice:**
  - In SWI Prolog, queries are terminated by a full stop.
  - To answer this query, Prolog consults its database to see if this is a known fact.
  - In example dialogues with Prolog, the text in *green italics* is what the user types.

# Query

?- *lectures(codd, 9020).*

No/fail

- if answer is Yes/true, the query *succeeded*
- if answer is No/fail, the query *failed*
- The use of lower case for codd is critical.
- Prolog is not being intelligent about this - it would not see a difference between this query and *lectures(fred, 9020).* or *lectures(xyzzy, 9020).* though a person inspecting the database can see that fred is a student, not a lecturer, and that xyzzy is neither student nor lecturer.



# Variables

- Question: "What course does Turing teach"?
- This could be written as:  
Is there a course, X, that Turing teaches?
- The variable X stands for an object which the questioner does not know about yet.
- Prolog has to find out the value of X, if it exists.
- As long as we do not know the value of a variable it is said to be *unbound*.
- When a value is found, the variable is said to *bound* to that value.
- The name of a variable must begin with a capital letter or an underscore character, "\_".

# Variables

- To ask Prolog to find the course which Turing teaches, the following query is entered:

?- *lectures(turing, Course)*.

Course = 9020 ← output from Prolog

- To ask which course(s) Prof. Codd teaches, we may ask,

?- *lectures(codd, Course)*.

Course = 9311 ; ← type ";" to get next solution

Course = 9314 ;

No

- Prolog can find all possible ways to answer a query, unless you explicitly tell it not to (see *cut*, later).



# Conjunctions of Goals

- How do we ask, "Does Turing teach Fred"?
- This means finding out if Turing lectures in a course that Fred studies.
- ?- *lectures(turing, Course), studies(fred, Course).*
- To answer this question, Prolog must find a single value for Course, that satisfies both goals.
- Read the comma, ",", as **and**.
- However, note that Prolog will evaluate the two goals left-to-right. This is sometimes referred to as "conditional-and".

# Backtracking

- Who does Codd teach?

?- *lectures(codd, Course), studies(Student, Course).*

Course = 9311 Student = jack ;

Course = 9314 Student = jill ;

Course = 9314 Student = henry ;

- Processes left to right and then *backtracking*.
- Prolog starts by trying to solve *lectures(codd, Course)*
- six lectures clauses, only two have codd as their first argument.
- Uses the first clause that refers to codd: *lectures(codd, 9311)*.
- It tries the next goal, *studies(Student, 9311)*.
- It finds the fact *studies(jack, 9311)*. and hence the first solution:  
(Course = 9311, Student = jack)

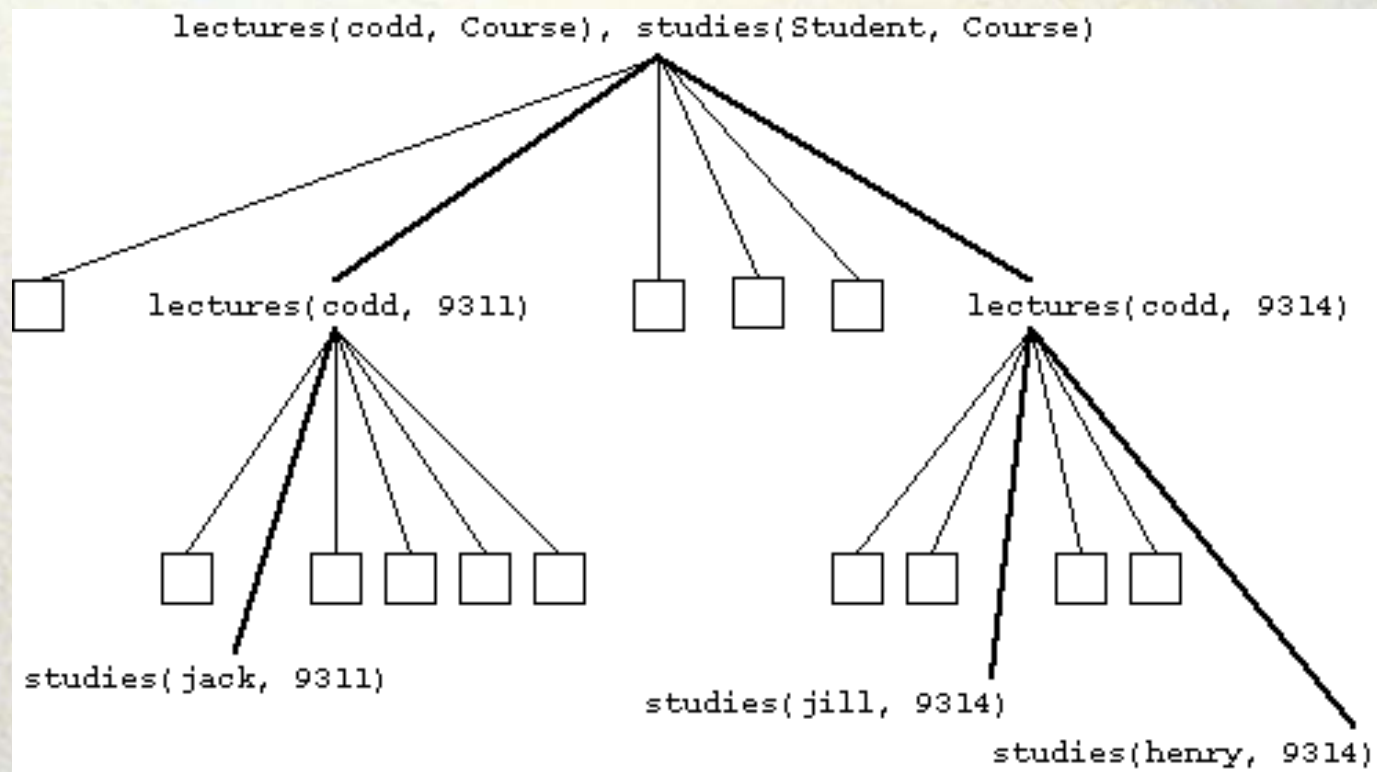


# Backtracking

- After the first solution is found, Prolog retraces its steps up the tree and looks for alternative solutions.
- First it looks for other students studying 9311 (but finds none).
- Then it
  - backs up
  - rebinds Course to 9314,
  - goes down the lectures(codd, 9314) branch
  - tries studies(Student, 9314),
  - finds the other two solutions:  
(Course = 9314, Student = jill)  
and (Course = 9314, Student = henry).

# Backtracking

*Proof tree:*





# Rules

- The previous question can be restated as a general rule:  
*One person, Teacher, teaches another person, Student if  
Teacher lectures in a course, Course and  
Student studies Course.*
- In Prolog this is written as:  
teaches(Teacher, Student) :-  
    lectures(Teacher, Course),  
    studies(Student, Course).  
?- teaches(codd, Student).
- Facts are *unit clauses* and rules are *non-unit clauses*.

# Clause Syntax

- ":-" means "if" or "is implied by". Also called the *neck* symbol.
- The left hand side of the neck is called the *head*.
- The right hand side of the neck is called the *body*.
- The comma, ",", separating the goals is stands for *and*.
- Another rule, using one of the *predefined predicate* ">".

more\_advanced(S1, S2) :-

year(S1, Year1),

year(S2, Year2),

Year1 > Year2.



# Tracing Execution

?- *trace*.

Yes

[trace] ?- *more\_advanced(henry, fred)*.

Call: more\_advanced(henry, fred) ?

Call: year(henry, \_L205) ?

Exit: year(henry, 4) ?

Call: year(fred, \_L206) ?

Exit: year(fred, 1) ?

^ Call: 4>1 ?

^ Exit: 4>1 ?

Exit: more\_advanced(henry, fred) ?

Yes

[debug] ?- *notrace*.

bind S1 to henry, S2 to fred  
test 1st goal in body of rule  
succeeds, binds Year1 to 4  
test 2nd goal in body of rule  
succeeds, binds Year2 to 1  
test 3rd goal: Year1 > Year2  
succeeds  
Succeeds

# More?

- Suppose we have the following facts and rule:

`bad_dog(fido).`

`bad_dog(Dog) :-`

`bites(Dog, Person),`

`is_person(Person),`

`is_dog(Dog).`

`bites(fido, postman).`

`is_person(postman).`

`is_dog(fido).`



# More?

There are two ways to prove `bad_dog(fido)`:

(a) it's there as a fact; and

(b) it can be proven using the `bad_dog` rule:

*?- bad\_dog(fido).*

More? ;

Yes

More? means Yes and prompts us to type ; if we want to check for another proof. The Yes that follows means that a second proof *was* found. Alternatively, we can just press the "return" key if we are not interested in whether there is another proof.

# Structures

- Functional terms can be used to construct complex data structures.
- If we want to say that John owns the novel Tehanu, we can write: `owns(john, 'Tehanu')`.
- Objects have a number of attributes:  
`owns(john, book('Tehanu', leguin))`.
- The author LeGuin has attributes too:  
`owns(john, book('Tehanu', author (leguin, ursula)))`.
- The arity of a term is the number of arguments it takes.
- all versions of `owns` have arity 2, but the detailed structure of the arguments changes.
- `gives(john, book, mary)`. is a term with arity 3.



# Asking Questions with Structures

- How do we ask, "What books does John own which were written by someone called LeGuin"?

?- *owns(john, book(Title, author(leguin, GivenName)))*.

Title = 'Tehanu' GivenName = ursula

- What books does John own?

?- *owns(john, Book)*.

Book = book('Tehanu', author(leguin, ursula))

- What books does John own?

?- *owns(john, book(Title, Author))*.

Title = 'Tehanu' Author = author(leguin, ursula)

- Prolog performs a complex matching operation between the structures in the query and those in the clause head.

# Library Database Example

- A database of books in a library contains facts of the form  
*book(CatalogNo, Title, author(Family, Given)).*  
*libmember(MemberNo, name(Family, Given), Address).*  
*loan(CatalogNo, MemberNo, BorrowDate, DueDate).*
- A member of the library may borrow a book.
- A "loan" records:
  - the catalogue number of the book
  - the number of the member
  - the date on which the book was borrowed
  - the due date



# Library Database Example

- Dates are stored as structures:  
date(Year, Month, Day)
- e.g. date(2008, 6, 16) represents 16 June 2008.
- which books has a member borrowed?

borrowed(MemFamily, Title, CatalogNo) :-

libmember(MemberNo, name(MemFamily, \_), \_),  
loan(CatalogNo, MemberNo, \_, \_),  
book(CatalogNo, Title, \_).

- The underscore or "don't care" variables (\_\_) are used because for the purpose of this query we don't care about the values in some parts of these structures.

# Comparing Two Terms

- we would like to know which books are overdue; how do we compare dates?

*%later(Date1, Date2) if Date1 is after Date2:*

*later(date(Y, M, Day1), date(Y, M, Day2)) :-  
Day1 > Day2.*

*later(date(Y, Month1, \_), date(Y, Month2, \_)) :-  
Month1 > Month2.*

*later(date(Year1, \_, \_), date(Year2, \_, \_)) :-  
Year1 > Year2.*

- This rule has three clauses: in any given case, only one clause is appropriate. They are tried in the given order.
- This is how disjunction (**or**) is often achieved in Prolog.



# Overdue Books

% overdue(Today, Title, CatalogNo, MemFamily):  
% given the date Today, produces the Title, CatalogNo,  
% and MemFamily of all overdue books.

overdue(Today, Title, CatalogNo, MemFamily) :-  
  loan(CatalogNo, MemberNo, \_, DueDate),  
  later(Today, DueDate),  
  book(CatalogNo, Title, \_),  
  libmember(MemberNo, name(MemFamily, \_), \_).

# Due Date

- Assume the loan period is one month, find the due date from today:

*%due\_date(Today, DueDate).*

*due\_date(date(Y, Month1, D), date(Y, Month2, D)) :-*

*Month1 < 12,*

*Month2 is Month1 + 1.*

*due\_date(date(Year1, 12, D), date(Year2, 1, D)) :-*

*Year2 is Year1 + 1.*



# The *is* operator

- The right hand argument of *is* must be an arithmetic expression that can be evaluated right now (no unbound variables).
- This expression is evaluated and bound to the left hand argument.
- *is* is not a C-style assignment statement:
  - $X \text{ is } X + 1$  won't work!
  - except via backtracking, variables can only be bound once, using *is* or any other way

# The *is* operator

- = does not cause evaluation of its arguments:

?-  $X = 2, Y = X + 1.$

$X = 2$

$Y = 2 + 1$

?-  $X = 2, Y \text{ is } X + 1.$

$X = 2$

$Y = 3$

- Use *is* if and only if you need to evaluate something:

$X \text{ is } 1$  **BAD!** - nothing to evaluate

$X = 1$  **GOOD!**




# Binary Trees

- In the library database example, some complex terms contained other terms, for example, book contained name.
- The following term also contains another term, this time one similar to itself:



`tree(tree(L1, jack, R1), fred, tree(L2, jill, R2))`

- The variables L1, L2, R1, and R2 should be bound to sub-trees.



# Recursive Structures


- A term that contains another term that has the same principal functor (in this case tree) is said to be recursive.
- Biological trees have leaves. For us, a *leaf* is a node with two empty branches:





# Another Tree Example

- `tree(tree(empty, 7, empty),  
 '+',  
 tree(tree(empty, 5, empty),  
 '-',  
 tree(empty, 3, empty)))`



# Recursive Programs for Recursive Structures

- A binary tree is either empty or contains some data and a left and right subtree which are also binary trees.

*is\_tree(empty).*

**trivial branch**

*is\_tree(tree(Left, Data, Right)) :-*

**recursive branch**

*is\_tree(Left),  
some\_data(Data),  
is\_tree(Right).*

- A non-empty tree is represented by a 3-arity term.
- Any recursive predicate must have:
  - (at least) one **recursive branch/rule** (or it isn't recursive :-)) and
  - (at least) one non-recursive or **trivial branch** (to stop the recursion going on for ever).



# Recursive Programs for Recursive Structures

- Let us define (or measure) the size of tree (i.e. number of nodes):

*tree\_size(empty, 0).*

*tree\_size(tree(L, \_, R), Total\_Size) :-*

*tree\_size(L, Left\_Size),*

*tree\_size(R, Right\_Size),*

*Total\_Size is Left\_Size + Right\_Size + 1.*

# Lists

- A list may be nil (i.e. empty) or it may be a term which has a head and a tail
- The head may be any term or atom.
- The tail is another list.
- We could define lists as follows:

`is_list(nil).`

`is_list(list(Head, Tail)) :-`

`is_list(Tail).`

- A list of numbers [1, 2, 3] would look like: `list(1, list(2, list(3, nil)))`
- Since lists are used so often, Prolog has a special notation:

`[1, 2, 3] = .(1, .(2, .(3, [])))`

`?- X = .(1, .(2, .(3, []))).`

`X = [1, 2, 3]`



# List Constructor I

- Within the square brackets [ ], the symbol | acts as an operator to construct a list from an item and another list.

?-  $X = [1 \mid [2, 3]]$ .

$X = [1, 2, 3]$ .

?-  $Head = 1$  ,  $Tail = [2, 3]$  ,  $List = [Head \mid Tail]$ .

$List = [1, 2, 3]$ .

# List Examples

?-  $[X, Y, Z] = [1, 2, 3]$ .

$X = 1 \ Y = 2 \ Z = 3$

?-  $[X \mid Y] = [1, 2, 3]$ .

$X = 1 \ Y = [2, 3]$

?-  $[X \mid Y] = [1]$ .

$X = 1 \ Y = []$

?-  $[X, Y \mid Z] = [\text{fred}, \text{jim}, \text{jill}, \text{mary}]$ .

$X = \text{fred} \ Y = \text{jim} \ Z = [\text{jill}, \text{mary}]$

?-  $[X \mid Y] = [[a, f(e)], [n, m, [2]]]$ .

$X = [a, f(e)] \ Y = [[n, m, [2]]]$



# List Membership

- A term is a member of a list if
  - the term is the same as the head of the list, or
  - the term is a member of the tail of the list.

- In Prolog:

```
member(X, [X | _]).
```

```
member(X, [_ | Y]) :-  
    member(X, Y).
```

- Member is actually predefined in Prolog.

# Concatenating Two Lists

- Suppose we want to take two lists, like [1, 3] and [5, 2] and concatenate them to make [1, 3, 5, 2]

`concat([], L, L).`

`concat([Item | Tail1], L, [Item | Tail2]) :-`

`concat(Tail1, L, Tail2).`



# An Application of Lists

- Find the total cost of a list of items:

*% cost data:*

*cost(cornflakes, 230).*

*cost(cocacola, 210).*

*cost(chocolate, 250).*

*cost(crisps, 190).*

?- *total\_cost([cornflakes, crisps], X).*

*X = 420*

# An Application of Lists

`total_cost([], 0).`

`total_cost([Item|Rest], Cost) :-`

`cost(Item, ItemCost),`

`total_cost(Rest, CostOfRest),`

`Cost is ItemCost + CostOfRest.`

- How about if we change the recursive branch:

`total_cost([Item|Rest], Cost) :-`

`total_cost(Rest, CostOfRest),`

`cost(Item, ItemCost),`

`Cost is ItemCost + CostOfRest.`



# Negation as Failure

- Build-in predicate not.

?- not(lectures(turing, 9020)).

Not/fail

# Remove duplicates

?- *remove\_dups*([1,2,3,1,3,4], X).

X = [2, 1, 3, 4]

% *remove\_dups*(+List, -NewList):

*remove\_dups*([], []).

*remove\_dups*([First | Rest], NewRest) :-

member(First, Rest),

*remove\_dups*(Rest, NewRest).

*remove\_dups*([First | Rest], [First | NewRest]) :-

not(member(First, Rest)),

*remove\_dups*(Rest, NewRest).



# Controlling Execution

## The Cut Operator

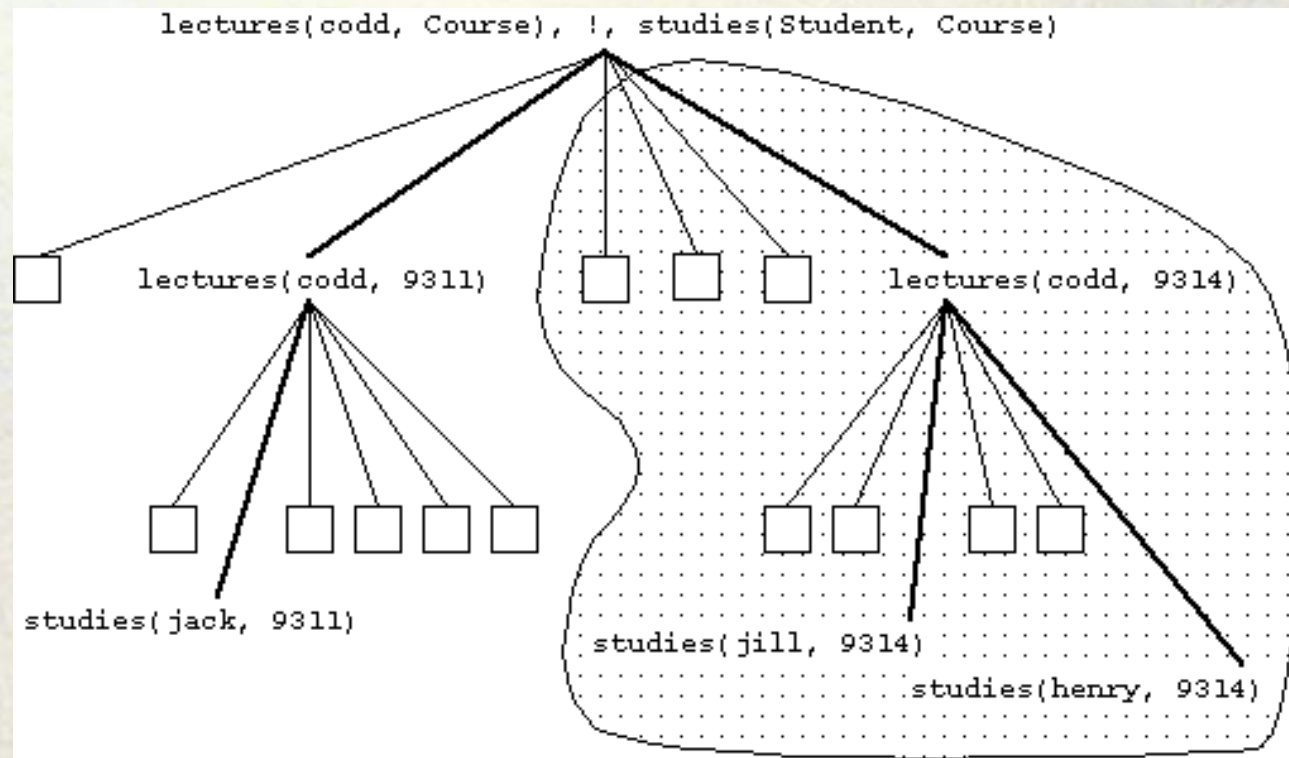
- Sometimes we need a way to prevent Prolog finding all solutions, i.e. a way to stop backtracking.
- The cut operator, written `!`, is a built-in goal that prevents backtracking.
- It turns Prolog from a nice declarative language into a hybrid monster.

# The Cut Operator !

- Cut prunes the search tree, prevents backtracking:
  - Once the cut operator has been passed when evaluating a predicate, no new variable instantiations are allowed to those variables which are bound at that point in time.
  - Uninstantiated variables can still be instantiated after the cut operator has been processed.
  - Backtracking can still take place, but only for those uninstantiated variables.
- If the goal(s) to the right of the cut fail then the entire clause fails and the goal which caused this clause to be invoked fails.
- In particular, alternatives for Course are not explored.



# The Cut operator



# Cut example

?- *lectures(codd, X).*

X = 9311 ;

X = 9314 ;

No

?- *lectures(codd, X), !.*

X = 9311.



# Cut example

- max, without cut:

% max(A, B, C) binds C to the larger of A and B.

max(A, B, A) :-

A > B.

max(A, B, B) :-

A =< B.

- max, with cut:

max(A, B, A) :-

A > B, !.

max(A, B, B).

# Cut example

- remove\_dups, with cut:

remove\_dups([], []).

remove\_dups([First | Rest], NewRest) :-

member(First, Rest),

!,

remove\_dups(Rest, NewRest).

remove\_dups([First | Rest], [First | NewRest]) :-

remove\_dups(Rest, NewRest).



# Exercises

- Reversing Lists:
  - `reverse(A, B)`: B is the reversed list of A.
- Viết chương trình duyệt theo chiều rộng xem có đường đi từ 1 đỉnh này đến 1 đỉnh khác (adjacency list representation `adjacent(1, [2, 4, 5]).`):
  - `reachable(A, B)`: Có đường đi từ A đến B không?
  - `path(A, B, Path)`: in ra đường đi từ A đến B.
- Missionaries and Cannibals