
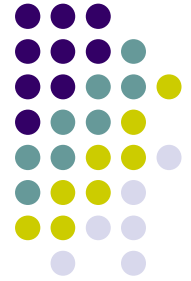


Operating System

Nguyen Tri Thanh
ntthanh@vnu.edu.vn





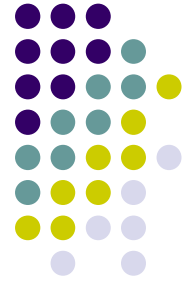
Review

- A system uses FCFS process (arrived_time, duration)
 - $P_1(0,20)$, $P_2(30,10)$, $P_3(20,40)$, $P_4(50,15)$
- Which of the following is the correct running order of the above processes?
 - A. P_1, P_2, P_3, P_4
 - B. P_1, P_3, P_2, P_4
 - C. P_1, P_4, P_2, P_3
 - D. P_5, P_2, P_3, P_1




Review

- A system uses SJF process (arrived_time, duration)
 - $P_1(0,20)$, $P_2(30,10)$, $P_3(20,40)$, $P_4(50,15)$
- Which of the following is the correct running order of the above processes?
 - A. P_1, P_2, P_3, P_4
 - B. P_1, P_4, P_2, P_3
 - C. P_1, P_3, P_2, P_4
 - D. P_4, P_2, P_3, P_1



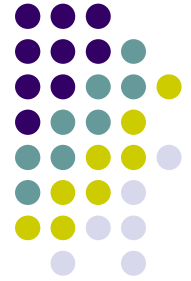
Review

- A system uses SRTF process (arrived_time, duration)
 - $P_1(0,20)$, $P_2(30,10)$, $P_3(20,40)$, $P_4(40,15)$
- Which of the following is the correct running order of the above processes?
 - A. P_1, P_3, P_2, P_4, P_3
 - B. P_1, P_2, P_3, P_4, P_4
 - C. P_1, P_4, P_2, P_3, P_2
 - D. P_1, P_2, P_3, P_1, P_4



Review

- A system uses RR process (arrived_time, duration)
 - $P_1(0,20)$, $P_2(30,10)$, $P_3(20,40)$, $P_4(40,25)$
 - Time quantum = 15
- Which of the following is the correct **running order** of the above processes?
 - A. $P_1, P_2, P_3, P_1, P_2, P_3, P_4, P_3$
 - B. $P_1, P_3, P_1, P_3, P_2, P_3, P_4, P_3$
 - C. $P_1, P_1, P_2, P_3, P_2, P_3, P_4, P_3$
 - D. $P_1, P_1, P_3, P_2, P_4, P_3, P_4, P_3$



Review

- A system uses RR process (arrived_time, duration)
 - $P_1(0,20)$, $P_2(30,10)$, $P_3(20,40)$, $P_4(40,25)$
 - Time quantum 15
- Which of the following is the correct total **waiting time** of the above processes?
 - A. 40
 - B. 50
 - C. 60
 - D. 70




Inter-process Communication (IPC)




Objectives

- Present what IPC is
- Write a simple IPC program in Linux
- Present why we need synchronization
 - Methods of synchronization
- Write a simple synchronization program

Reference



- Chapter 2, 6 of **Operating System Concepts**




Introduction

- In some situations, processes need to communicate with each other
 - To send/receive data (web browser – web server)
 - To control the other process
 - To synchronize with each other
- This can be done by IPC
- IPC is implemented differently among OSes
 - Linux: message queue, semaphore, shared segment, ...



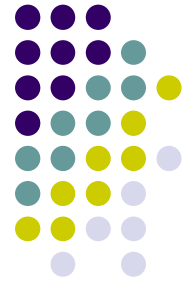
Introduction (cont'd)

- IPC can be divided into 2 categories
 - IPC among processes within the same system
 - Linux: pipe, named pipe, file mapping, ...
 - IPC among processes in different systems
 - Remote Procedure Call (RPC), Socket, Remote Method Invocation (RMI), ...




Process Synchronization

Synchronization definition




- **Process synchronization** refers to the idea that multiple processes are to join up or **handshake at a certain point**, in order to reach an agreement or **commit to a certain sequence of actions**.
 - [http://en.wikipedia.org/wiki/Synchronization_\(computer_science\)](http://en.wikipedia.org/wiki/Synchronization_(computer_science))

Team Work



another example ...



Problem

Write process *P*:

```
while (true) {  
    val=buf;  
    val += count();//Take time  
    buf=val  
}
```

buf: Buffer

```
UPDATE A SET  
    buf=buf+count();
```

What if more than
one *P* are
running?



Problem (cont'd)

- Two concurrent processes

```
val=buf;  
val += count();  
buf=val
```

```
val=buf;  
val += count();  
buf=val
```

Do we always get the expected value of **buf**? Why?



Problem (cont'd)

- Suppose buf=5

val=buf;

//val=5

val+=count();

//val=10

val=buf;

//val=5

val+=count();

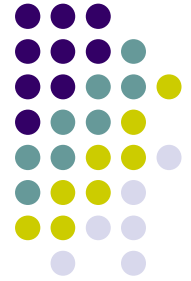
//val=10

buf=val

//buf=10

buf=buf

//buf=10



Problem (cont'd)


- Cause: P and Q simultaneously operate on global variable **buf**
- Solution: Let them operate **separately**

val=buf;	//val=5
val+=count();	//val=10
buf=val	//buf=10

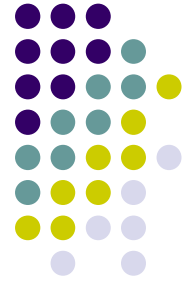
val=buf;	//val=10
val+=count();	//val=15
buf=val	//buf=15

Race condition

- Happen when many processes simultaneously work with shared data




To avoid “trouble”, processes need to be controlled



Critical section

- In [concurrent programming](#) a **critical section** is a piece of [code](#) that accesses a shared resource (data structure or device) that must not be concurrently accessed by more than one [thread of execution](#). A critical section will usually terminate in fixed time, and a thread, task or process will have to wait a fixed time to enter it (aka bounded waiting). Some [synchronization](#) mechanism is required at the entry and exit of the critical section to ensure exclusive use, for example a [semaphore](#).
- (http://en.wikipedia.org/wiki/Critical_section)





Critical section

- Suppose n processes P_1, \dots, P_n share a global variable v
 - v can also be other resource, e.g, file
- Each process has a segment of code CS_i which operates on v
 - CS_i is called **critical section**
 - Because it is critical to prone errors
 - CS_i should be the **smallest** code segment
- Need to make the critical section safe



Critical section





Question

Process *P*:

```
while (true) {  
    waitForNewRequest();  
    if(found){  
        hit+=1;  
        val=hit;  
    }  
    Respond();  
}
```

hit: a global variable

Which is the critical section of the code when multiple processes of *P* run?

```
while (true) {  
    waitForNewRequest();  
    if(found){  
        hit+=1;  
        val=hit;  
    }  
    Respond();  
}
```

}



Critical section (cont'd)

- Common structure

do {

 Enter_Section (CS_i);

 Run CS_i ;

 Exit_Section(CS_i);

 Run ($REMAIN_i$); // Remainder section

} while (TRUE);



Critical section (cont'd)

- Short description

do {

ENTRY_i; // Enter section

Run *CS_i*; // Critical section

EXIT_i; // Exit section

REMAIN_i; // Remainder section

} while (TRUE);

Implementation of Critical section



Implementation must satisfy 3 conditions

1. Mutual Exclusion

- If a process is in its critical section, then **no other** processes can be in their critical sections

2. Progress

- If **no** process is in its critical section
- other processes waiting to enter their critical section,
- then the selection of the process to enter the critical section cannot be postponed **indefinitely**

3. Bounded Waiting


- No process has to wait **indefinitely** to enter its critical section

Question



Which is the purpose of the first condition?

- A. It supports the priority of process
- B. It ensures the correct use of the shared resource
- C. It tries to utilize the shared resource effectively
- D. It makes the implementation of OS simpler



Question

Which is the consequence of the second condition?


- A. It reduces the waiting time of requested processes
- B. It ensures the correct use of the shared resource
- C. It supports the priority of processes
- D. It makes the implementation of OS simpler



Question

Which is the consequence of the second condition?

- A. It supports the priority of processes
- B. It ensures the correct use of the shared resource
- C. It utilizes the shared resource effectively
- D. It makes the algorithm complicated to implement



Question

Which is the consequence of the 3rd condition?

- A. It supports the priority of processes
- B. It ensures the correct use of the shared resource
- C. It utilizes the shared resource effectively
- D. It makes sure no process can never enter its critical section


Question



Which is the correct conditions of critical section?

- A. mutual exclusion, protection, bounded using
- B. mutual exclusion, protection, bounded waiting
- C. mutual exclusion, progressive, bounded waiting
- D. mutual exclusion, bounded waiting, progress


Question



Which is the correct purpose the 2nd condition of critical section?

- A. maximize CPU utilization
- B. maximize the shared resource utilization
- C. maximize disk utilization
- D. maximize RAM utilization

Question



Which is the consequence of the 3rd condition?

- A. It supports the priority of processes
- B. It ensures the correct use of the shared resource
- C. It ensures the relative fairness of processes to use the shared resource
- D. It utilizes the shared resource effectively

The fairness



The fair exam today is to swim

Critical section (cont'd)




- Each process has to
 - request to run (enter section) its critical section CS_i
 - and announce its completion (exit section) of its CS_i .

Peterson's Solution



- Solution for two processes
- The two processes share two variables:
 - int **turn**; // with the value of 0 or 1
 - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section
 - If **turn==i** then P_i is in turn to run its CS_i
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process P_i is ready!



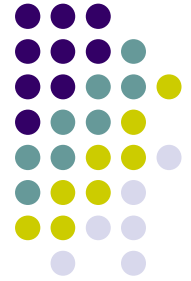
Peterson's solution (cont'd)

- Program P_i :
do {
 flag[i] = TRUE;
 turn = j;
 while (flag[j] && turn == j) ;
 CS_i;
 flag[i] = FALSE;
 REMAIN_i;
} while (1);

Peterson's solution (cont'd)





- The proof of this solution is provided on page 196 of the textbook
- Comments
 - Complicated when the number of processes increases
 - Difficult to control



Question

- Which code snippet is Enter_Section?
 - A. `flag[i] = TRUE;`
 `turn = j;`
 `while (flag[j] && turn == j) ;`
 - B. `flag[i] = TRUE;`
 `while (flag[j] && turn == j) ;`
 - C. `flag[i] = TRUE;`
 `turn = j;`
 - D. `turn = j;`
 `while (flag[j] && turn == j) ;`

Semaphore



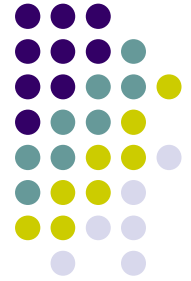


Reference information

- Semaphore is proposed by Edsger Wybe Dijkstra (Dutch) for Computer Science in 1972
- Semaphore was firstly used in his book “The operating system”



Edsger Wybe Dijkstra
(1930-2002)



Semaphore

- Semaphore is **an integer**, can be only access through two **atomic operators** wait (or P) and signal (or V).
 - P: proberen – check (in Dutch)
 - V: verhogen – increase (in Dutch)
- Processes can share a semaphore
- **Atomic operators** guarantee the **consistency**



wait and signal operators

wait(S) // or P(S)

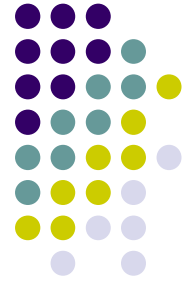
```
{  
    while (S<=0);  
    S--;  
}
```

- Wait if semaphore $S \leq 0$ else decrease S by 1

signal(S) // or V(S)

```
{  
    S++;  
}
```

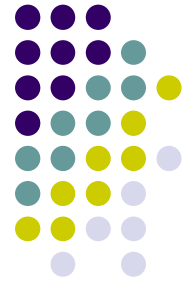
- Increase S by 1

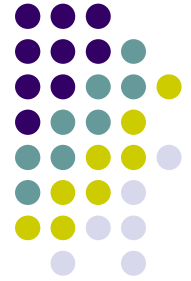


Using semaphore

- Apply for critical section
do {
 wait(s); // s is a semaphore initialized by 1
 CS_i;
 signal(s);
 REMAIN_i;
} while (1);

Semaphore





Question

Process *P*:

```
while (true) {  
    waitForNewRequest();  
    if(found){  
        hit+=2;  
        val=hit;  
    }  
    Respond();  
}
```

hit: a global variable


Use semaphore to make the code safe?

```
if(found){  
    wait(mutex);  
    hit+=2;  
    val=hit;  
    signal(mutex);  
}
```



Using semaphore (cont'd)

- P_1 needs to do O_1 ; P_2 need to do O_2 ; O_2 can only be done **after** O_1
- Solution: use a semaphore $synch = 0$



Semaphore implementation



- In the above semaphore implementation
 - Use **busy waiting** (while loop)
 - Resource wasting
- Atomic operators
 - When a process called wait(), it will be **blocked** if the semaphore is not free
 - This type of semaphore is called *spinlock*
 - Other wait() implementation just returns true/false and **does not block** the calling process

Semaphore implementation (cont'd)



- Remove the busy waiting loop by using **block**
- To restored a blocked process, use **wakeup**
- Semaphore data structure

```
typedef struct {  
    int value; // value of semaphore  
    struct process *L; //waiting process list  
} semaphore;
```


Semaphore implementation (cont'd)




```
void wait(semaphore *S)
{
    S->value--;
    if (S->value < 0) {
        Add the requested
        process  $P$  into s->L;
        block( $P$ );
    }
}
```

```
void signal(semaphore *S)
{
    S->value++;
    if (S->value <= 0) {
        remove a process  $P$ 
        from s->L;
        wakeup( $P$ );
    }
}
```


Semaphore implementation (cont'd)

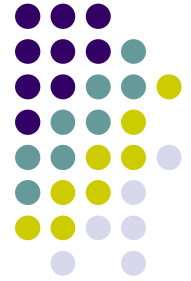




Binary semaphore

- Semaphore only has the value of 0 or 1
- Other semaphore type is **counting semaphore**






Question

- When counting semaphores are suitable to use?
 - A. When 2 processes share a single variable/resource
 - B. When 3 processes share a single variable/resource
 - C. When n processes share a single variable/resource
 - D. When n processes share m variables/resources of the same type

Classical synchronization problems





Bounded-Buffer Problem

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value N .

Bounded-buffer problem (cont'd)



Write process *P*:

```
do {  
    wait(empty);  
    wait(mutex);  
    Write (item);  
    signal(mutex);  
    signal(full);  
} while (TRUE);
```

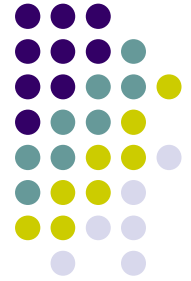
Read process *Q*:

```
do {  
    wait(full);  
    wait(mutex);  
    Read(item);  
    signal(mutex);  
    signal(empty);  
} while (TRUE);
```




Question

- Which is the initialized value of the *full* variable in the above algorithm?
 - A. -1
 - B. 0
 - C. 1
 - D. NULL

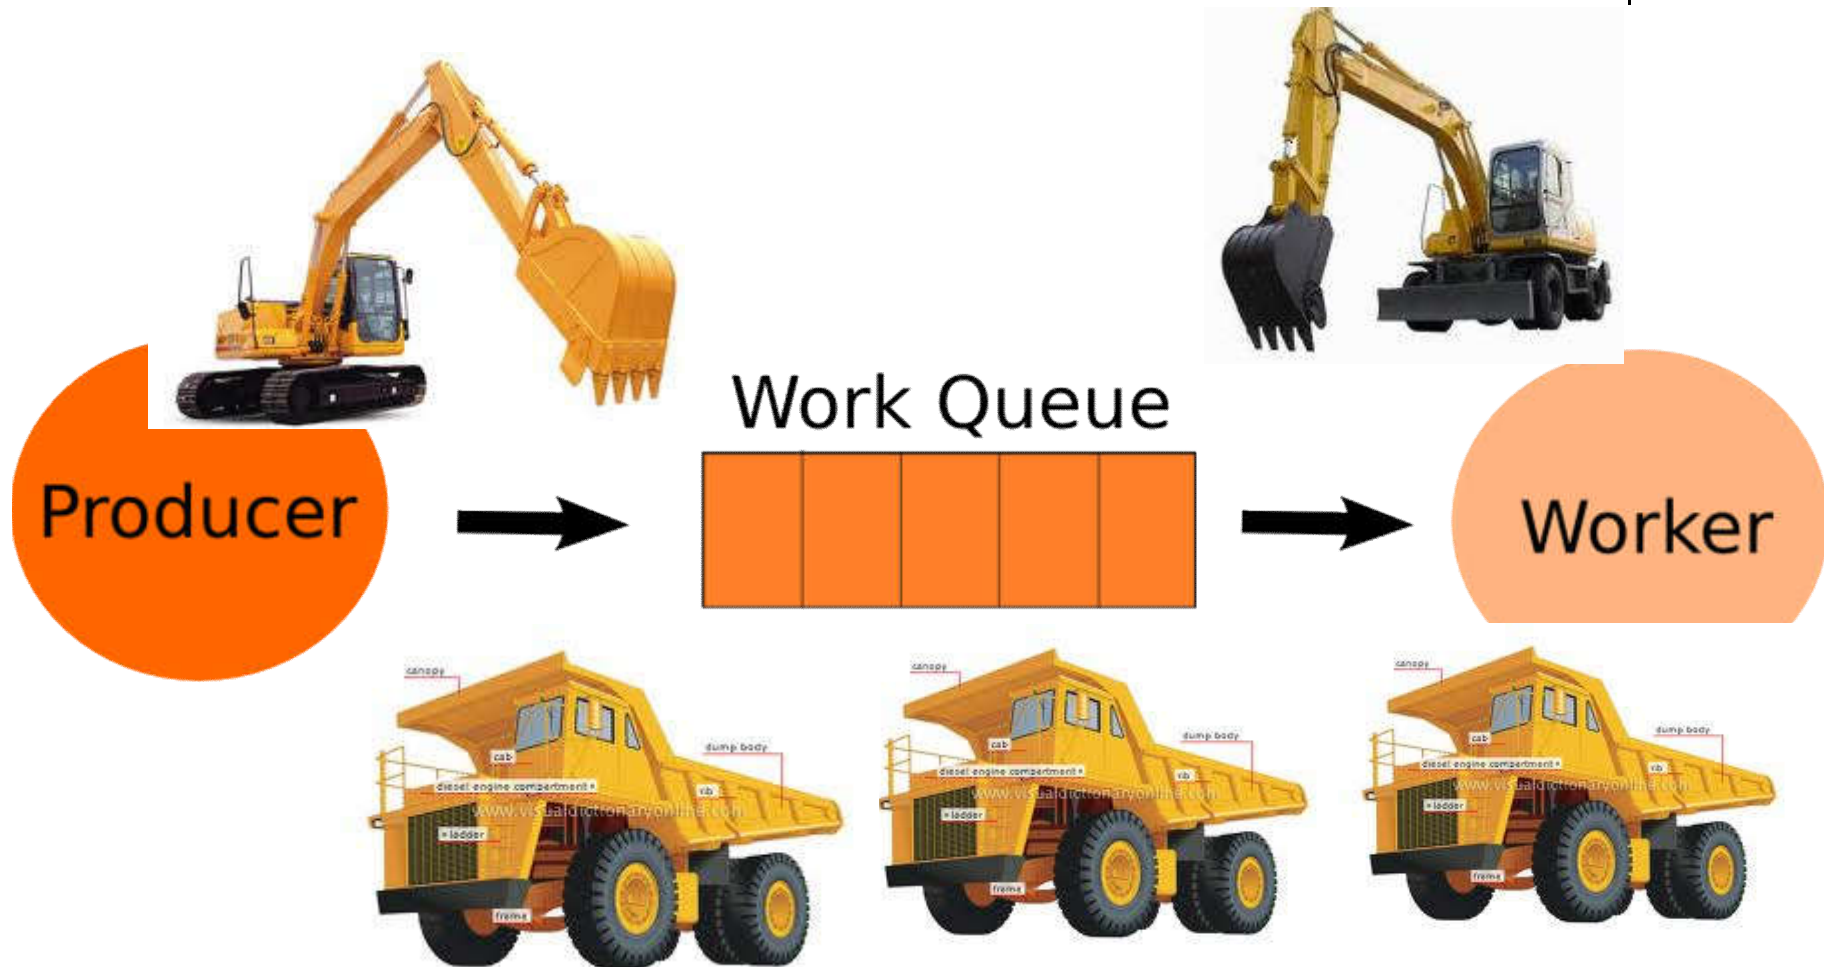


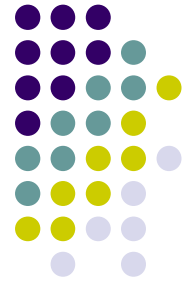
Question

What will be the problem if the initialized value of the *full* variable is 1?

- A. no problem at all
- B. the writer process can not run
- C. the reader process can not run
- D. the reader can read an invalid value

Bounded-buffer problem (cont'd)

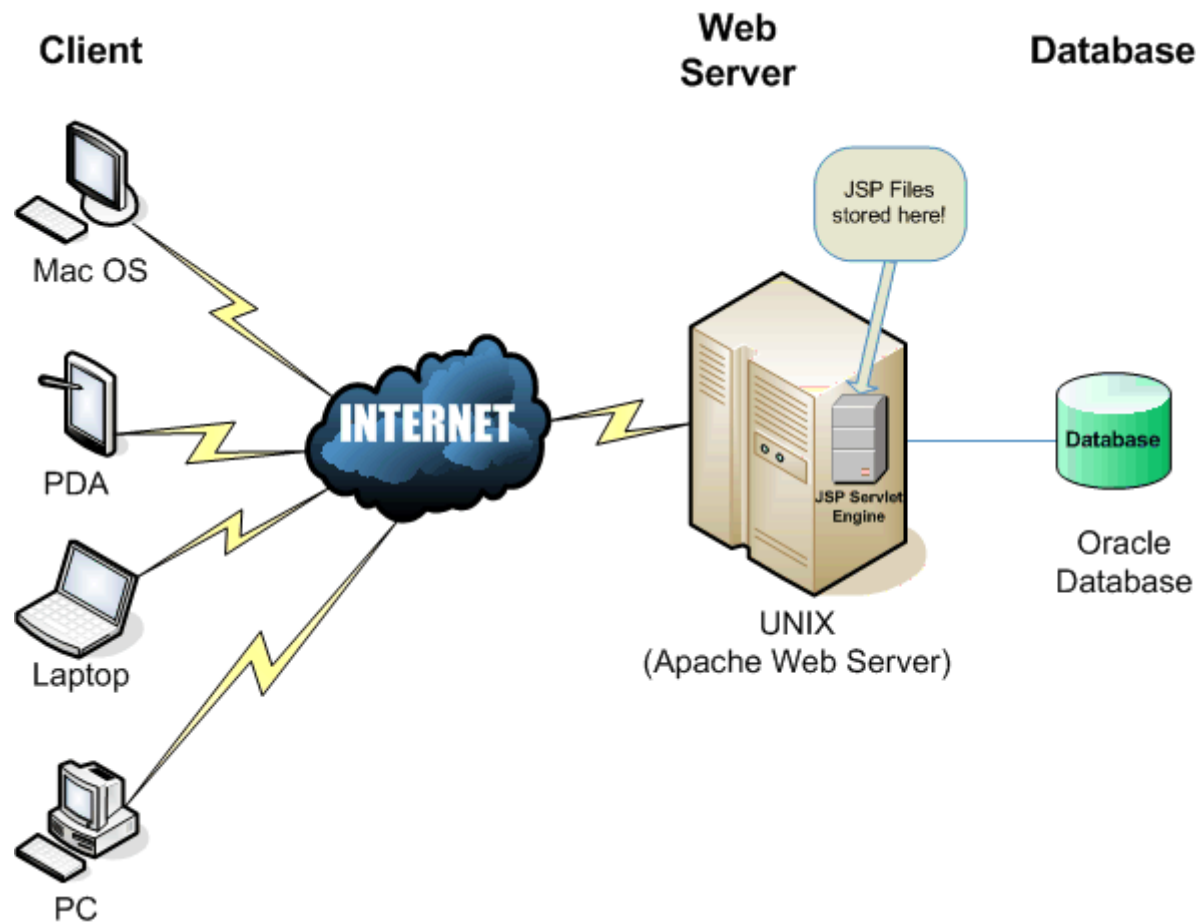




Readers-writers problem

- A **data set** is shared among a number of concurrent processes
 - Readers – only **read**
 - Writers – can both **read and write**
- Problem
 - allow **multiple** readers to read at the same time
 - Only **one** writer can access the shared data at the a time

Readers-writers problem



Readers-writers problem (cont'd)



- Shared data
 - Data set
 - Semaphore *wrt* initialized by 1
 - Used to manage **write** access
 - Integer *readcount* initialized by 0 to count the number of readers that are **reading**
 - Semaphore *mutex* initialized by 1
 - Used to manage readcount access

Readers-writers problem (cont'd)



- **Process writer P_w :**

```
do {  
    wait(wrt);  
    write(data_set);  
    signal(wrt);  
}while (TRUE);
```

- **Process reader P_r :**

```
do {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1) wait(wrt);  
    signal(mutex);  
    read(data_set);  
    wait(mutex);  
    readcount--;  
    if (readcount == 0) signal(wrt);  
    signal(mutex);  
} while (TRUE);
```



Question

Why do we need *readcount* variable?

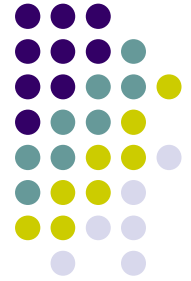
- A. We may remove this variable
- B. To make sure there is one reader at a time
- C. To make sure no readers are reading
- D. To make sure no readers are reading before writing



Question

Which is the initialized value of the *readcount* variable in the above algorithm?

- A. -1
- B. 0
- C. 1
- D. NULL



Question

Which is the purpose of *mutex* variable?

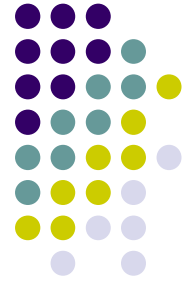
- A. To safely access the `data_set`
- B. We may remove this variable without affecting the program
- C. To safely access the *readcount* variable
- D. To safely access the *wrt* variable



Question

Which is the initialized value of the *mutex* variable in the above algorithm?

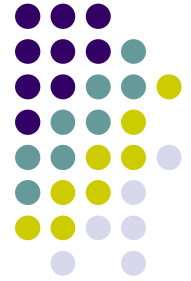
- A. -1
- B. 0
- C. 1
- D. NULL



Question

Which is the purpose of *wrt* variable?

- A. To safely access the *mutex* variable
- B. To safely write the *data_set*
- C. To safely write the *readcount* variable
- D. To safely read the *data_set*



Question


Which is the initialized value of the *wrt* variable in the above algorithm?

- A. -1
- B. 0
- C. 1
- D. NULL

Dining-Philosophers Problem



- Five philosophers at a table having 5 chopsticks, 5 bowls and a rice cooker
- A philosopher just eats or thinks
- How to make sure philosophers correctly use the “shared data” – the chopsticks




Dining-philosophers problem (cont'd)




- Use semaphore to handle chopstick access
 - semaphore chopstick[5];
 - Solution is provided as in the next text box
- Code of philosopher i :
do {
 wait(chopstick[i]);
 wait(chopstick[(i+1)%5]);
 Eat(i);
 signal(chopstick[i]);
 signal(chopstick[(i+1)%5]);
 Think(i);
} while (TRUE);

Question




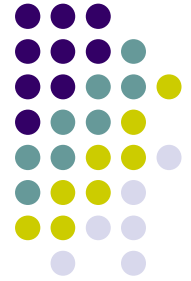
- What value chopstick[i] is initialized?
 - A. 1
 - B. 2
 - C. 0
 - D. 5



Question

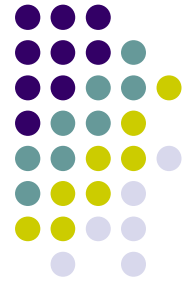
- Is there any problem with the solution?
 - A. No problem
 - B. Only one philosopher can eat at a time
 - C. Only three philosophers can eat at a time
 - D. No philosopher could eat in case each takes a chopstick and waits for the second one





Question

- Which of the following is **incorrect** about the solution to the above problem?
 - A. No solution available
 - B. Create an order of philosophers to eat
 - C. Create an order of philosophers to think
 - D. Allow at most 4 philosophers to request to eat at a time



Limitations of semaphore (cont'd)

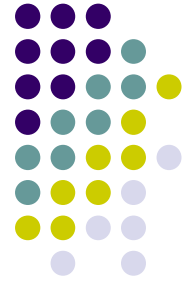
- Compare the two code snippets

- Snippet 1

```
...  
wait(mutex);  
//Critical section  
signal(mutex);  
...
```

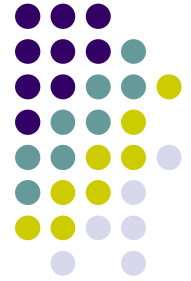
- Snippet 2

```
...  
signal(mutex);  
//Critical section  
wait(mutex);  
...
```



Question

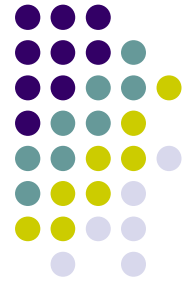
- What is the problem of the two code snippets?
 - A. Snippet 1 has problem
 - B. Snippet 2 has problem
 - C. Both snippets have problem
 - D. No problem at all




Question

- Which is the problem of the incorrect use of semaphore in the above code snippet?
 - A. No process can enter its critical section
 - B. No problem at all
 - C. The mutual exclusion condition may be violated
 - D. No process can exit its critical section

Limitations of semaphore



- Semaphores need correct calls to **wait** and **signal**
- Incorrect use of semaphore may lead to **deadlock**
- Even **correct use** of semaphores may lead to deadlock, in some cases



Limitations of semaphore (cont'd)

- Compare the two code snippets

- Snippet 1

...

```
wait(mutex);
```

```
    CS1;
```

```
wait(mutex);
```

...

- Snippet 2

...

```
wait(mutex);
```

```
CS2;
```

```
signal(mutex);
```

...



Question

- Which of the two code snippets has problem?
 - A. Snippet 1 has problem
 - B. Snippet 2 has problem
 - C. Both snippets have problem
 - D. No problem at all



Question

- Which is the consequence of the above problem?
 - A. One process will be blocked
 - B. There will be a deadlock
 - C. No consequences if only two processes are involved
 - D. No consequences



Limitations of semaphore (cont'd)

- Process P_1

...

wait(S);

wait(Q);

...

signal(S);

signal(Q);

- Process P_2

...

wait(Q);

wait(S);

...

signal(Q);

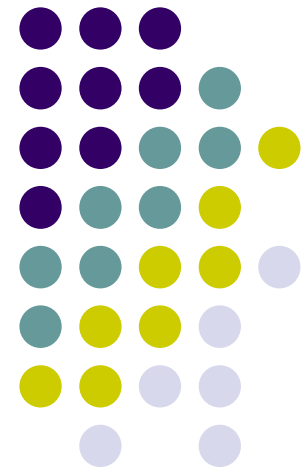
signal(S);

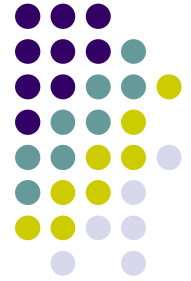


Question

- What is the problem of the above two processes?
 - A. There is deadlock
 - if P_1 got S and waits for Q and
 - P_2 got Q and waits for S
 - B. The exclusive condition is violated
 - C. The order of semaphore calls is incorrect
 - D. No problem at all

Monitor



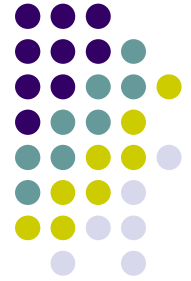


Reference information

- Per Brinch Hansen (Dennish) proposed the concept and implemented in 1972
- Monitor was firstly used in Concurrent Pascal programming language



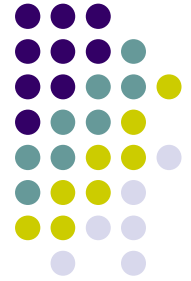
Per Brinch Hansen
(1938-2007)



What is monitor?

- Monitor means to **supervise**
- It is a type of construct in a high level programming language for synchronization purpose
 - C# programming language
 - <http://msdn.microsoft.com/en-us/library/hf5de04k.aspx>
 - Java programming language
 - <http://www.artima.com/insidejvm/ed2/threadsynch.html>
 - <http://journals.ecs.soton.ac.uk/java/tutorial/java/threads/monitors.html>
- Monitor was studied and developed to overcome the limitations of semaphores

Monitor




- A monitor usually has
 - Member variables as shared resources
 - A set of procedures which operate on the shared resources
 - Exclusive lock
 - Constraints to manage race condition
- This description of monitor is like a class

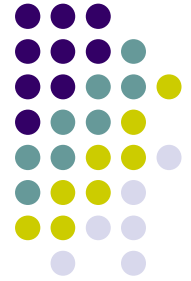


A sample monitor type

```
monitor monitor_name {  
    //Shared resources  
    procedure P1(...) { ...  
    }  
    procedure P2(...) { ...  
    }  
    ...  
    procedure Pn(...) { ...  
    }  
    initialization_code (..) { ...  
    }  
}
```


Schematic view of a Monitor





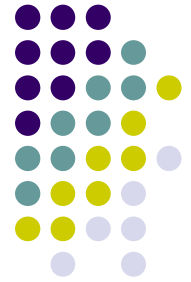
Monitor implementation

- Monitor must be implemented so that
 - only one process can enter the monitor at a time (mutual exclusive)
 - programmer do not need to write code for this
- Other monitor implementation
 - have more synchronization mechanism
 - add *condition variable*

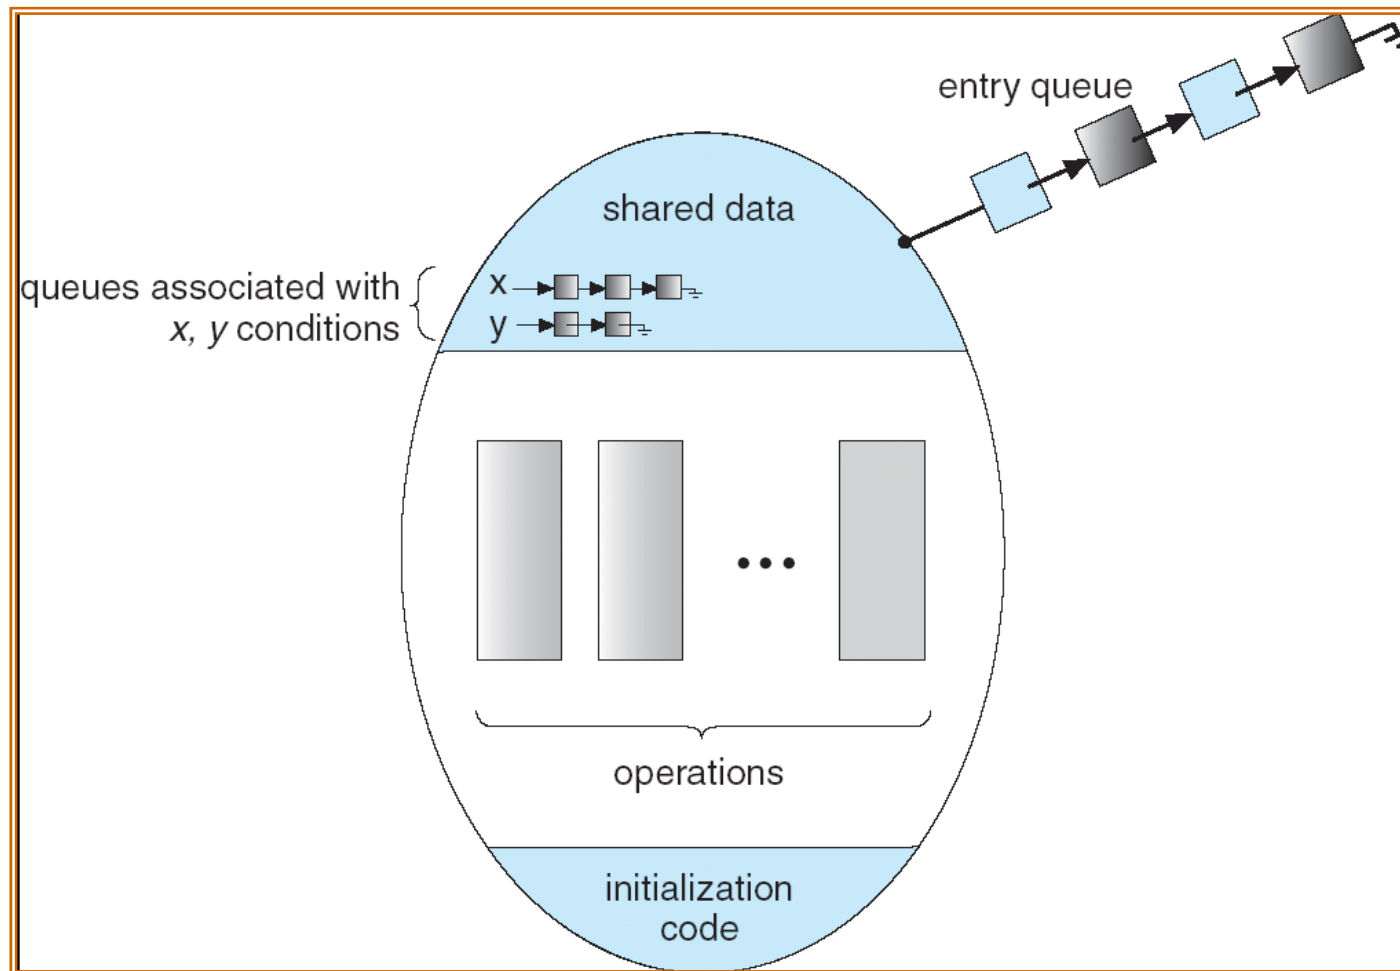


Condition type

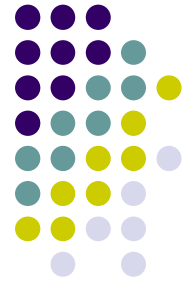
- Declaration
 - condition x, y;
- Use condition variable
 - there are two operators: wait and signal
 - x.wait():
 - process calls x.wait() will have to wait or suspend
 - x.signal():
 - process calls x.signal() will wakeup a waiting process
 - the one that called x.wait()



Monitor with condition



x.signal() characteristics



- x.signal() wakeup only one waiting process
- If no waiting process, it does nothing
- x.signal() is different from that of classical semaphore
 - signal in classical semaphore always change the state (value) of semaphore

Solution to Dining Philosophers



monitor DP

```
{  
    enum { THINKING; HUNGRY, EATING) state [5] ;  
    condition self [5];
```

```
void pickup (int i) {  
    state[i] = HUNGRY;  
    test(i);  
    if (state[i] != EATING) self [i].wait;  
}
```

```
void putdown (int i) {  
    state[i] = THINKING;  
    // test left and right neighbors  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}
```

Solution to Dining Philosophers (cont)



```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

```
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
}
```

Solution to Dining Philosophers (cont)



- Each philosopher invokes the operations `pickup()` and `putdown()` in the following sequence

`dp.pickup (i)`

EAT

`dp.putdown (i)`

Monitor Implementation Using Semaphores



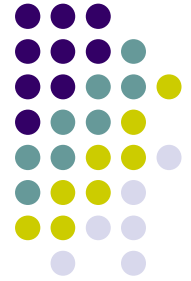
- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
```

- Each procedure F will be replaced by

```
wait(mutex);
...
body of  $F$ ;
...
if (next-count > 0)
    signal(next)
else
    signal(mutex);
```

- Mutual exclusion within a monitor is ensured.



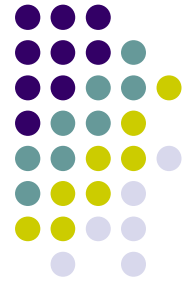
Monitor Implementation

- For each condition variable x , we have:

```
semaphore x-sem; // (initially = 0)
int x-count = 0;
```

- The operation $x.\text{wait}$ can be implemented as:

```
x-count++;
if (next-count > 0)
    signal(next);
else
    signal(mutex);
wait(x-sem);
x-count--;
```




Monitor Implementation

- The operation `x.signal` can be implemented as:

```
if (x-count > 0) {  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```

Linux Synchronization



- Linux:
 - disables interrupts to implement short critical sections
- Linux provides:
 - semaphores
 - spin locks

Pthreads Synchronization



- pthread API is OS-independent
- It provides:
 - mutex locks
 - condition variables
- Non-portable extensions include:
 - read-write locks
 - spin locks



Question?