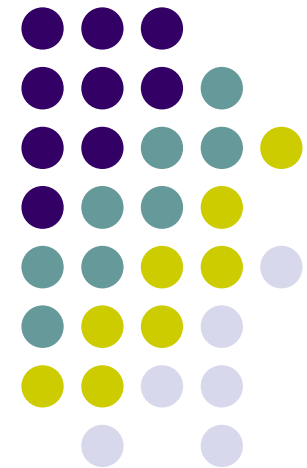


# Operating System

---

**Nguyen Tri Thanh**  
**ntthanh@vnu.edu.vn**

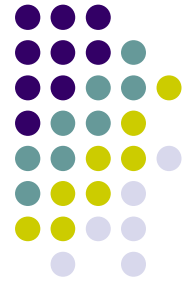


# Review



Which is correct about race condition?

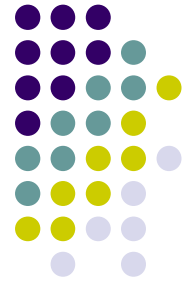
- A. Happen even when there is only once process
- B.** Happen when multiple processes use a shared resource concurrently
- C. Happen when multiple processes use a resource sequentially
- D. Happen when there are multiple processes in the system



# Review

Which is incorrect about the Peterson's solution?

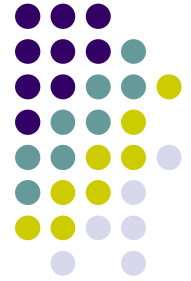
- A. It satisfies all the conditions of critical section
- B. It is easy to control even the number of processes is above 2
- C. It is difficult to control
- D. It is complicated when the number of processes is above 2



# Review

Which of the following is the most correct about critical section?

- A. A code snippet that operates on a global variable
- B. A code snippet that operates on a resource
- C. A code snippet that operates on a global resource
- D.** A code snippet that operates on a shared resource



# Review

How many conditions for resolving critical section are there ?

- A. 1
- B. 2
- C. 3**
- D. 4



# Review

Which is incorrect about the conditions of critical section?

- A. The progress condition utilizes the resource effectively
- B. The exclusive condition removes race condition**
- C. The exclusive condition ensures processes to use a shared resource sequentially
- D. The bounded waiting condition allows a process to use a shared resource several consecutive times

# Question



Which is the purpose of the second condition of critical section?

- A. It reduces the waiting time of requested processes
- B. It ensures the correct use of the shared resource
- C. It makes the algorithm more complicated to implement
- D. It makes the algorithm less complicated to implement

# Question



Which is the purpose of the third condition of critical section?

- A. It supports the priority of processes
- B. It ensures the correct use of the shared resource
- C. It utilizes the shared resource effectively
- D.** It makes sure no process is in its critical section forever



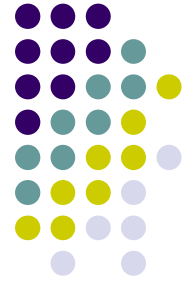


# Review

Which is incorrect about the semaphore?

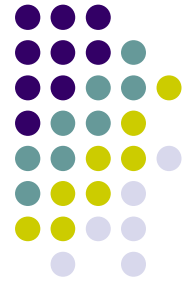
- A. Semaphore is an implementation of critical section
- B.** Semaphore does not guarantee the conditions of critical section
- C. A semaphore usually includes an integer variable
- D. Semaphore has atomic operators

# Review



How many types the semaphore are there?

- A. 1
- ☒ B. 2
- C. 3
- D. 4



# Review

Which of the following is correct about counting semaphore?

- A. The value of the semaphore is 0 or 1
- B. The same as binary semaphore
- C. The value of the semaphore variable can be above 1
- D. The value of the semaphore variable can never be below 0

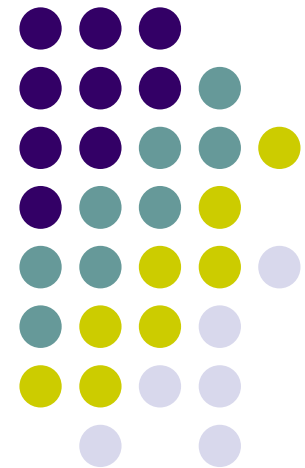
# Review

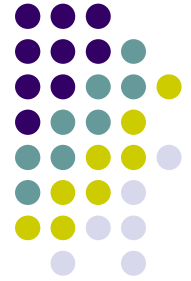


Which of the following is the most suitable use for counting semaphore?

- A. Use for shared resources with a single instance
- B. Use for shared resources with 2 instances
- C. Use for shared resources with any instances
- D.** Use for shared resources with multiple instances

# Deadlock






# Objectives

- Introduce what a deadlock is
- Introduce methods of handling deadlocks
- Implement deadlock handling algorithms

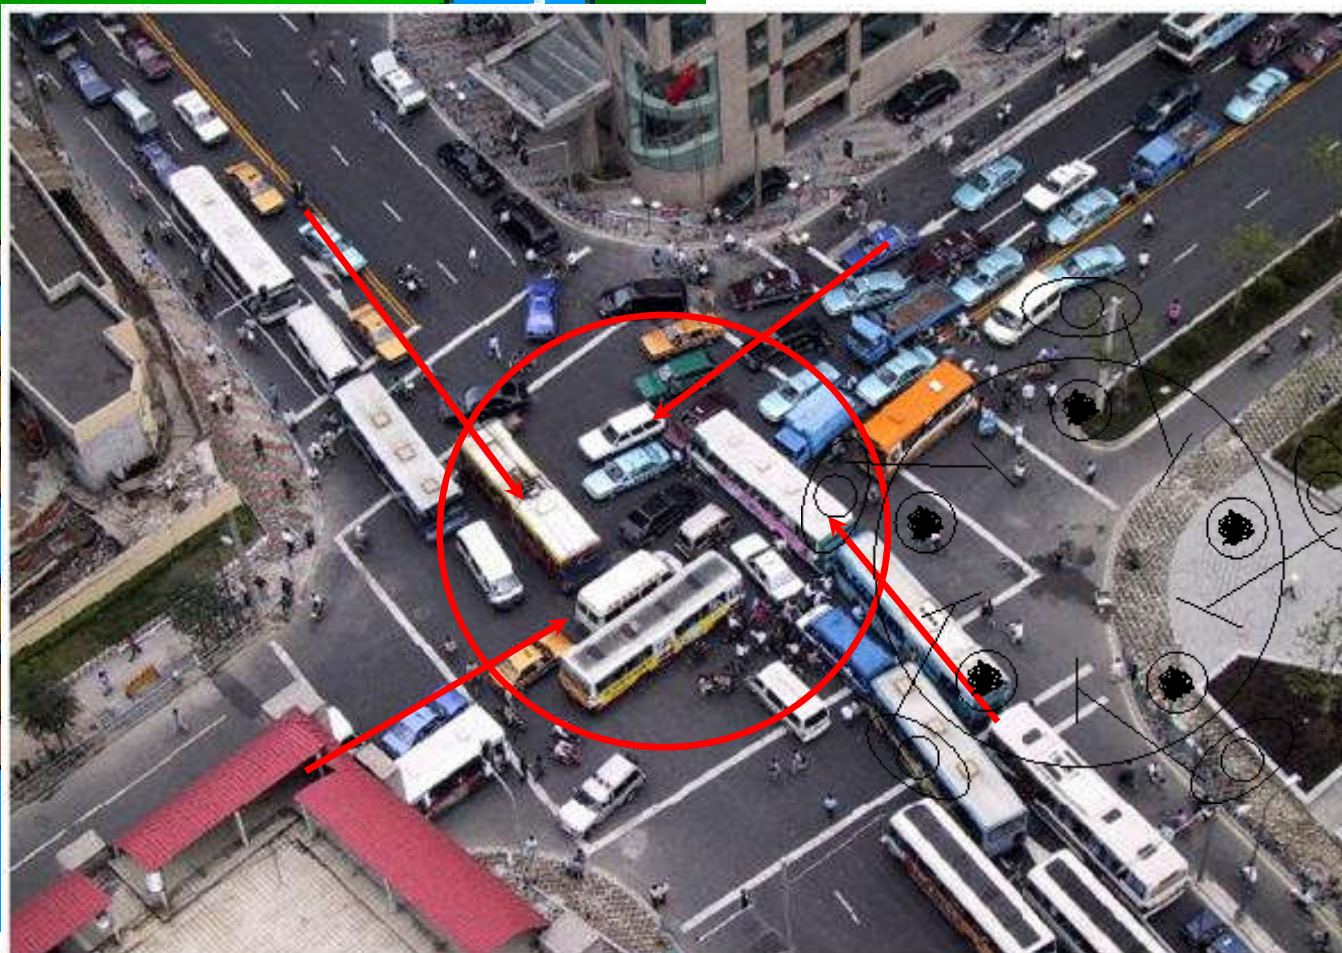
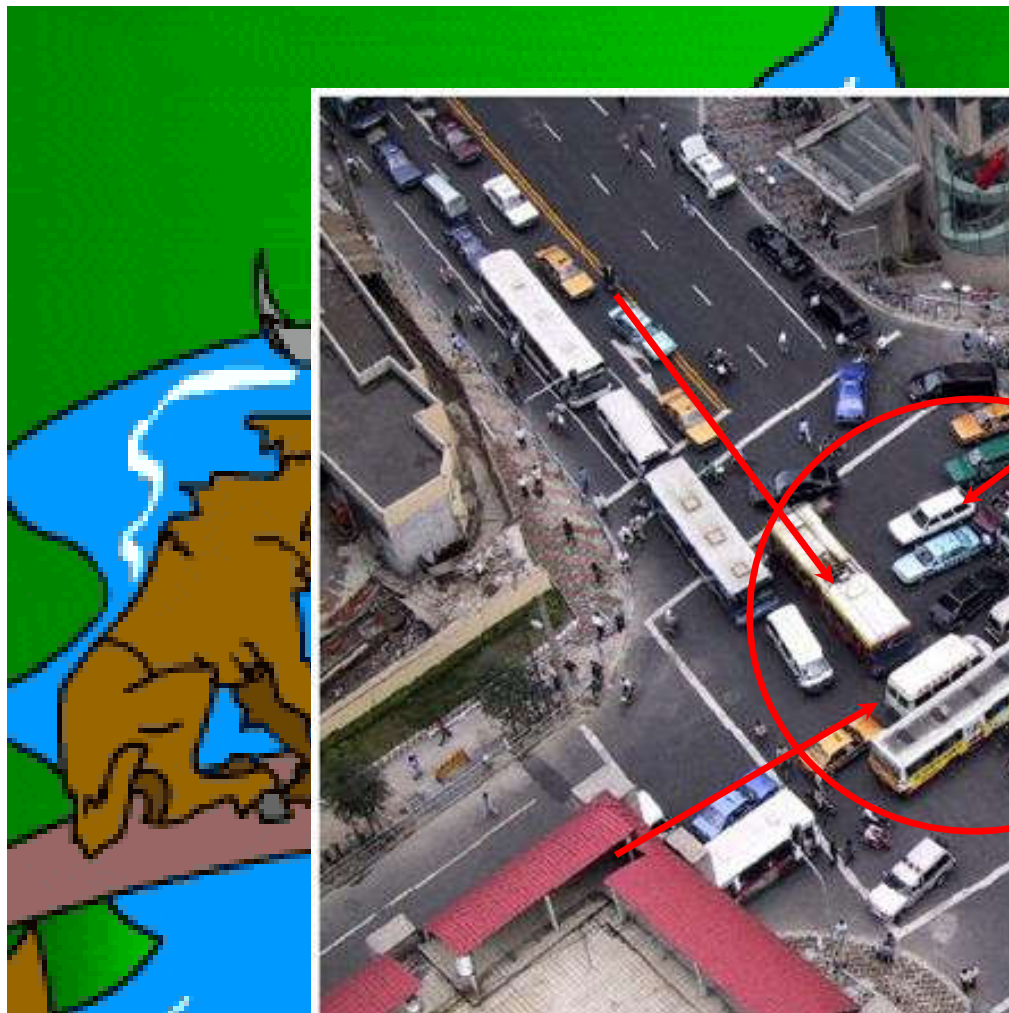
# Reference



- Chapter 7 of **Operating System Concepts**



# Deadlock examples



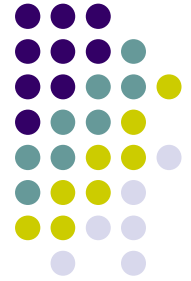
2/21/2017





# Definition of deadlock

- A **set** of blocked processes each
  - **holding** a resource and
  - **waiting** to acquire a resource held by another process in the set
- There must be a **circular wait** in this set



## Deadlock example (cont'd)

- Process A:

{

...

Lock file  $F_1$ ;

...

Open file  $F_2$ ;

...

Unlock  $F_1$ ;

}

- Process B

{

...

Lock file  $F_2$ ;

...

Open file  $F_1$ ;

...

Unlock  $F_1$ ;

}



## Question

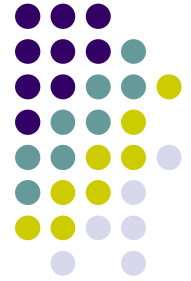
When does the deadlock happen?

- A. A gets F1 and waits for F2
- B. A gets F2 and waits for F1 and B waits for F1
- C. A gets F1 and waits for F2 and B gets F2 and waits for F1
- D. A gets F1 and F2 and B waits for F2

# Deadlock Characterization



- Deadlock can arise if four conditions hold simultaneously
  - **C1: Mutual exclusion**
  - **C2: Hold and wait** holding one resource, waiting other resources held by another
  - **C3: No preemption** only process has right to release its holding resources
  - **C4: Circular wait** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of processes:
    - $P_0$  is waiting for a resource that is held by  $P_1$ ,
    - $P_1$  is waiting for a resource that is held by  $P_2$ , ...
    - $P_n$  is waiting for a resource that is held by  $P_0$ .



# System Model

- Resource types  $R_1, R_2, \dots, R_m$ 
  - *shared variables, memory space, I/O devices,*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - request
  - use
  - release

# Resource-Allocation Graph



A set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the **processes** in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all **resource types** in the system.
- request edge – directed edge  $P_i \rightarrow R_j$
- assignment edge – directed edge  $R_j \rightarrow P_i$

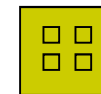
# Resource-Allocation Graph (Cont'd)




- Process



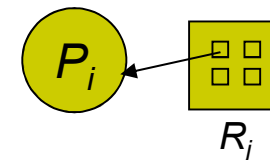
- Resource Type with 4 instances



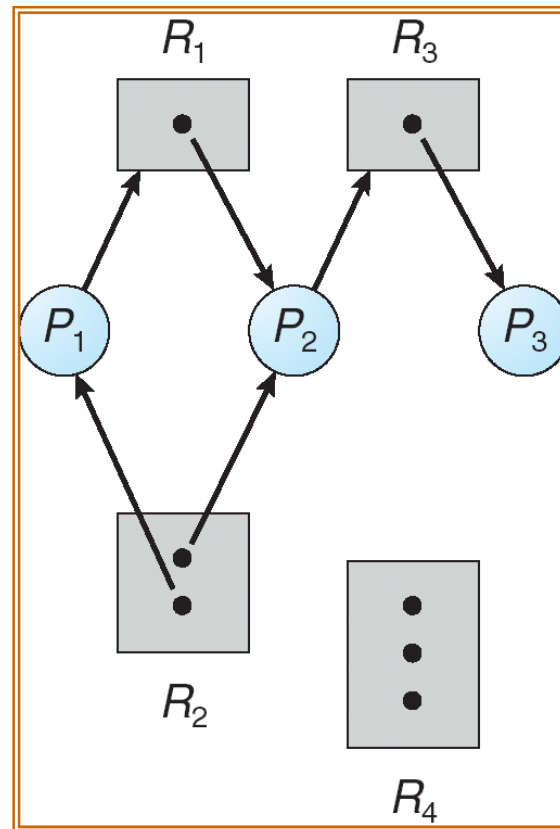
- $P_i$  requests instance of  $R_j$



- $P_i$  is holding an instance of  $R_j$

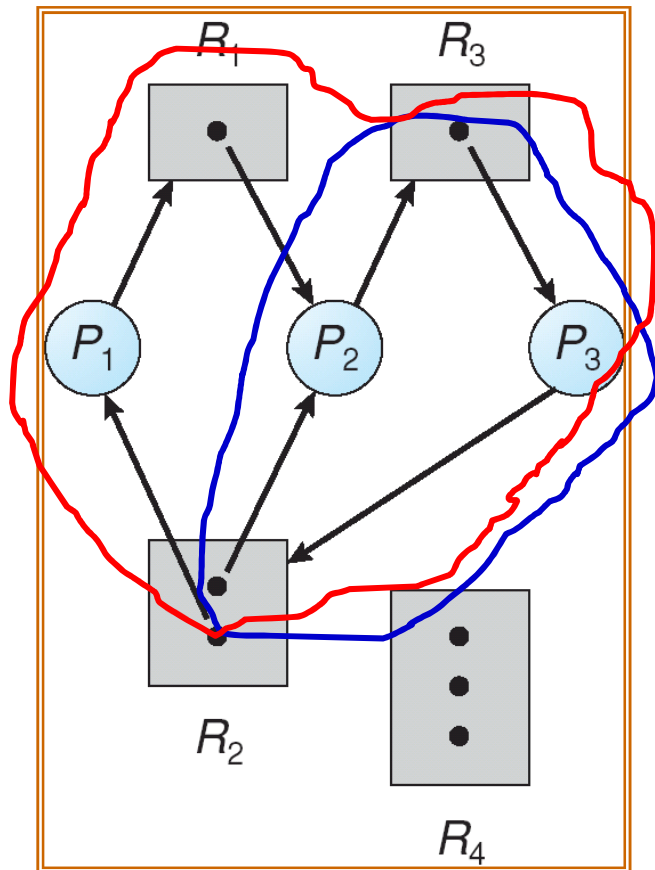


# Example of a RAG



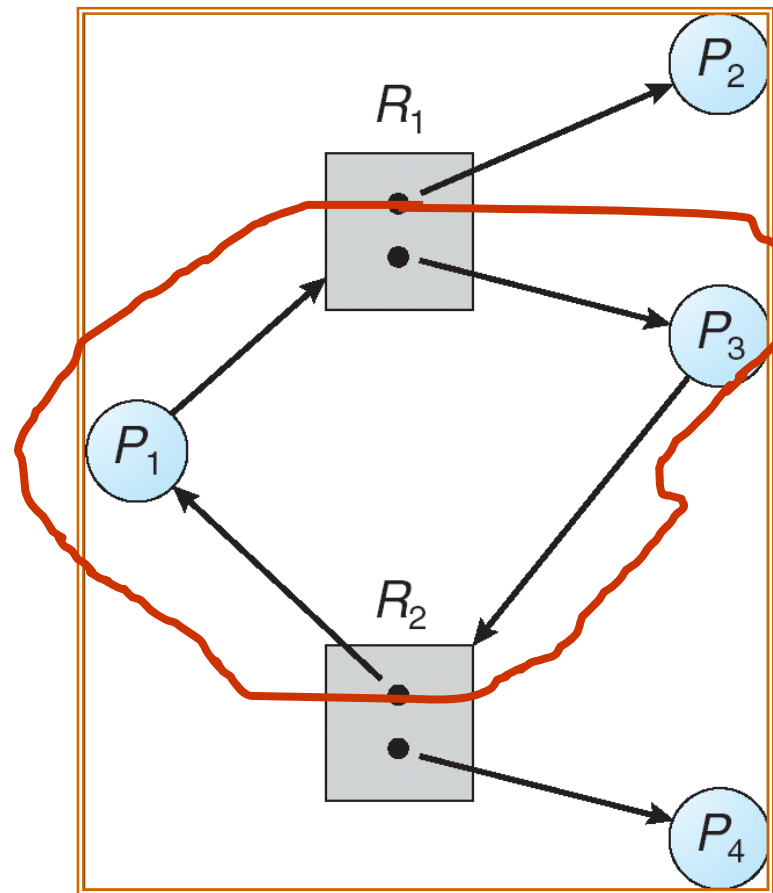


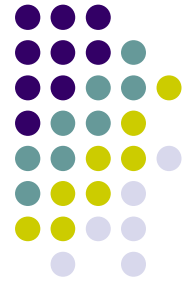
# RAG With A Deadlock



- When  $P_3$  asks for  $R_2$
- There are two circulars
  - $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_2 \rightarrow P_3 \rightarrow R_3 \rightarrow P_1$
  - $P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$
- Set of  $P_1, P_2, P_3$  is deadlock

# Graph With A Cycle But No Deadlock



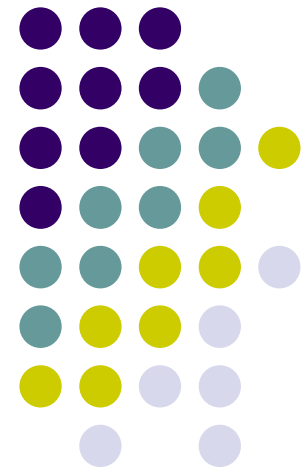


## Basic Facts

- If graph contains no cycles  $\Rightarrow$  no deadlock.  
If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock.
  - if several instances per resource type, possibility of deadlock.

# Deadlock handling

---

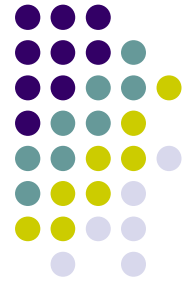


# Methods for Handling Deadlocks



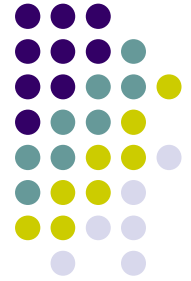
- Ensure that the system will **never** enter a deadlock state
  - Deadlock prevention, deadlock avoidance
- Allow the system to enter a deadlock state and then **recover**
  - Deadlock detection and recovery
- **Ignore** the problem and pretend that deadlocks never occur in the system
  - used by most operating systems, including UNIX.

# Deadlock Prevention



- The method prevents at least **one** of the **four deadlock conditions** from occurring
- This method is classified as a **static** method

# Deadlock Prevention



- **C1: Mutual Exclusion**
  - In some situations, this condition is required
  - Not feasible to make this NOT to happen

# Deadlock Prevention



- **C2: Hold and Wait**
  - Solution
    - must guarantee that whenever a process requests a resource, it does not hold any other resources, or
    - require process to request and be allocated all its resources before it begins execution
  - low resource utilization; starvation possible.



# Deadlock Prevention (Cont'd)



- **C3: No Preemption**

- If a process holding some resources requests another resource that cannot be immediately allocated to it,
  - then all resources currently being held are **released**
  - released resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources and the new requesting ones

# Deadlock Prevention (Cont'd)



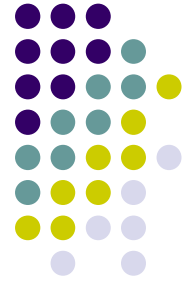
- **C4: Circular Wait**
  - impose a total ordering of all resource types and
  - require that each process requests resources in an increasing order of enumeration



## Question

How many conditions for a dead lock to happen are there?

- A. 2
- B. 3
- C. 4
- D. 5




## Question

When does a deadlock happen?

- A. any of the 4 conditions occur
- B. any two of the 4 conditions occur
- C. any 3 of the 4 conditions occur
- D. all the 4 conditions occur

# Deadlock avoidance

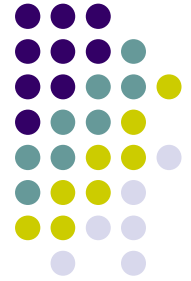
---





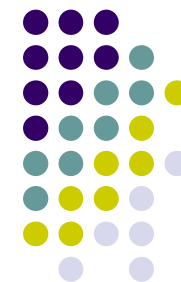
# Deadlock Avoidance

- This method requires **additional information** to decide resource allocation so that deadlock will not happen
  - each process has to register the **number of each required resource types** as additional information
- The deadlock-avoidance algorithm **dynamically** examines the resource-allocation state to ensure that there can never be a **circular-wait** condition



# Deadlock Avoidance

- Deadlock avoidance algorithms check the **state** of resource-allocation to decide allocation
- Resource-allocation **state** is defined by the number of **available** and allocated resources, and the **maximum demands** of the processes



## Safe State

- System is in **safe state** if a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes exists
  - $P_i$  can be satisfied by currently **available** resources + resources held by all the  $P_j$ , with  $j < i$
  - processes terminate in the above order

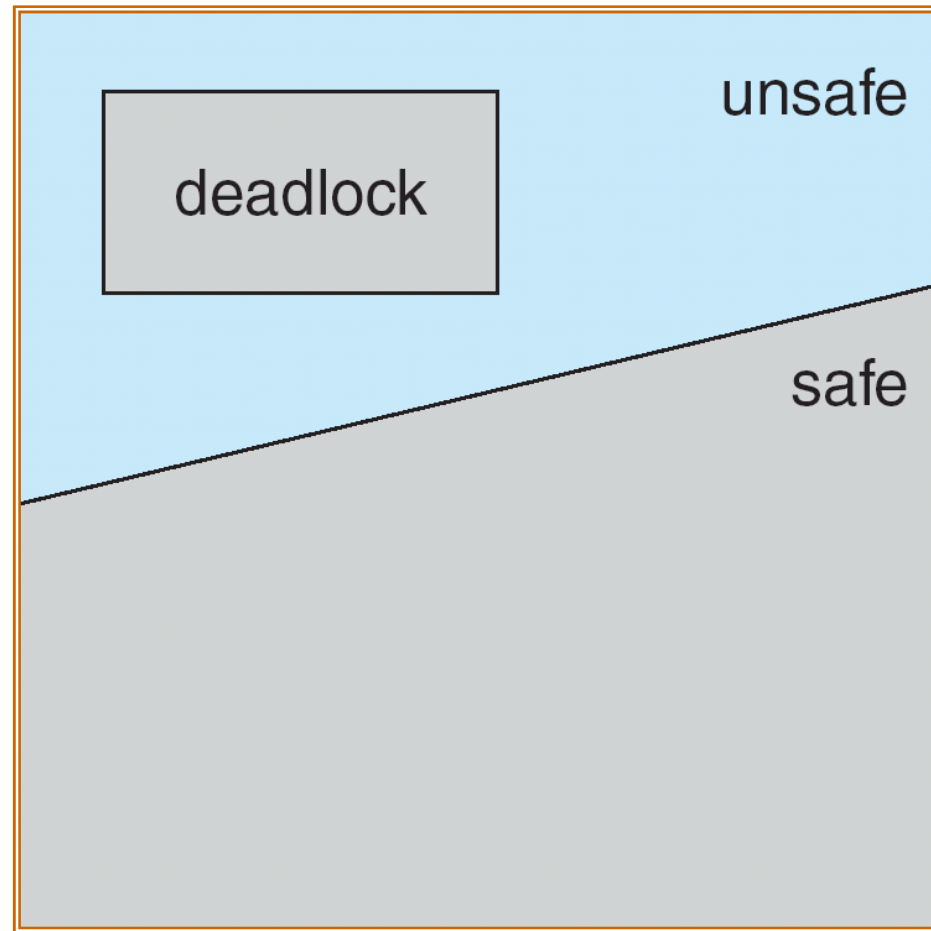




## Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will **never** enter an **unsafe state**

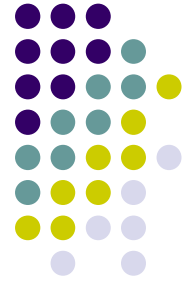
# Safe, Unsafe , Deadlock State





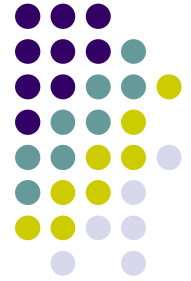
## Example

- A system has 12 tapes, and 3 processes  $P_0$ ,  $P_1$ ,  $P_2$  with corresponding requests:
  - $P_0$  requests at most 10 tapes
  - $P_1$  requests at most 4 tapes
  - $P_2$  requests at most 9 tapes
- At  $t_0$ ,  $P_0$  has 5 tapes,  $P_1$  and  $P_2$  each has 2 tapes
  - 3 tapes available



# Avoidance algorithms

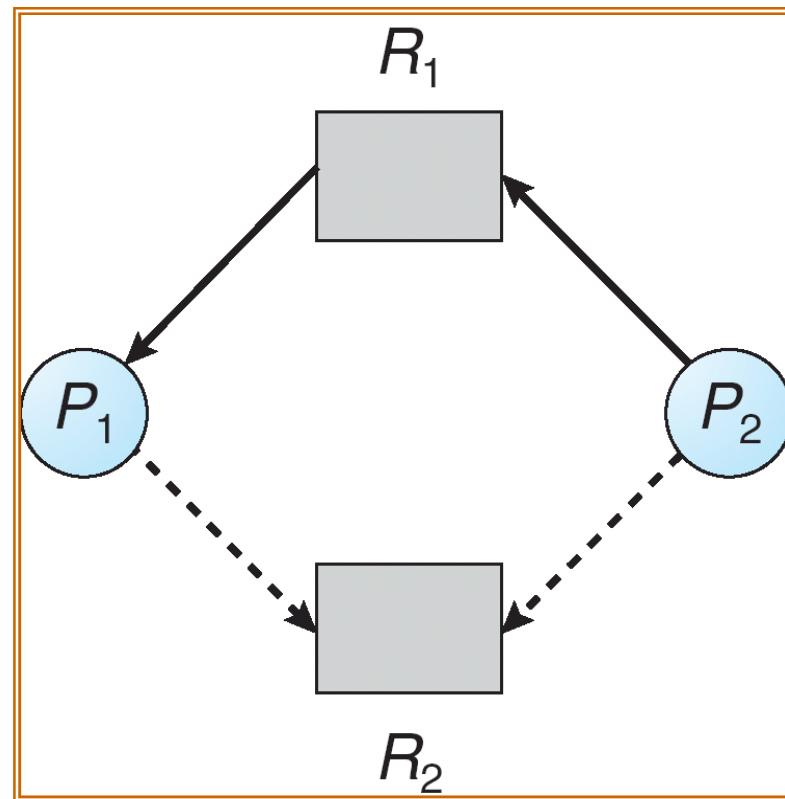
- Single instance of a resource type
  - Use a resource-allocation graph
- Multiple instances of a resource type
  - Use the banker's algorithm



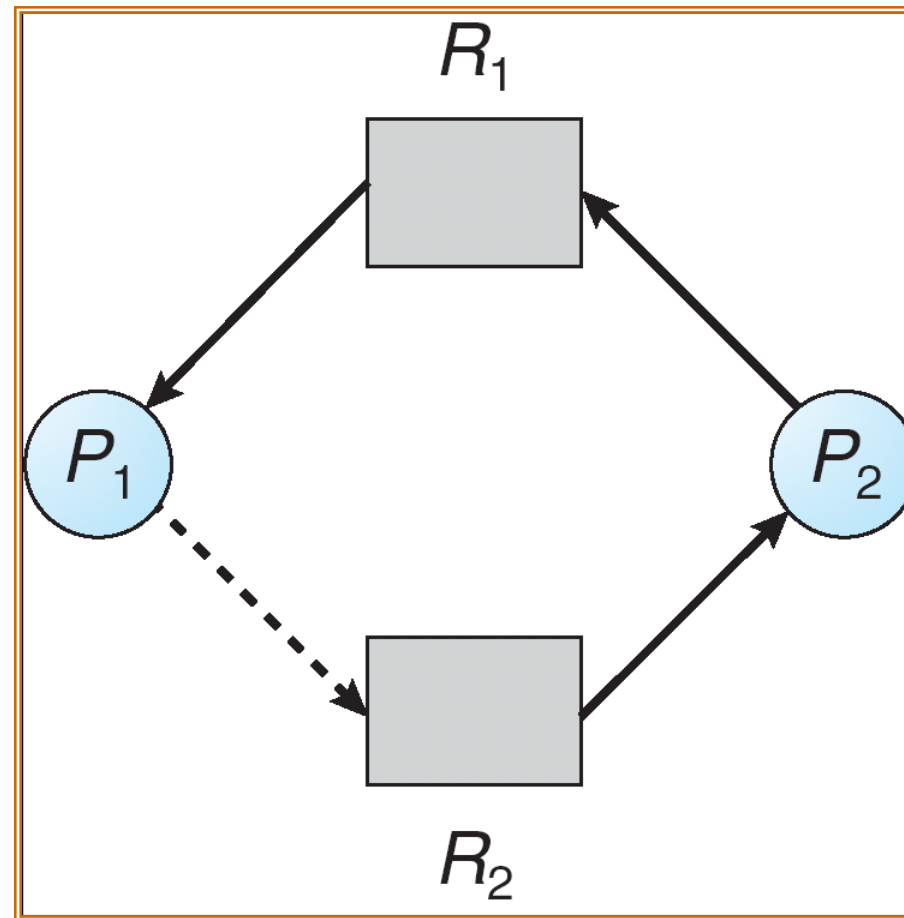
# RAG Algorithm convention

- **Claim** edge  $P_i \rightarrow R_j$ 
  - process  $P_j$  may request resource  $R_j$
  - presented as a dash line
  - Claim edge converts to **request** edge when a process requests a resource
- Request edge becomes an **assignment** edge when the resource is assigned to it
- When a resource is released by a process, assignment edge reconverts to a claim edge
  - Resources must be claimed *a priori* in the system.

# Resource-Allocation Graph



# Unsafe State In Resource-Allocation Graph



# Resource-Allocation Graph Algorithm



- Suppose that process  $P_i$  requests  $R_j$
- The request can be granted only if
  - converting the request edge to an assignment edge does not result in a **cycle** in the RAG





# Banker's Algorithm

- Multiple instances
  - Each process must a priori claim **maximum** use
- When a process requests a resource it may have to **wait**
- When a process gets all its resources, it must return them in a **finite** amount of time

# Data Structures for the Banker's Algorithm



Let  $n$  = number of processes, and  $m$  = number of resources types

- **Available:** Vector of length  $m$ 
  - Available  $[j] = k$ : there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix
  - Max  $[i,j] = k$ :  $P_i$  may request at most  $k$  instances of  $R_j$
- **Allocation:**  $n \times m$  matrix
  - Allocation  $[i,j] = k$ :  $P_i$  is allocated  $k$  instances of  $R_j$

# Data Structures for the Banker's Algorithm



Let  $n$  = number of processes, and  $m$  = number of resources types

- **Need**:  $n \times m$  matrix
  - $Need[i,j] = k$ :  $P_i$  may need  $k$  more instances of  $R_j$
  - $Need[i,j] = Max[i,j] - Allocation[i,j]$
- Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively
- Let  $A=(A_1, A_2, \dots, A_n)$ ,  $B=(B_1, B_2, \dots, B_n)$
- Define  $A \leq B$  if only if  $A_i \leq B_i, \forall 1 \leq i \leq n$

# Safety/Banker Algorithm



## 1. Initialize

$Work = Available$

$Finish[i] = false$  for  $i = 0, 1, \dots, n-1$

## 2. Find an $i$ that satisfies both

(a)  $Finish[i] = false$

(b)  $Need[i] \leq Work$

If no such  $i$  exists, go to step 4

## 3. $Work = Work + Allocation_i$

$Finish[i] = true$

go to step 2

## 4. If $Finish[i] == true$ for all $i$ , then the system is in a **safe** state

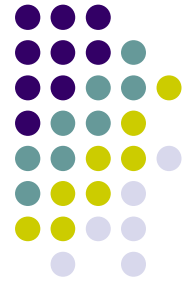
## Question



Which of the following is correct about the *Work* variable in the algorithm?

- A. It stores the available resources when each process finishes
- B. It is a redundant variable
- C. It stores the state of the system
- D. It stores possible resources for each process

# Question



Which of the following is the most correct about banker's algorithm?

- A. it detects the unsafe state of the system
- B. it detects the deadlock state of the system
- C. it detects the safe sequence of the system
- D. it detects the available resources

# Example of Banker's Algorithm



- 5 processes:  $P_0 - P_4$ ; 3 resource types
  - A (10 instances), B (5 instances), and C (7 instances)
- At time  $T_0$ :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	3	3	2
$P_1$	2	0	0	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			



## Example (Cont'd)

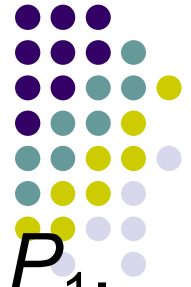
- Matrix *Need* = *Max* – *Allocation*

	<u>Need</u>		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

- The system is in a safe state or not?
  - sequence  $\langle P_1, P_3, P_0, P_2, P_4 \rangle$  satisfies safety criteria.



# Example



A system has 12 tapes, and 3 processes  $P_0$ ,  $P_1$ ,  $P_2$  with corresponding requests:

	<u>Max request</u>	<u>Current Allocation</u>
$P_0$	10	5
$P_1$	4	2
$P_2$	9	2

- At  $t_0$ , the system is in safe state
  - The sequence  $\langle P_1, P_0, P_2 \rangle$  is a safe sequence
- At  $t_1$ ,  $P_2$  requests 1 more tape (is it safe?)
  - the system is in unsafe state
  - it is **wrong** to allocate a tape for  $P_2$

# Resource-Request Algorithm



- Resource-request algorithm
  - another algorithm to avoid unsafe state
- Additional data structure
  - *Request* = request vector for process  $P_i$
  - $Request_i[j] = k$ : process  $P_i$  wants  $k$  instances of  $R_j$

# Resource-Request Algorithm



1. If  $Request_i \leq Need_i$  go to step 2 Otherwise, raise error condition

- since process has exceeded its maximum claim

2. If  $Request_i \leq Available$ , go to step 3 Otherwise  $P_i$  must wait

- since resources are not available

3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$ .
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Question



Which of the following is correct about resource-request algorithm?

- A. it detects the unsafe state of the system
- B. it detects the deadlock state of the system
- C. it detects the safe sequence of the system
- D. it detects the safe sequence of the system if the request is granted**



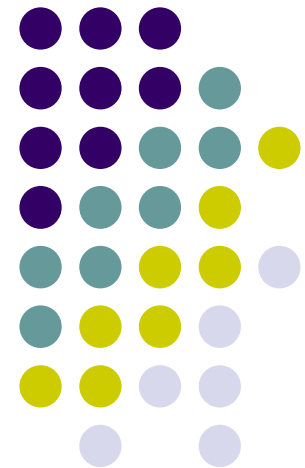
## Example: $P_1$ Request (1,0,2)

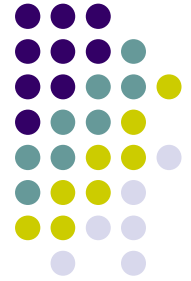
- Request  $\leq$  Available  $((1,0,2) \leq (3,3,2) \Rightarrow \text{true})$

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	1	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

- $\langle P_1, P_3, P_0, P_2, P_4 \rangle$  is a safe sequence
  - Can request for (1,0,0) by  $P_4$  be granted?
  - Can request for (0,2,0) by  $P_0$  be granted?

# Deadlock detection





# Deadlock Detection

- Allow system to enter deadlock state
- Use detection algorithms
- Recover from deadlock

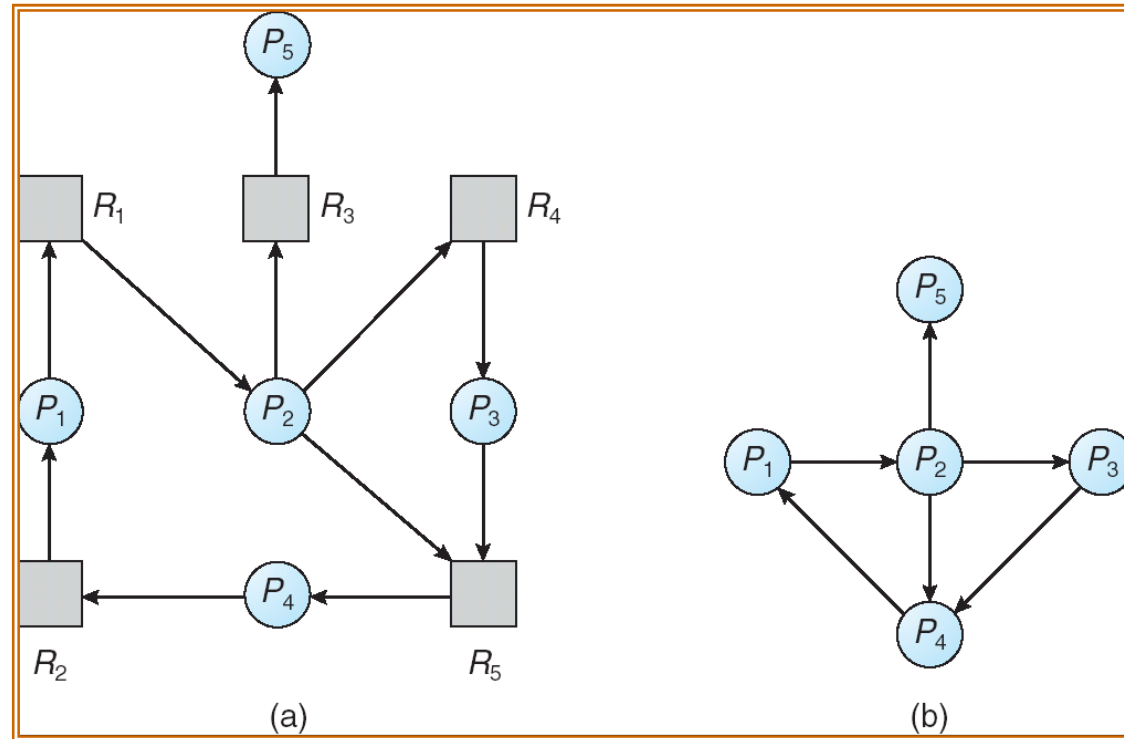
# Single Instance of Each Resource Type



- Maintain *wait-for* graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph
  - If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations
  - where  $n$  is the number of vertices in the graph



# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

# Several Instances of a Resource Type



- **Available:** A vector of length  $m$ 
  - number of available resources of each type
- **Allocation:** An  $n \times m$  matrix
  - number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix
  - current request of each process
  - If  $Request_i[j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$

# Detection Algorithm



Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively

1. Initialize:

(a) *Work* = *Available*

(b) For  $i = 1, 2, \dots, n$ ,

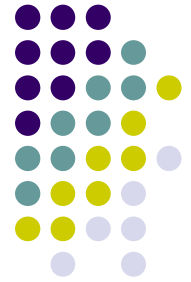
if  $Allocation_i \neq 0$  OR *Request* <sub>$i$</sub>   $\neq 0$ , then *Finish*[ $i$ ] = false;  
otherwise, *Finish*[ $i$ ] = true.

2. Find an index  $i$  such that both

(a) *Finish*[ $i$ ] == false

(b) *Request* <sub>$i$</sub>   $\leq$  *Work*

2/21/2017 If no such  $i$  exists, go to step 4.



## Detection Algorithm (Cont'd)

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2

4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state

- Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked.

Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state.

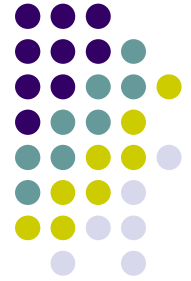
# Example of Detection Algorithm



- Processes  $P_0 - P_4$ ; resources (numbers)
  - A (7), B (2), and C (6)
- Snapshot at time  $T_0$  (deadlock?)

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Sequence  $\langle P_0, P_2, P_1, P_3, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .



## Example (Cont'd)

- $P_2$  requests an additional instance of type C

Request

A B C

$P_0$  0 0 0

$P_1$  2 0 1

$P_2$  0 0 1

$P_3$  1 0 0

$P_4$  0 0 2

- State of system (deadlock? processes in deadlock?)
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$

# Question



- Which of the following is correct about deadlock detection algorithm?
  - A. it only detects the unsafe state of the system
  - B. all the processes in the system are in the deadlock when it detects a deadlock
  - C. it can only detect the deadlock not the processes involved in the deadlock
  - D.** it can detect deadlock as well as the involved processes


# Detection-Algorithm Usage



- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - one for each disjoint cycle
- If detection algorithm is invoked arbitrarily
  - there may be many cycles in the resource graph
  - would not be able to tell which of the many deadlocked processes “caused” the deadlock.



# Recovery from Deadlock




# Recovery from Deadlock: Process Termination



- Abort all deadlocked processes
- Abort each process until the deadlock is removed
- In which order should we choose to abort?
  - Priority of the process.
  - How long process has computed, and/or how much longer to completion.
  - Resources the process has used.
  - Resources process needs to complete.
  - How many processes will need to be terminated.
  - Is process interactive or batch?

# Recovery from Deadlock: Resource Preemption

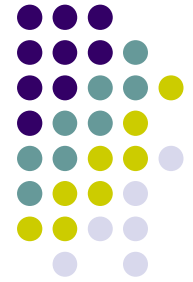


- Selecting a **victim**
  - minimize cost
- **Rollback**
  - return to some safe state, restart process for that state
- Starvation
  - same process may always be picked as victim, include number of rollback in cost factor

# Discussion



- For each abort condition, discuss which process will be selected to be cancelled
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?



# End of chapter





Windows is shutting down...

**Question?**