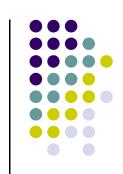
Operating Systems

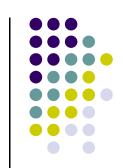
Nguyen Tri Thanh ntthanh@vnu.edu.vn





What is incorrect about overlays?

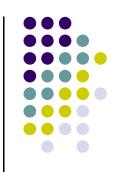
- A. overlays allows a large program to run in a smaller MEM
- B. Overlays only loads codes on demand (when they are used)
- Programmers need to split the program into modules
- Overlays is supported in all high level programming languages



What is incorrect about swapping?

- A. swapping is the same as overlays
- B. swapping uses hard disk as the backing store
- swapping allows many processes whose size is even larger than MEM to run
- a lower priority process is rolled out for a higher priority one to run (when needed)

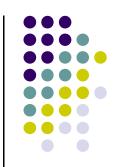
Review



Which is incorrect about non-contiguous MEM allocation?

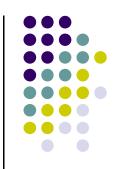
- A. split logical memory into parts
- utilize MEM more effectively in comparison with contiguous allocation method
- c. need a Memory Management Unit
- D. only suitable for some types of processes

Review



Which is correct about MMU of paging and segmentation allocation methods?

- A. they are the same
- B. MMU of paging needs more information than that of segmentation
- c. they use different resolution methods
- MMU of segmentation is faster than that of paging

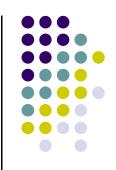


Suppose a process in contiguous allocation:

- the base address is 10400
- the limit register is 1200
- the reference is 246;

Which of the following is the correct physical address of the reference?

- A. 10154
- в. 10646
- c. 1446
- D. 954

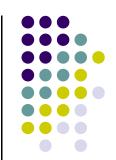


A system uses paging

- the frame size of 2KB;
- the address register is 32 bits

Which of the following is correct about register segmentation?

- A. (page:offset) = (19:13)
- B. (page:offset) = (21:11)
- c. (page:offset) = (22:10)
- D. (page:offset) = (20:12)



A system uses paging

- the frame size of 4KB;
- the address register is 32 bits
- Which of the following is the correct physical address of the reference (2,1296)?
 - A. 560*4096+1296
 - B. 120*4096+1296
 - c. 3*4096+1296
 - D. 120*1024+1296

Frame
56
120
3

Virtual Memory

Paging on demand
Page replacement
Frame allocation
Thrashing

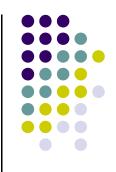


Objectives

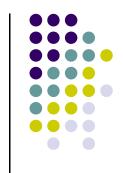


- Introduce paging method
- Introduce segmentation method





Chapter 9 of Operating System Concepts



Virtual memory

Virtual memory

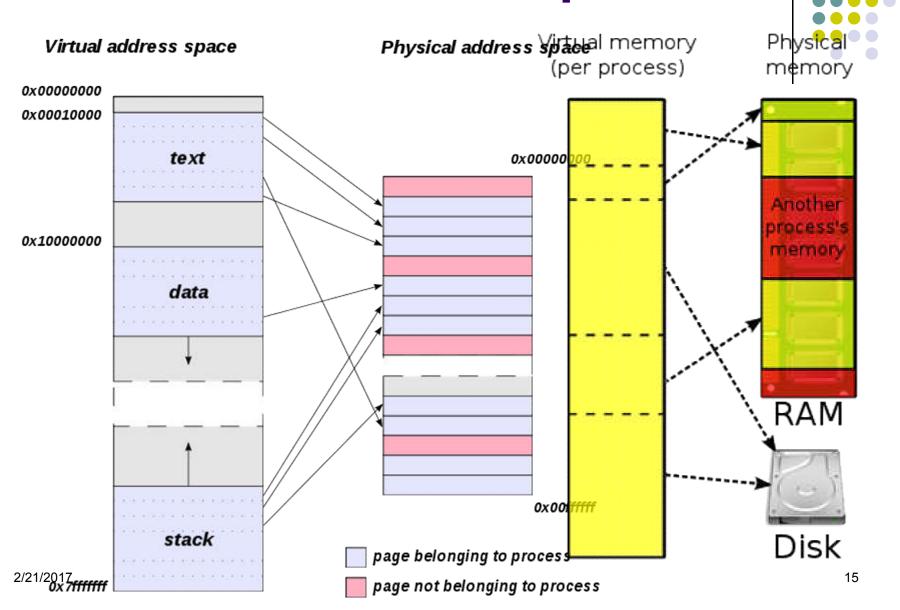
- Separation of user logical memory from physical memory.
 - Only a part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
- Virtual memory can be implemented via
 - Paging on demand
 - Segmentation on demand





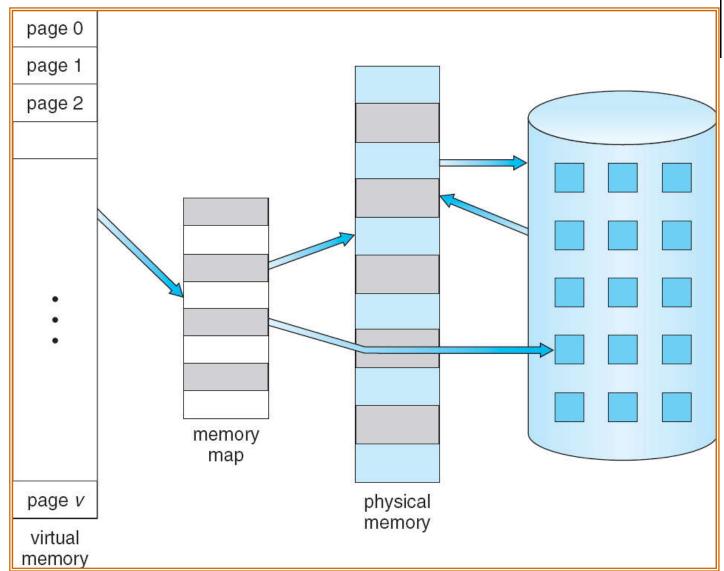
- Linux is, of course, a virtual memory system, meaning that the addresses seen by user programs do not directly correspond to the physical addresses used by the hardware. Virtual memory introduces a layer of indirection that allows a number of nice things. With virtual memory, programs running on the system can allocate far more memory than is physically available; indeed, even a single process can have a virtual address space larger than the system's physical memory. Virtual memory also allows the program to play a number of tricks with the process's address space, including mapping the program's memory to device memory.
- http://www.makelinux.net/ldd3/chp-15-sect-1

Virtual-address Space



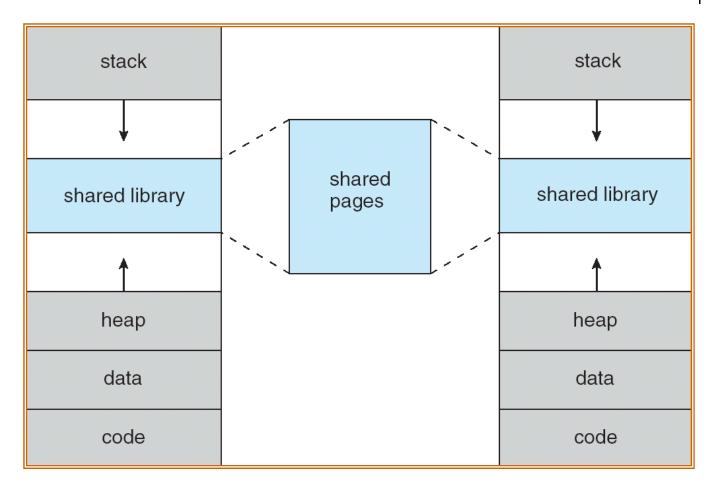
Virtual Memory That is Larger Than Physical Memory





Shared Library Using Virtual Memory







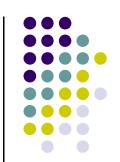


- Virtual memory allows other benefits
 - Copy-on-Write during process creation
 - Memory-Mapped Files

Copy-on-Write

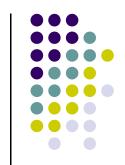
- Copy-on-Write (COW) allows both parent and child processes to initially share the same pages in memory
 - If either process modifies a shared page, only then is the page copied
 - refer to Sect. 2, Chapter 3 of "Lập trình C/C++ ..."
- COW allows more efficient process creation as only modified pages are copied

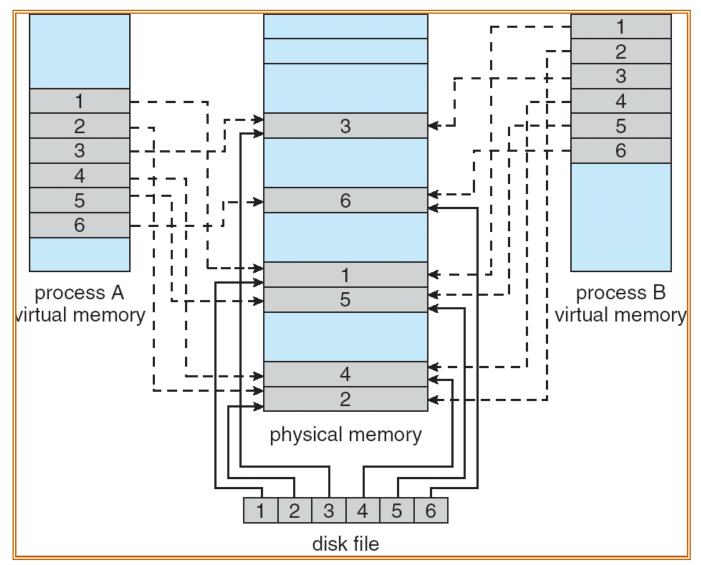
Memory Mapped Files



- A file is considered as a memory segment
- Read/write operations are performed via memory
 - not read/write file system calls
- Allow multiple processes shared a file
- refer to Sect. 5, Chapter 4 of "Lập trình /C++ trên Linux"

Memory Mapped Files





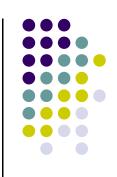


- What is the correct advantage of memory mapped file?
 - A. reduces the task of the system's OS
 - B. treats as the buffer for manipulating the file
 - c. allows programmers to organize the file
 - uses as shared resource among processes



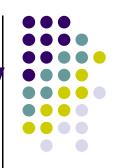
- Which of the following is incorrect about virtual memory?
 - A. it is separated from physical memory
 - B. it is mapped into physical memory during process execution
 - c. it gives additional benefits, e.g., COW, file mapping
 - an address in virtual memory is preserved when mapped into physical memory

Dynamic loading

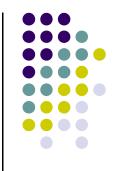


- Routine is not loaded until it is called
- Better memory-space utilization
 - unused routine is never loaded
- Useful when
 - large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system
 - refer to Section 3.4, Chapter 7 of "Lập trình C/C++ ..."

Dynamic linking and shared library



- Linking postponed until execution time
- Small piece of code, stub, used to locate the appropriate memory-resident library routine
 - Stub replaces itself with the address of the routine, and executes the routine
 - Operating system needed to check if routine is in processes' memory address
- Dynamic linking is particularly useful for libraries
- Also known as shared libraries in Linux
 - refer to Sect. 3, Chapter 7 of "Lập trình C/C++ trên Linux"



Paging on demand



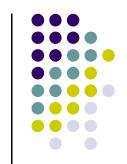
Paging on demand

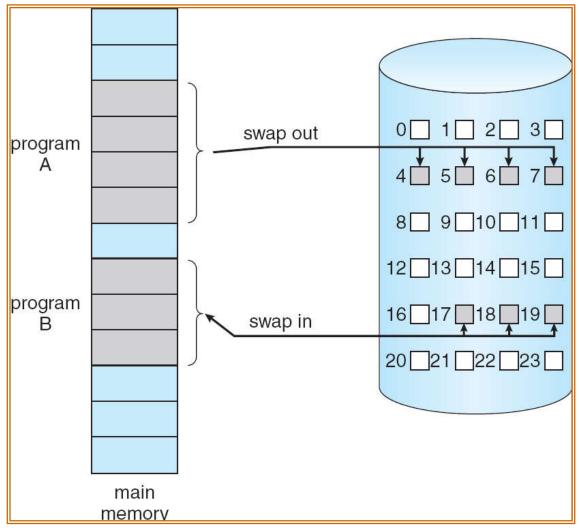
- Bring a page into memory only when needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed ⇒ reference to it
 - invalid reference ⇒ abort
 - not-in-memory ⇒ bring to memory
- Lazy swapper never swaps a page into memory unless page is needed
 - Swapper that deals with pages is a pager



- Why it is possible when only a part of a process is loaded into MEM?
 - A. Because instructions of a process are independent
 - B. Because we can indicate which instructions to run
 - Because only one instruction is executed at a time
 - Because related instructions are always in the same group

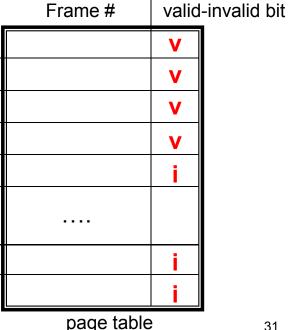
Transfer of a Paged Memory to Contiguous Disk Space



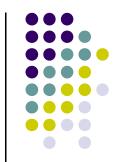


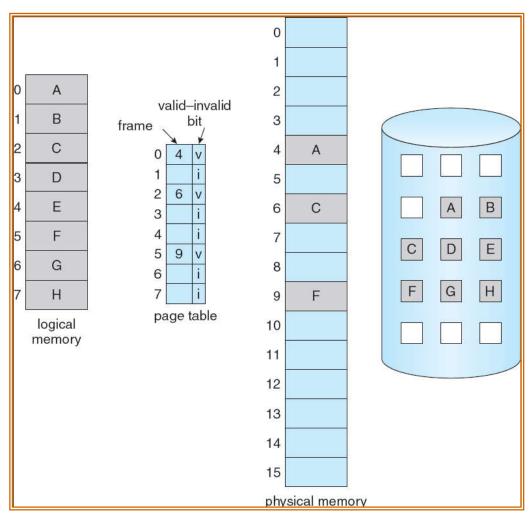
Valid-Invalid Bit

- With each page table entry a valid—invalid bit is associated
 - ($\lor \Rightarrow$ in-memory, $i \Rightarrow$ not-in-memory)
- Initially valid—invalid bit is set to i on all entries
- Example of a page table snapshot:
- During address translation, if valid-invalid bit in page table entry is $I \Rightarrow$ page fault



Page Table When Some Pages Are Not in Main Memory





2/21/2017 32

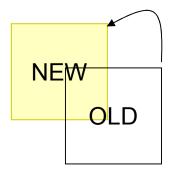
Page Fault

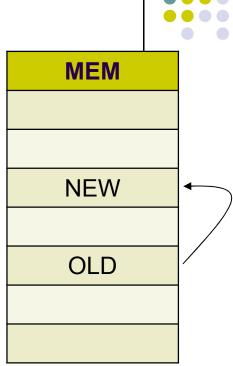
- If there is a reference to a page,
 - first reference to that page will trap to operating system: page fault
 - 1. Operating system looks at another table to decide:
 - Invalid reference ⇒ abort
 - Just not in memory
 - 2. Get empty frame
 - 3. Swap page into frame
 - 4. Update the page table
 - 5. Set validation bit = v
 - 6. Restart the instruction that caused the page fault



Page Fault (Cont.)

- Restart instruction
 - block move

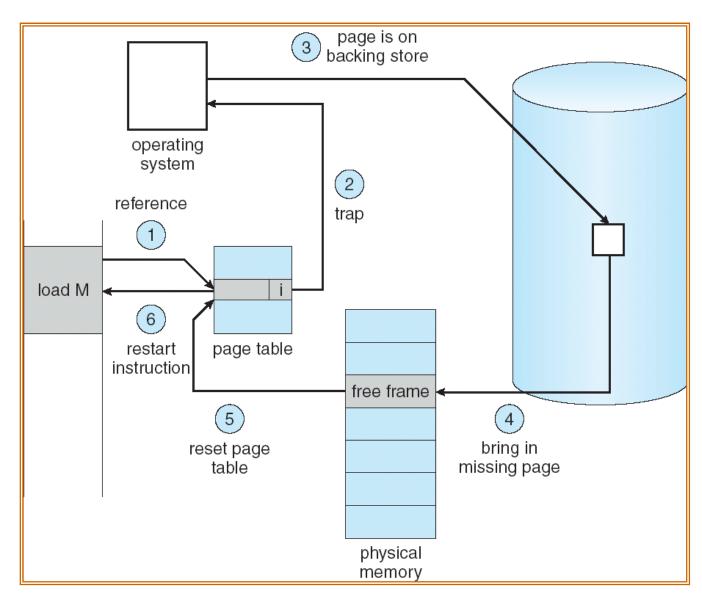


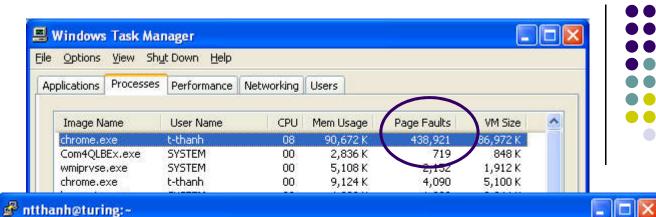


auto increment/decrement location

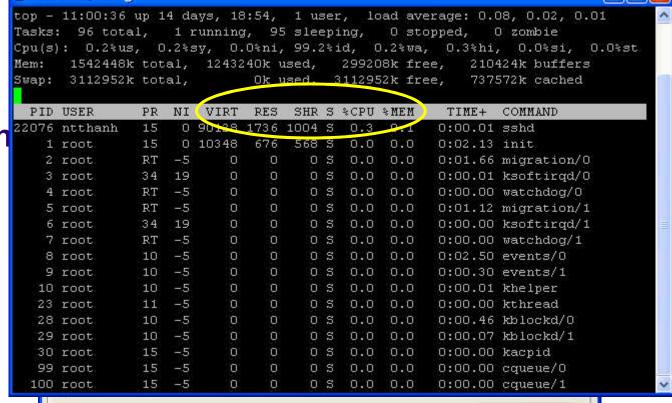
Steps in Handling a Page Fault



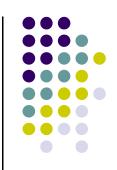




Process information



Question



- Which of the following is incorrect about a page fault?
 - A. it happens in paging on demand
 - B. it happens when a reference to a page that is not in MEM
 - when a page fault occurs the corresponding process will be terminated
 - a page fault handler is called whenever it occurs

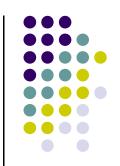
Question

- Which of the following is incorrect order of steps in handling a page fault?
 - A. check if the valid bit is invalid ⇒ raise a page fault
 - B. a page fault is raised ⇒ find the page in backing store
 - c. a page fault is raised \Rightarrow find a free frame
 - D. load the page into memory ⇒ restart the instruction

2/21/2017

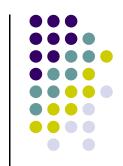
38

If no free frame available



- Call page replacement procedure
 - swap out an unused page from MEM
- Algorithms
 - FIFO, Optimal, LRU, LRU-approximation
- Performance of the algorithm
 - page-fault rate
 - which algorithm is better?

Performance of paging on demand



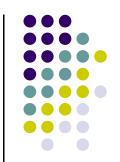
- Page Fault Rate $0 \le p \le 1.0$
 - if p = 0 no page faults
 - if p = 1, every reference is a page fault
- Effective Access Time (EAT)

```
EAT = (1 - p) * memory access
```

+ p (page fault overhead

- + swap page out
- + swap page in
- + restart overhead)

Paging on Demand Example

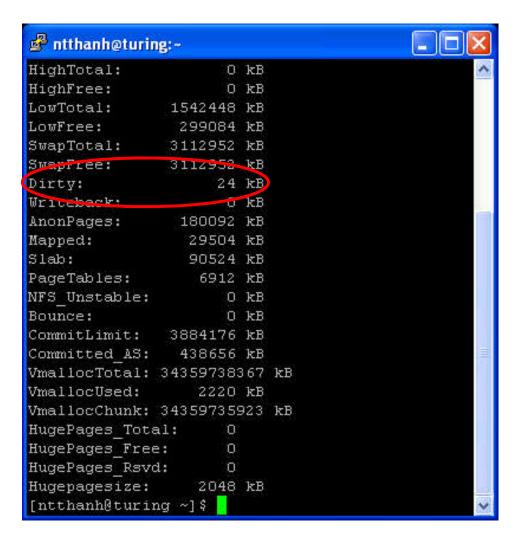


- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
 - EAT = $(1 p) \times 200 + p (8 \text{ milliseconds})$ = $(1 - p) \times 200 + p \times 8,000,000$ = $200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault
 - EAT = 8.2 microseconds.
 - slowdown by a factor of 40!!

Page Replacement

- Prevent over-allocation of memory
 - include page replacement in page-fault service routine
- Use modify (dirty) bit to reduce overhead of page transfers
 - only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory
 - large virtual memory can be provided on a smaller physical memory

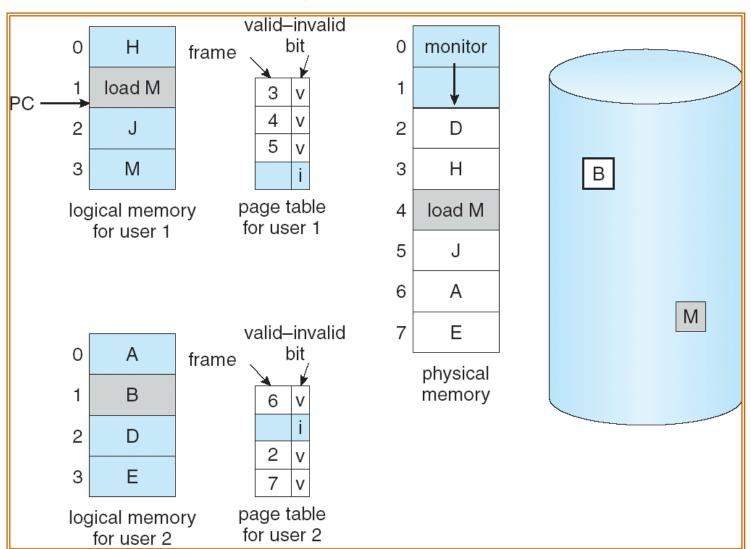
Page Replacement





2/21/2017 43

Need For Page Replacement





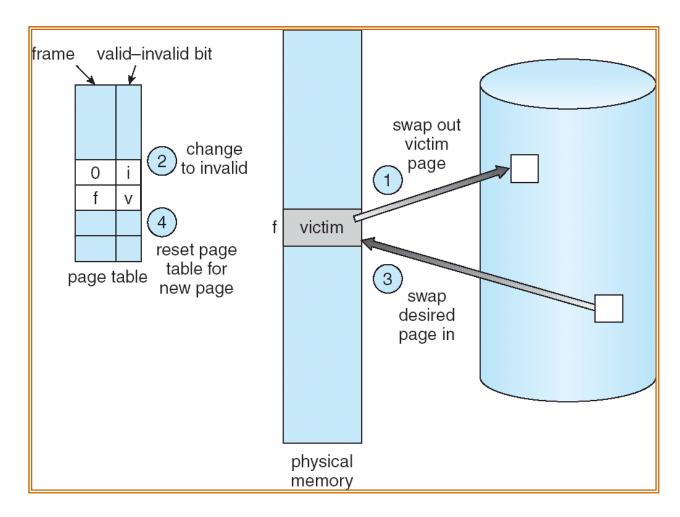
2/21/2017 44

Basic Page Replacement

- 1. Find the location of the desired page on disk
- 2. Find a free frame
 - If there is a free frame, use it
 - Else use a page replacement algorithm to select a victim frame
- 3. Bring the desired page into the (newly) free frame; update the page table
- 4. Resume the process

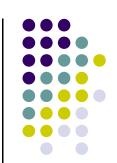
Page Replacement





2/21/2017 46

Question



Which of the following is incorrect about page replacement?

- A. a victim frame is selected to be swapped out
- B. the page table which is the victim will be updated
- the victim frame is always written into the backing store
- the victim frame is only written into the backing store if it is dirty

Page Replacement Algorithms



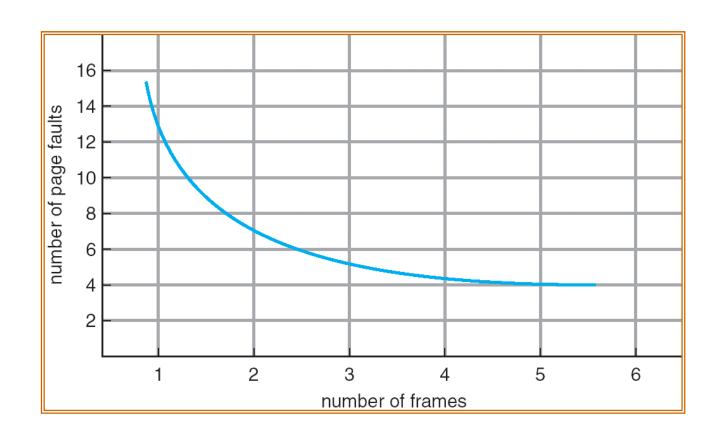
- Want lowest page-fault rate
- Evaluate algorithm
 - run it on a particular string of memory references (reference string)
 - compute the number of page faults on that string
- In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

2/21/2017 48

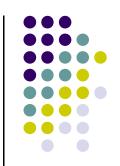
Graph of Page Faults Versus The Number of Frames





2/21/2017 49

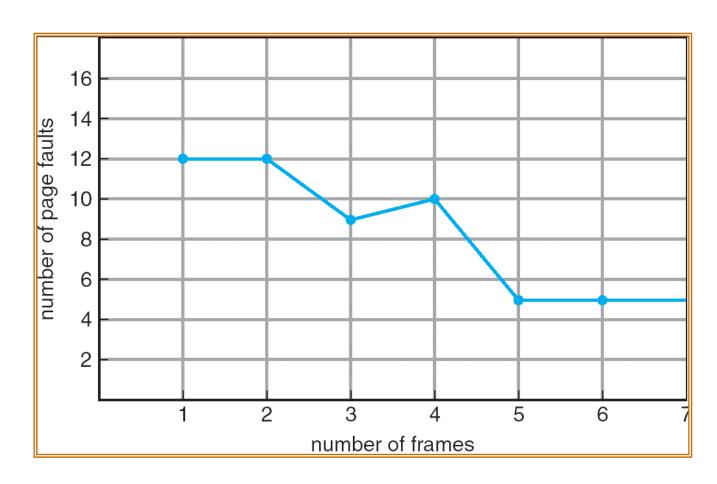
First-In-First-Out (FIFO) Algorithm



Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

FIFO Illustrating Belady's Anomaly

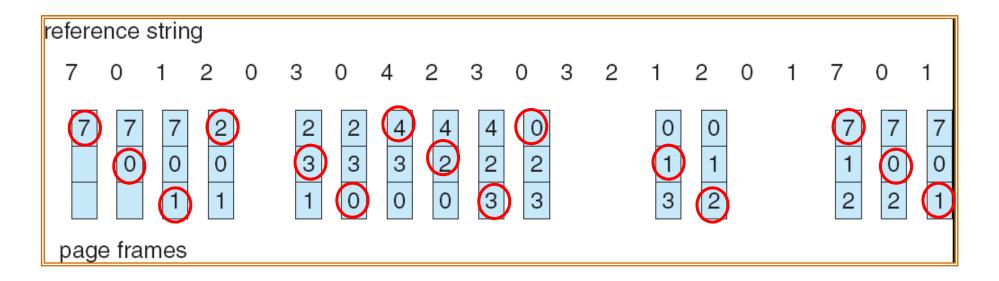




2/21/2017 51

FIFO Page Replacement









 A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, FIFO is used with 3 frames. Which of the following is the correct order of swapped out pages?

- A 701230423012
- B. 701230432012
- c. 701320423012
- D. 701230423102

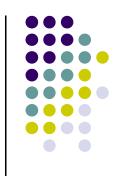




 A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, FIFO is used with 3 frames. Which of the following is the correct number of page faults?

- A. 13
- в. 14
- c. 15
- D. 16

Optimal Algorithm

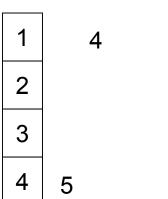


- Replace page that will not be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

How do you know this?

 Used for measuring how well the algorithm performs



6 page faults





 A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, Optimal algorithm is used with 3 frames. Which of the following is the correct order of swapped out pages?

- A. 710342
- B. 701432
- c. 710432
- D. 714132

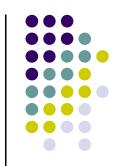




 A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, Optimal algorithm is used with 3 frames. Which of the following is the correct number of page faults?

- A. 8
- B. 9
- c. 10
- D. 11

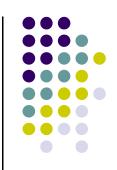
Least Recently Used (LRU) Algorithm



- Least recently used page is swapped out first
 - Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - 4 frames

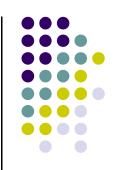
1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3





- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, LRU is used with 3 frames.
- Which of the following is the correct order of swapped out pages?
 - A. 712314132
 - B. 721304232
 - c. 712304123
 - D. 712304032





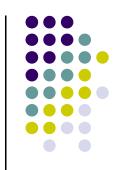
- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, LRU is used with 3 frames.
- Which of the following is the correct number of page faults?
 - A. 13
 - В. 12
 - c. 11
 - D. 10

LRU Algorithm (Cont.)



- Stack implementation
 - keep a stack of page numbers in a double link form
 - Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
 - No search for replacement

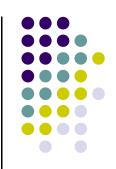
Use Of A Stack to Record The Most Recent Page References



reference string												
4	7	0	7	1	0	1	2	1	2	7	1	2
	2					-	7		Î	Î		
	1					2	2		а	b		
	0					-	1					
	7					(0					
	4					2	4					
	stack before			stack after								
a)								

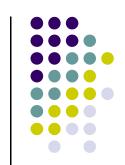
2/21/2017 62

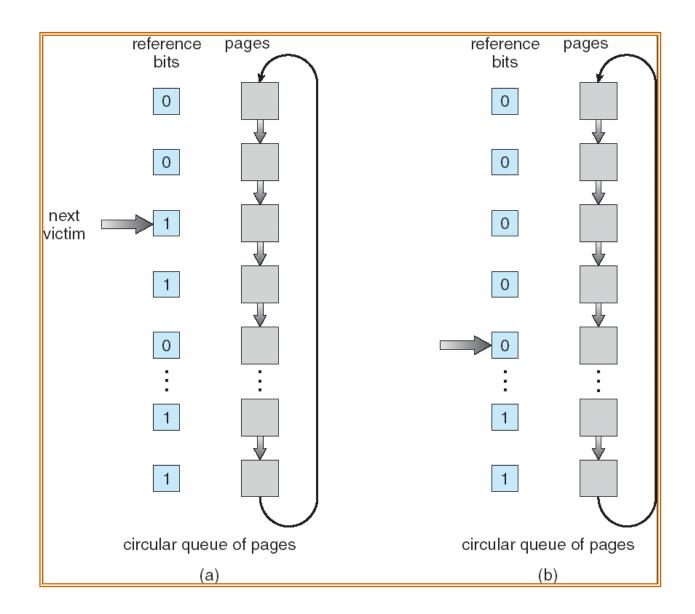
LRU Approximation Algorithms



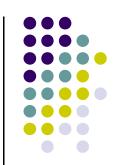
- Reference bit
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace the one which is 0 (if one exists)
 - We do not know the order
- Second chance (follow clock order)
 - Need reference bit
 - Clock replacement
 - If page to be replaced has reference bit = 1 then:
 - set reference bit 0
 - leave page in memory
 - replace next page, subject to same rules

Second-Chance (clock) Page-Replacement Algorithm



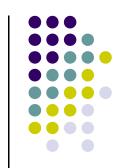


Question



- Suppose the second chance is used;
 - the reference bits of frames are: 1 1 0 1 1 0
 - the head is at second frame
- Which of the following are the reference bits after a page replacement is done
 - A. 000110
 - B. 101110
 - c. 100110
 - D. 101010

Least frequently Used (LFU) Algorithm



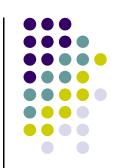
- Counter implementation
 - Every page entry has a counter;
 - every time page is referenced, copy the clock into the counter
 - When a page needs to be swapped
 - look at the counters to determine

Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- Least Frequently Used (LFU) Algorithm
 - replaces page with smallest count
- Most Frequently Used (MFU) Algorithm
 - based on the argument that the page with the smallest count was probably just brought in and has yet to be used



Allocation of Frames



- Each process needs minimum number of pages
- Example: IBM 370 6 pages to handle SS MOVE instruction:
 - instruction is 6 bytes, might span 2 pages
 - 2 pages to handle from
 - 2 pages to handle to
- Two major allocation schemes
 - fixed allocation
 - priority allocation

Fixed Allocation



- Equal allocation
 - For example, if there are 100 frames and 5 processes, give each process 20 frames.
- Proportional allocation
 - Allocate according to the size of process

$$-s_i = \text{size of process } p_i$$

$$-S = \sum S_i$$

-m = total number of frames

$$-a_i = \text{allocation for } a_i = \frac{S_i}{S} \times m$$

$$m = 64$$
 $s_1 = 10$
 $s_2 = 127$
 $a_1 = \frac{10}{137} \times 64 \approx 5$
 $a_2 = \frac{127}{137} \times 64 \approx 59$

Question



- A system uses proportional allocation and has
 - 90 frames x 2KB
 - 3 processes with size of (138KB, 96KB, 164KB)
- Which of the following is the correct number of allocated frames of (P₁, P₂, P₃)
 - A. 32, 21, 37
 - в. 31, 22, 37
 - c. 30, 22, 38
 - D. 33, 22, 35

Priority Allocation



71

- Use a proportional allocation scheme using priorities rather than size
- If process P_i generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

Global vs. Local Allocation



Global replacement

- process selects a replacement frame from the set of all frames;
- one process can take a frame from another

Local replacement

 each process selects from only its own set of allocated frames

Thrashing

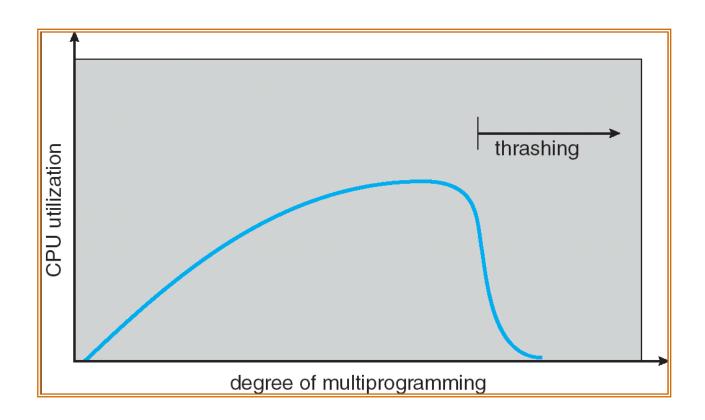
 If a process does not have "enough" frames, the page-fault rate is very high.

Thrashing

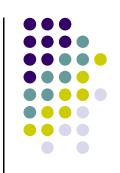
- a process is busy swapping pages in and out
- This caused by:
 - low CPU utilization
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system

Thrashing (Cont.)



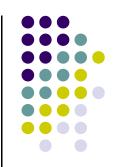


Solutions to Thrashing



- Use local allocation
- Use priority allocation
 - not good solution
- Working set model
 - A suitable solution

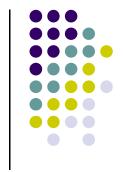
Question



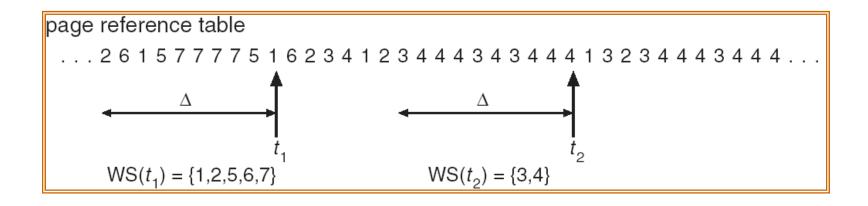
- Which of the following is incorrect about priority allocation?
 - A. higher priority process is allocated first
 - B. it prevents thrashing from happening
 - c. frames are allocated globally
 - it does not prevent thrashing from happening

Working-Set Model

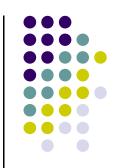
- $\Delta \equiv$ working-set window
 - a number of page references, e.g. 10,000
- Working set of Process P_i
 - WSS_i =total number of pages referenced in the most recent Δ (varies in time)
 - if ∆ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \Sigma WSS_i \equiv \text{total demand frames}$
- if D > m (total of frames)⇒ Thrashing
- Policy if D > m, then suspend one process



Working-set model

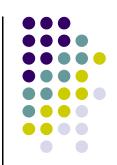


Keeping Track of the Working Set



- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page
 - Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - If one of the bits in memory = $1 \Rightarrow$ page in working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units

Question



- Suppose a delta =10; reference string
 - 261577775162344434413234443
 444...
- Which of the following is the correct WSS at 20th reference?
 - A. {2 3 4 6}
 - B. {2 3 4 5 6}
 - c. {1 2 3 4 6}
 - D. {7 1 2 3 4 6}



Question?