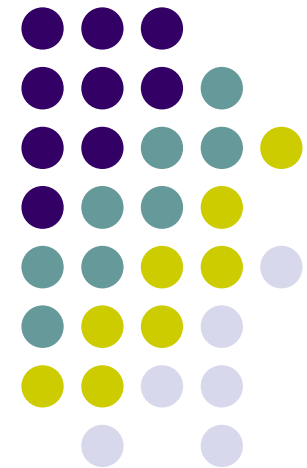


Operating Systems

Nguyen Tri Thanh
ntthanh@vnu.edu.vn



Question



What is **incorrect** about overlays?

- A. overlays allows a large program to run in a smaller MEM
- B. Overlays only loads codes on demand (when they are used)
- C. Programmers need to split the program into modules
- D. Overlays is supported in all high level programming languages

Question



What is **incorrect** about swapping?

- A. swapping is the same as overlays
- B. swapping uses hard disk as the *backing store*
- C. swapping allows many processes whose size is even larger than MEM to run
- D. a lower priority process is rolled out for a higher priority one to run (when needed)

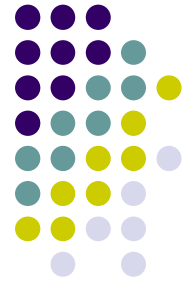
Review



Which is incorrect about non-contiguous MEM allocation?

- A. split logical memory into parts
- B. utilize MEM more effectively in comparison with contiguous allocation method
- C. need a Memory Management Unit
- D. only suitable for some types of processes

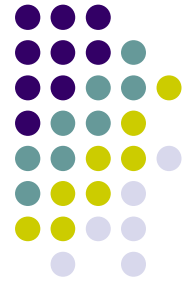
Review



Which is correct about MMU of paging and segmentation allocation methods?

- A. they are the same
- B. MMU of paging needs more information than that of segmentation
- C. they use different resolution methods
- D. MMU of segmentation is faster than that of paging

Question



Suppose a process in contiguous allocation:

- the base address is 10400
- the limit register is 1200
- the reference is 246;

Which of the following is the correct physical address of the reference ?

- A. 10154
- B. 10646
- C. 1446
- D. 954

Question




A system uses paging

- the frame size of 2KB;
- the address register is 32 bits

Which of the following is correct about register segmentation?

- A. (page:offset) = (19:13)
- B. (page:offset) = (21:11)
- C. (page:offset) = (22:10)
- D. (page:offset) = (20:12)

Question




A system uses paging

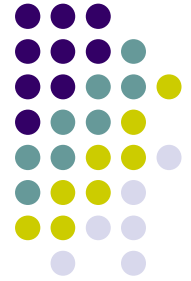
- the frame size of 4KB;
- the address register is 32 bits
- Which of the following is the correct physical address of the reference (2,1296)?
 - A. $560 \times 4096 + 1296$
 - B. $120 \times 4096 + 1296$
 - C. $3 \times 4096 + 1296$
 - D. $120 \times 1024 + 1296$

Frame
56
120
3

Virtual Memory

Paging on demand
Page replacement
Frame allocation
Thrashing






Objectives

- Introduce paging method
- Introduce segmentation method

Reference



- Chapter 9 of **Operating System Concepts**



Virtual memory

Virtual memory



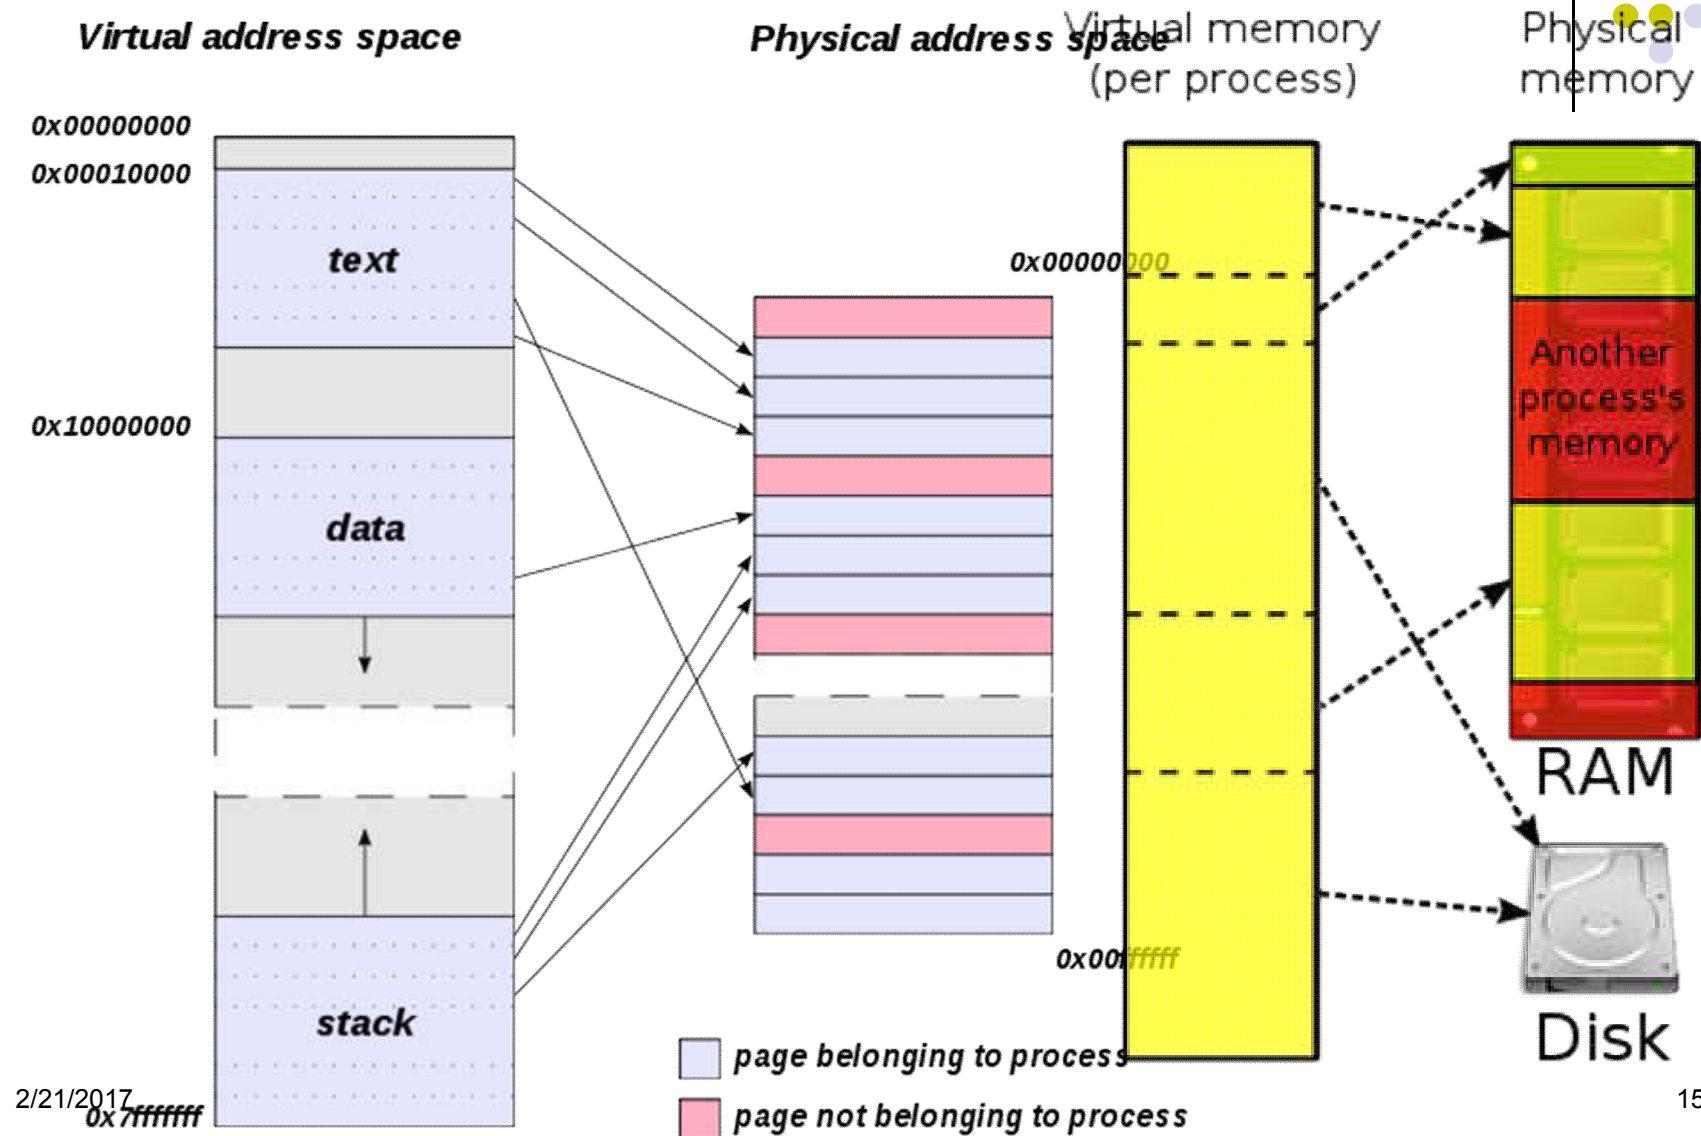
- Separation of user logical memory from physical memory.
 - Only a **part** of the program needs to be in memory for execution
 - Logical address space can therefore be much **larger** than physical address space
 - Allows address spaces to be **shared** by several processes
 - Allows for more efficient process creation
- Virtual memory can be **implemented** via
 - Paging on demand
 - Segmentation on demand




Virtual memory

- Linux is, of course, a **virtual memory system**, meaning that the addresses seen by user programs do not directly correspond to the physical addresses used by the hardware. Virtual memory introduces a layer of indirection that allows a number of nice things. With virtual memory, programs running on the system can allocate **far more memory than is physically available**; indeed, even a single process can have a virtual address space larger than the system's physical memory. Virtual memory also allows the program to play a number of tricks with the process's address space, including mapping the program's memory to device memory.
- <http://www.makelinux.net/ldd3/chp-15-sect-1>

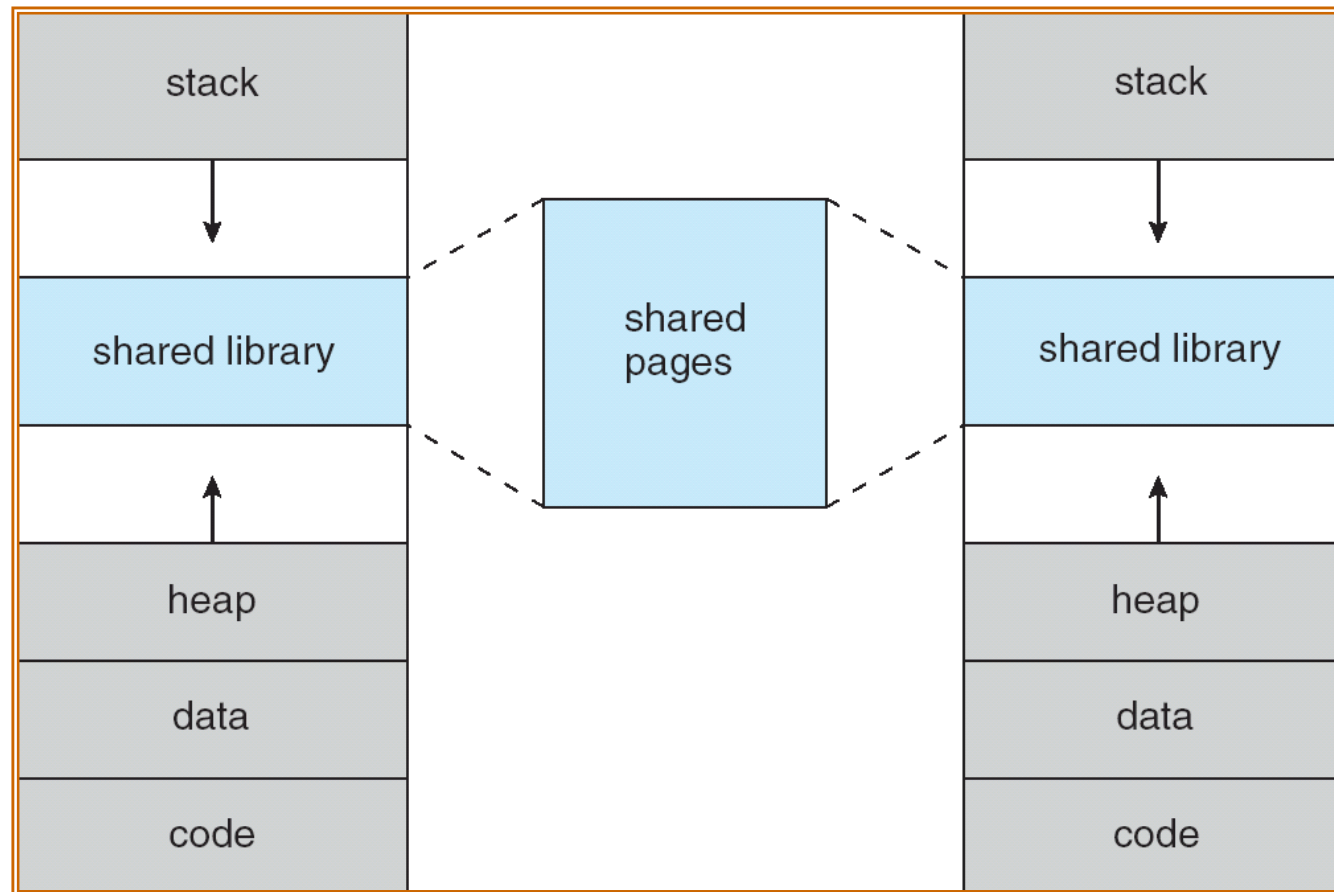
Virtual-address Space

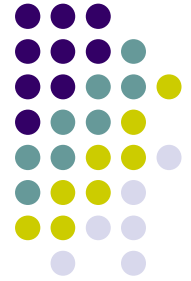


Virtual Memory That is Larger Than Physical Memory



Shared Library Using Virtual Memory






Process Creation

- Virtual memory allows other benefits
 - Copy-on-Write during process creation
 - Memory-Mapped Files

Copy-on-Write




- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory
 - If either process modifies a shared page, only then is the page copied
 - refer to Sect. 2, Chapter 3 of “Lập trình C/C++ ...”
- COW allows more efficient process creation as only modified pages are copied

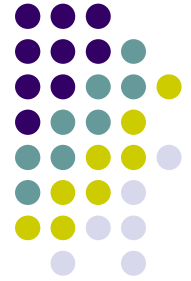
Memory Mapped Files



- A file is considered as a memory segment
- Read/write operations are performed via memory
 - not read/write file system calls
- Allow multiple processes shared a file
- refer to Sect. 5, Chapter 4 of “Lập trình /C++ trên Linux”

Memory Mapped Files

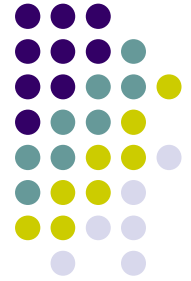




Question

- What is the correct advantage of memory mapped file?
 - A. reduces the task of the system's OS
 - B. treats as the buffer for manipulating the file
 - C. allows programmers to organize the file
 - D. uses as shared resource among processes

Question



- Which of the following is incorrect about virtual memory?
 - A. it is separated from physical memory
 - B. it is mapped into physical memory during process execution
 - C. it gives additional benefits, e.g., COW, file mapping
 - D. an address in virtual memory is preserved when mapped into physical memory

Dynamic loading

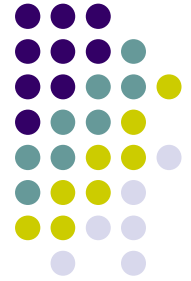


- Routine is not loaded until it is called
- Better memory-space utilization
 - unused routine is never loaded
- Useful when
 - large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system
 - refer to Section 3.4, Chapter 7 of “Lập trình C/C++ ...”

Dynamic linking and shared library




- Linking postponed until execution time
- Small piece of code, *stub*, used to locate the appropriate memory-resident library routine
 - Stub replaces itself with the address of the routine, and executes the routine
 - Operating system needed to check if routine is in processes' memory address
- Dynamic linking is particularly useful for libraries
- Also known as **shared libraries** in Linux
 - refer to Sect. 3, Chapter 7 of “Lập trình C/C++ trên Linux”



Paging on demand

Food order




2/21/2017

Each person prefers different food




Paging on demand




- Bring a page into memory only when needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
- **Lazy swapper** – never swaps a page into memory unless page is needed
 - Swapper that deals with pages is a **pager**


Question



- Why it is possible when only a part of a process is loaded into MEM?
 - A. Because instructions of a process are independent
 - B. Because we can indicate which instructions to run
 - C. Because only one instruction is executed at a time
 - D. Because related instructions are always in the same group

Transfer of a Paged Memory to Contiguous Disk Space






Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated
 - (v \Rightarrow in-memory, i \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to i on all entries
- Example of a page table snapshot:
- During address translation, if valid–invalid bit in page table entry is i \Rightarrow page fault

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

Page Table When Some Pages Are Not in Main Memory




Page Fault




- If there is a reference to a page,
 - first reference to that page will trap to operating system: *page fault*
 - 1. Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
 - 2. Get empty frame
 - 3. Swap page into frame
 - 4. Update the page table
 - 5. Set validation bit = **v**
 - 6. Restart the instruction that caused the page fault

Page Fault (Cont.)

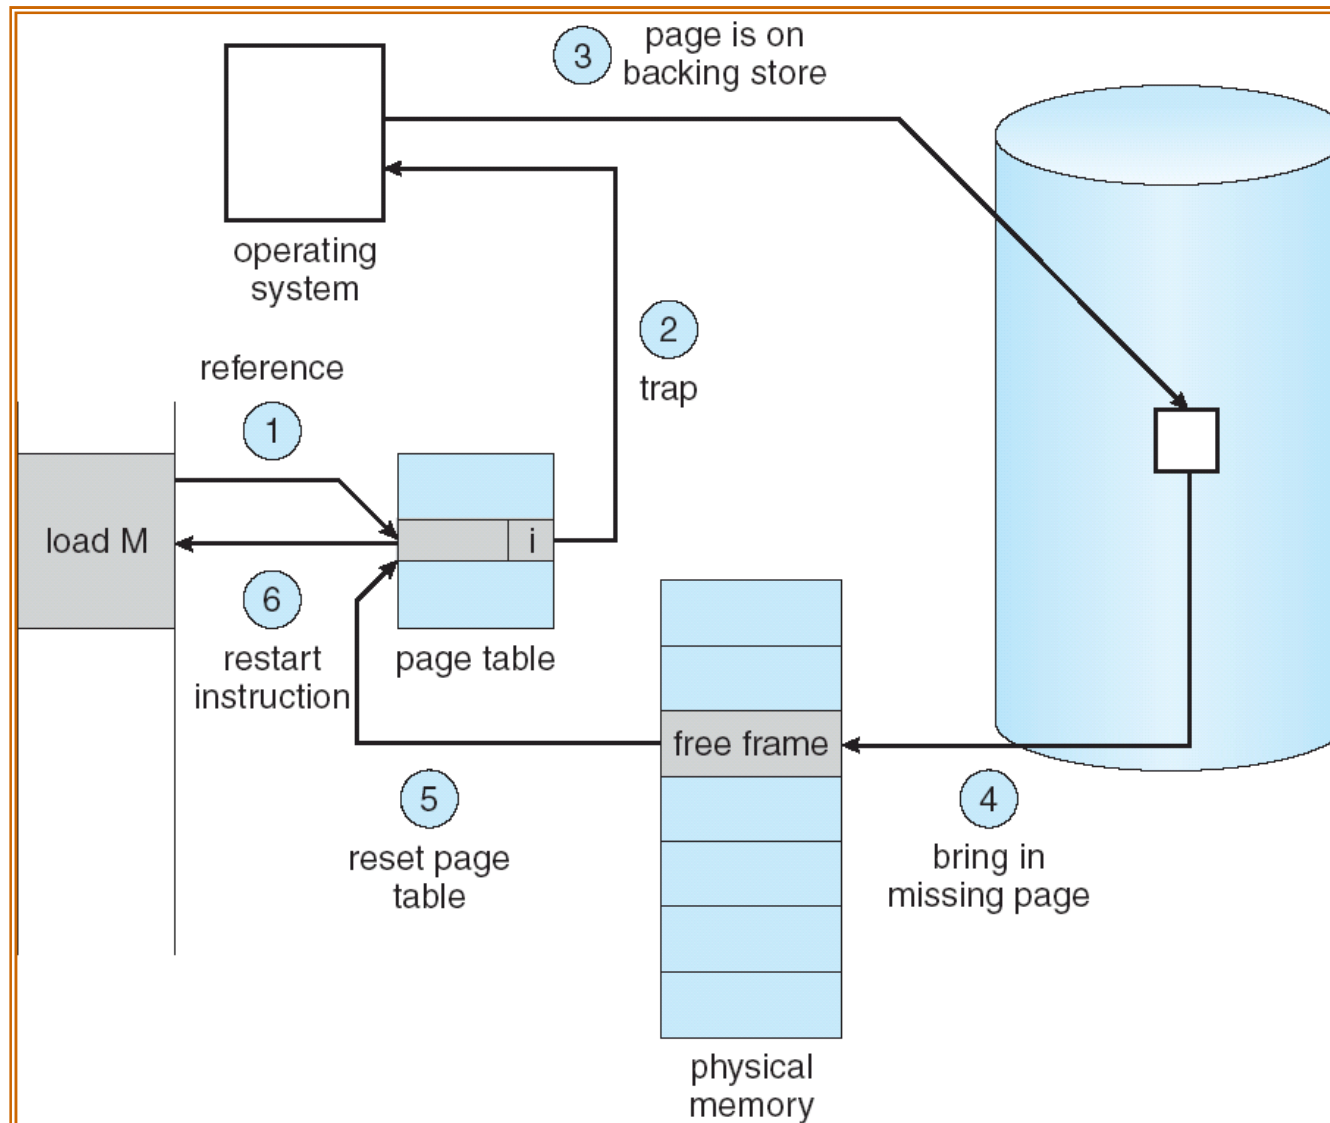
- Restart instruction
 - block move



- auto increment/decrement location



Steps in Handling a Page Fault



Process information

Windows Task Manager

File Options View Shut Down Help

Applications Processes Performance Networking Users

Image Name	User Name	CPU	Mem Usage	Page Faults	VM Size
chrome.exe	t-thanh	08	90,672 K	438,921	86,972 K
Com4QLBEx.exe	SYSTEM	00	2,836 K	719	848 K
wmiprvse.exe	SYSTEM	00	5,108 K	2,132	1,912 K
chrome.exe	t-thanh	00	9,124 K	4,090	5,100 K

ntthanh@turing:-

```
top - 11:00:36 up 14 days, 18:54, 1 user, load average: 0.08, 0.02, 0.01
Tasks: 96 total, 1 running, 95 sleeping, 0 stopped, 0 zombie
Cpu(s): 0.2%us, 0.2%sy, 0.0%ni, 99.2%id, 0.2%wa, 0.3%hi, 0.0%si, 0.0%st
Mem: 1542448k total, 1243240k used, 299208k free, 210424k buffers
Swap: 3112952k total, 0k used, 3112952k free, 737572k cached
```


PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
22076	ntthanh	15	0	90128	1736	1004	S	0.3	0.1	0:00.01	sshd
1	root.	15	0	10348	676	568	S	0.0	0.0	0:02.13	init.
2	root.	RT	-5	0	0	0	S	0.0	0.0	0:01.66	migration/0
3	root.	34	19	0	0	0	S	0.0	0.0	0:00.01	ksoftirqd/0
4	root.	RT	-5	0	0	0	S	0.0	0.0	0:00.00	watchdog/0
5	root.	RT	-5	0	0	0	S	0.0	0.0	0:01.12	migration/1
6	root.	34	19	0	0	0	S	0.0	0.0	0:00.00	ksoftirqd/1
7	root.	RT	-5	0	0	0	S	0.0	0.0	0:00.00	watchdog/1
8	root.	10	-5	0	0	0	S	0.0	0.0	0:02.50	events/0
9	root.	10	-5	0	0	0	S	0.0	0.0	0:00.30	events/1
10	root.	10	-5	0	0	0	S	0.0	0.0	0:00.01	khelper
23	root.	11	-5	0	0	0	S	0.0	0.0	0:00.00	kthread
28	root.	10	-5	0	0	0	S	0.0	0.0	0:00.46	kblockd/0
29	root.	10	-5	0	0	0	S	0.0	0.0	0:00.07	kblockd/1
30	root.	15	-5	0	0	0	S	0.0	0.0	0:00.00	kacpid
99	root.	15	-5	0	0	0	S	0.0	0.0	0:00.00	cqueue/0
100	root.	15	-5	0	0	0	S	0.0	0.0	0:00.00	cqueue/1

☐ Show processes from all users

End Process

Processes: 62 CPU Usage: 14% Commit Charge: 1510M / 4948M

Question




- Which of the following is incorrect about a page fault?
 - A. it happens in paging on demand
 - B. it happens when a reference to a page that is not in MEM
 - C. when a page fault occurs the corresponding process will be terminated
 - D. a page fault handler is called whenever it occurs

Question



- Which of the following is **incorrect order** of steps in handling a page fault?
 - A. check if the valid bit is invalid \Rightarrow raise a page fault
 - B. a page fault is raised \Rightarrow find the page in backing store
 - C. a page fault is raised \Rightarrow find a free frame
 - D. load the page into memory \Rightarrow restart the instruction



If no free frame available

- Call page replacement procedure
 - swap out an unused page from MEM
- Algorithms
 - FIFO, Optimal, LRU, LRU-approximation
- Performance of the algorithm
 - page-fault rate
 - which algorithm is better?

Performance of paging on demand



- Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a page fault
- Effective Access Time (EAT)
$$\text{EAT} = (1 - p) * \text{memory access} \\ + p (\text{page fault overhead} \\ + \text{swap page out} \\ + \text{swap page in} \\ + \text{restart overhead})$$

Paging on Demand Example



- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
 - $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$
 $= (1 - p) \times 200 + p \times 8,000,000$
 $= 200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault
 - $EAT = 8.2 \text{ microseconds.}$
 - slowdown by a factor of 40!!

Page Replacement




- Prevent over-allocation of memory
 - include page replacement in page-fault service routine
- Use **modify (dirty) bit** to reduce overhead of page transfers
 - only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory
 - large virtual memory can be provided on a smaller physical memory

Page Replacement



```
ntthanh@turing:~$ cat /proc/meminfo
HighTotal:        0 kB
HighFree:         0 kB
LowTotal:       1542448 kB
LowFree:        299084 kB
SwapTotal:       3112952 kB
SwapFree:       3112952 kB
Dirty:           24 kB
Writeback:        0 kB
AnonPages:       180092 kB
Mapped:          29504 kB
Slab:            90524 kB
PageTables:       6912 kB
NFS_Unstable:    0 kB
Bounce:          0 kB
CommitLimit:    3884176 kB
Committed_AS:   438656 kB
VmallocTotal: 34359738367 kB
VmallocUsed:      2220 kB
VmallocChunk: 34359735923 kB
HugePages_Total: 0
HugePages_Free:  0
HugePages_Rsvd:  0
Hugepagesize:   2048 kB
[ntthanh@turing ~]$
```

Need For Page Replacement




Basic Page Replacement



1. Find the location of the desired page on disk
2. Find a free frame
 - If there is a free frame, use it
 - Else use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page table
4. Resume the process



Page Replacement



Question



Which of the following is **incorrect** about page replacement?


- A. a victim frame is selected to be swapped out
- B. the page table which is the victim will be updated
- C. the victim frame is always written into the backing store
- D. the victim frame is only written into the backing store if it is dirty

Page Replacement Algorithms



- Want lowest page-fault rate
- Evaluate algorithm
 - run it on a particular string of memory references (reference string)
 - compute the number of page faults on that string
- In all our examples, the reference string is
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Graph of Page Faults Versus The Number of Frames



First-In-First-Out (FIFO) Algorithm



- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

- 3 frames


1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

- 4 frames

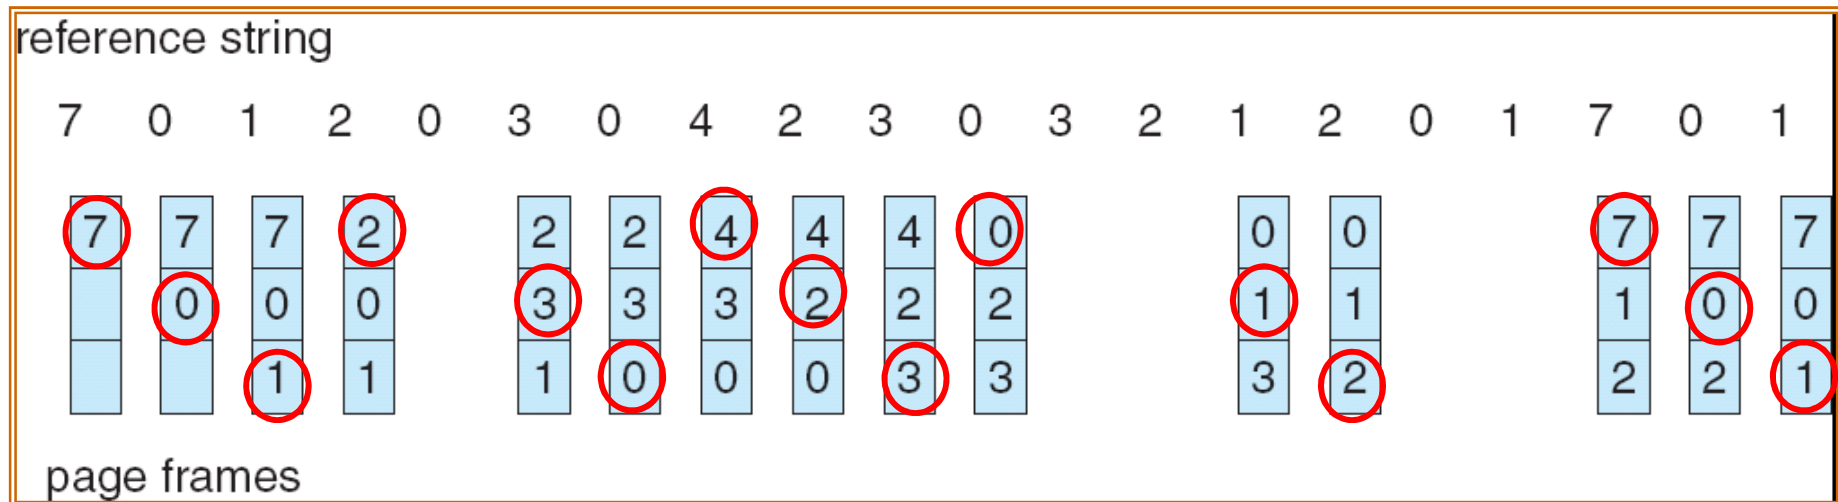
1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		⇒ Belady's anomaly


Belady's Anomaly: more frames ⇒ more page faults

FIFO Illustrating Belady's Anomaly



FIFO Page Replacement





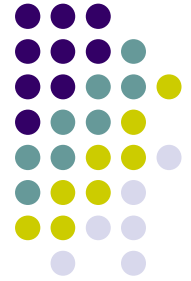
Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, FIFO is used with 3 frames. Which of the following is the correct order of swapped out pages?
 - A. 7 0 1 2 3 0 4 2 3 0 1 2
 - B. 7 0 1 2 3 0 4 3 2 0 1 2
 - C. 7 0 1 3 2 0 4 2 3 0 1 2
 - D. 7 0 1 2 3 0 4 2 3 1 0 2



Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, FIFO is used with 3 frames. Which of the following is the correct number of page faults?
 - A. 13
 - B. 14
 - C. 15
 - D. 16



Optimal Algorithm

- Replace page that **will not** be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

How do you know this?

- Used for measuring **how well** the algorithm performs


1	4
2	
3	
4	5

6 page faults



Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, Optimal algorithm is used with 3 frames. Which of the following is the correct order of swapped out pages?
 - A. 7 1 0 3 4 2
 - B. 7 0 1 4 3 2
 - C. 7 1 0 4 3 2
 - D. 7 1 4 1 3 2



Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, Optimal algorithm is used with 3 frames. Which of the following is the correct number of page faults?
 - A. 8
 - B. 9**
 - C. 10
 - D. 11

Least Recently Used (LRU) Algorithm



- Least recently used page is swapped out first
 - Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**
 - 4 frames

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3



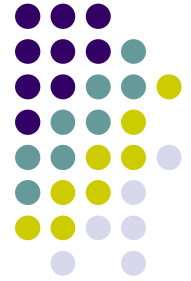
Question

- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, LRU is used with 3 frames.
- Which of the following is the correct order of swapped out pages?
 - A. 7 1 2 3 1 4 1 3 2
 - B. 7 2 1 3 0 4 2 3 2
 - C. 7 1 2 3 0 4 1 2 3
 - D. 7 1 2 3 0 4 0 3 2**



Question


- A reference string 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1, LRU is used with 3 frames.
- Which of the following is the correct number of page faults?
 - A. 13
 - B. 12**
 - C. 11
 - D. 10



LRU Algorithm (Cont.)

- Stack implementation
 - keep a stack of page numbers in a double link form
 - Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
 - No search for replacement

Use Of A Stack To Record The Most Recent Page References

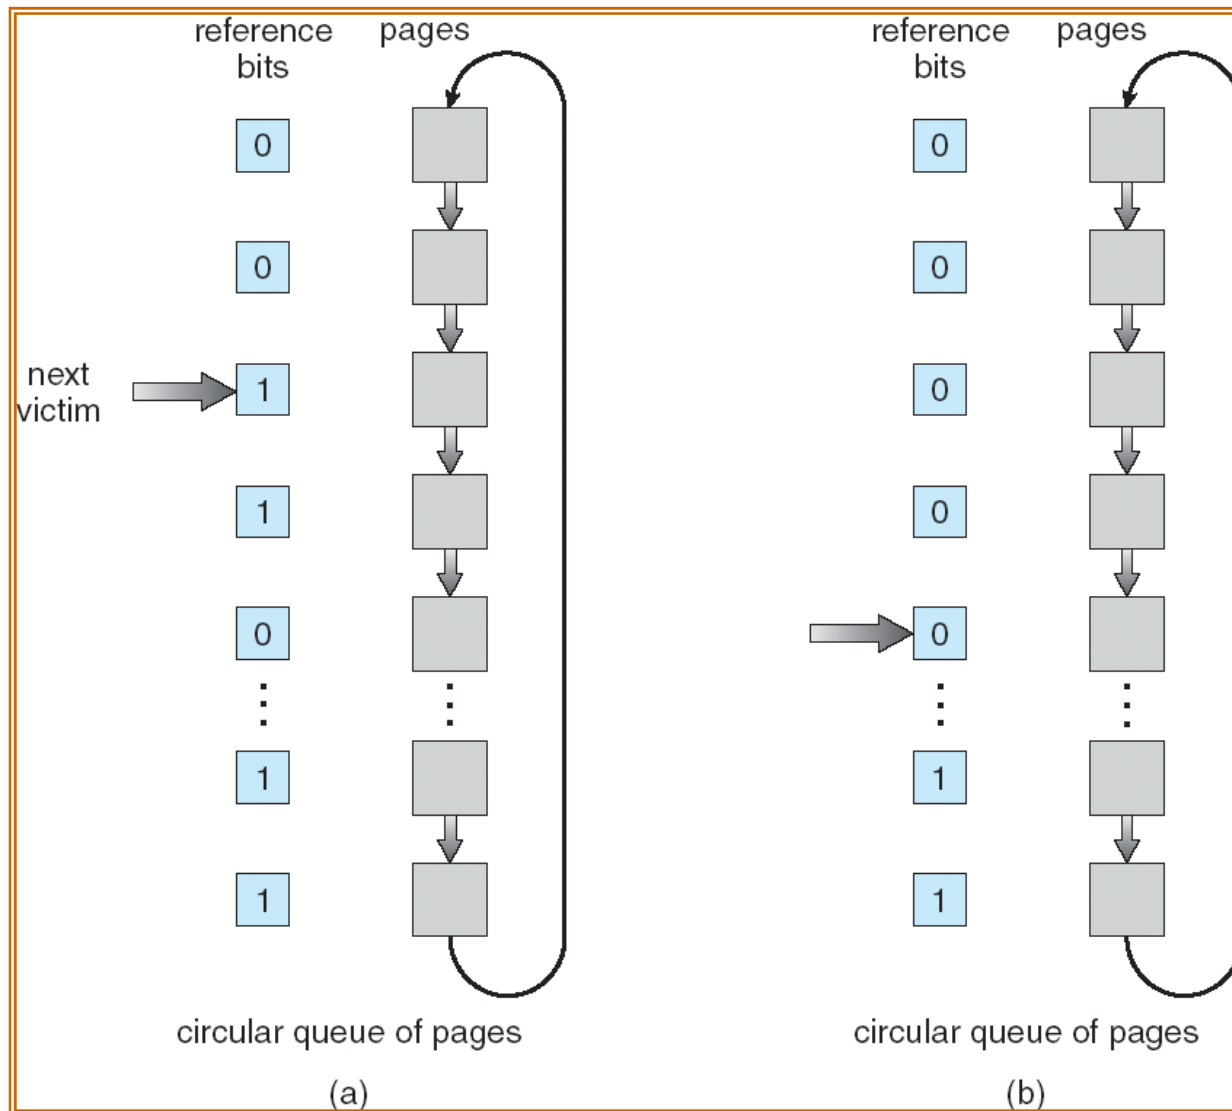



LRU Approximation Algorithms



- Reference bit
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace the one which is 0 (if one exists)
 - We do not know the order
- Second chance (follow clock order)
 - Need reference bit
 - Clock replacement
 - If page to be replaced has reference bit = 1 then:
 - set reference bit 0
 - leave page in memory
 - replace next page, subject to same rules

Second-Chance (clock) Page-Replacement Algorithm





Question


- Suppose the second chance is used;
 - the reference bits of frames are: 1 1 0 1 1 0
 - the head is at second frame
- Which of the following are the reference bits after a page replacement is done
 - A. 0 0 0 1 1 0
 - B. 1 0 1 1 1 0
 - C. 1 0 0 1 1 0
 - D. 1 0 1 0 1 0

Least frequently Used (LFU) Algorithm



- Counter implementation
 - Every page entry has a counter;
 - every time page is referenced, copy the clock into the counter
 - When a page needs to be swapped
 - look at the counters to determine

Counting Algorithms



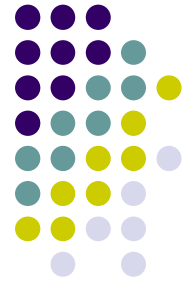
- Keep a counter of the number of references that have been made to each page
- **Least Frequently Used (LFU) Algorithm**
 - replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**
 - based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Allocation of Frames



- Each process needs *minimum* number of pages
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - instruction is 6 bytes, might span 2 pages
 - 2 pages to handle *from*
 - 2 pages to handle *to*
- Two major allocation schemes
 - fixed allocation
 - priority allocation

Fixed Allocation



- Equal allocation
 - For example, if there are 100 frames and 5 processes, give each process 20 frames.
- Proportional allocation
 - Allocate according to the size of process

- s_i = size of process p_i
- $S = \sum s_i$
- m = total number of frames
- a_i = allocation for $a_i = \frac{s_i}{S} \times m$

$$m = 64$$

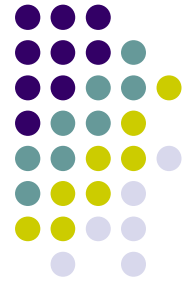
$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

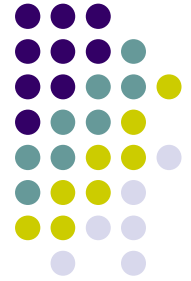
$$a_2 = \frac{127}{137} \times 64 \approx 59$$

Question



- A system uses proportional allocation and has
 - 90 frames x 2KB
 - 3 processes with size of (138KB, 96KB, 164KB)
- Which of the following is the correct number of allocated frames of (P_1 , P_2 , P_3)
 - A. 32, 21, 37
 - B. 31, 22, 37
 - C. 30, 22, 38
 - D. 33, 22, 35

Priority Allocation



- Use a proportional allocation scheme using priorities rather than size
- If process P_i generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

Global vs. Local Allocation




- **Global replacement**
 - process selects a replacement frame from the set of all frames;
 - one process can take a frame from another
- **Local replacement**
 - each process selects from only its own set of allocated frames

Thrashing



- If a process does not have “enough” frames, the page-fault rate is very high.
- **Thrashing**
 - a process is busy swapping pages in and out
- This caused by:
 - low CPU utilization
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system

Thrashing (Cont.)



Solutions to Thrashing

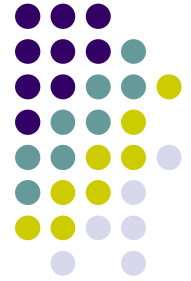


- Use local allocation
- Use priority allocation
 - not good solution
- Working set model
 - A suitable solution

Question



- Which of the following is incorrect about priority allocation?
 - A. higher priority process is allocated first
 - B. it prevents thrashing from happening
 - C. frames are allocated globally
 - D. it does not prevent thrashing from happening




Working-Set Model

- $\Delta \equiv$ working-set window
 - a number of page references, e.g. 10,000
- Working set of Process P_i
 - WSS_i = total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m$ (total of frames) \Rightarrow Thrashing
- Policy if $D > m$, then suspend one process



Working-set model



Keeping Track of the Working Set



- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page
 - Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - If one of the bits in memory = 1 \Rightarrow page in working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units

Question



- Suppose a delta = 10; reference string
 - 2 6 1 5 7 7 7 7 5 1 6 2 3 4 4 4 3 4 4 4 1 3 2 3 4 4 4 3
4 4 4 ...
- Which of the following is the correct WSS at 20th reference?
 - A. {2 3 4 6}
 - B. {2 3 4 5 6}
 - C. {1 2 3 4 6}
 - D. {7 1 2 3 4 6}



Question?