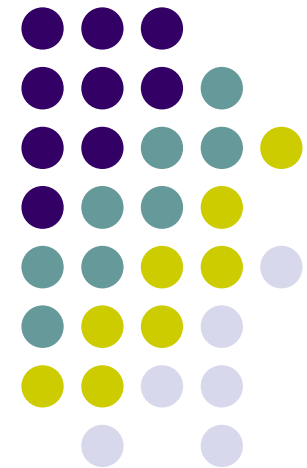


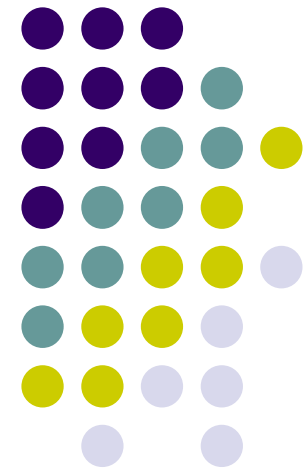
Operating System

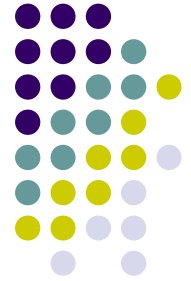
Nguyen Tri Thanh
ntthanh@vnu.edu.vn



File System Implementation

File-System Structure
File-System Implementation
Directory Implementation
Allocation Methods
Free-Space Management
Efficiency and Performance
Recovery






Objectives

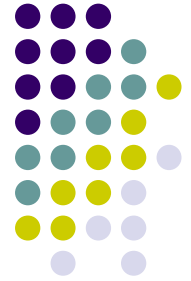
- Introduce what a file is
- Introduce file system implementation
- Introduce directory implementation
- Introduce 3 allocation methods
- Introduce free space management

Reference



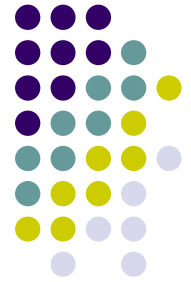
- Chapter 10, 11 of **Operating System Concepts**

File Concept



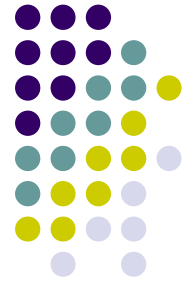
- Contiguous logical address space
- Types:
 - Data
 - numeric
 - character
 - binary
 - Program

File Structure



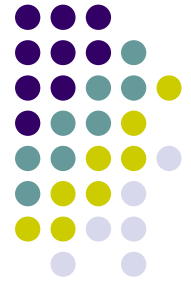
- None - sequence of words, bytes
- Simple record structure
 - Lines
 - Fixed length
 - Variable length
- Complex Structures
 - Formatted document
 - Relocatable load file
- Can simulate last two with first method by inserting appropriate control characters
- Who decides:
 - Operating system
 - Program

File Attributes



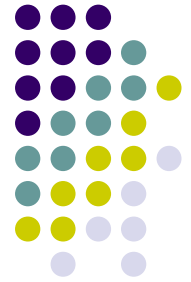
- Name
 - only information kept in human-readable form
- Identifier
 - unique (number) identifies file within file system
- Type
 - needed for systems that support different types
- Location
 - pointer to file location on storage device

File Attributes (cont'd)

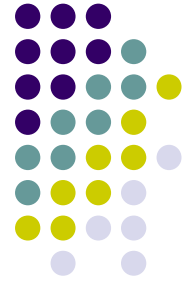


- Size
 - current file size
- Protection
 - controls who can do reading, writing, executing
- Time, date, and user identification
 - data for protection, security, and usage monitoring
- Information about files are kept in the **directory structure** on the disk

File System Structure



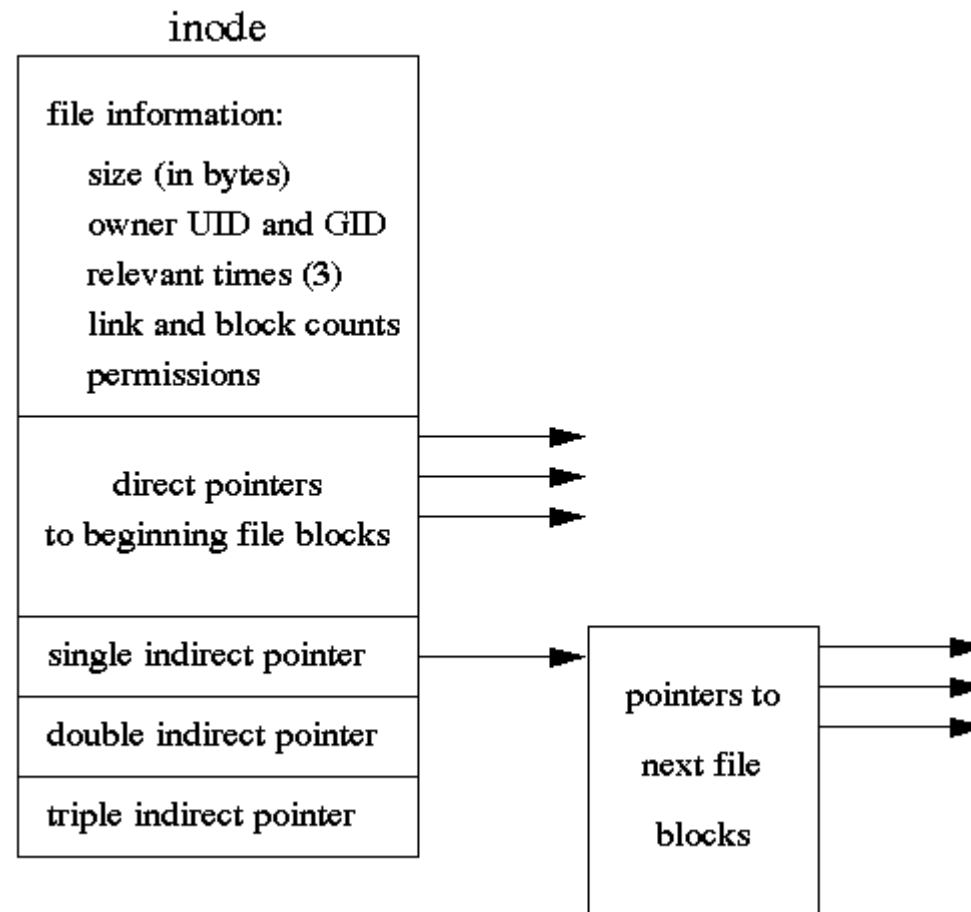
- File system resides on secondary storage
 - disks, CD ROM, DVD, flash drive,...
- File system organized into layers
- File control block (FCB)
 - storage structure of information about a file
 - *inode* (index node) is an implementation of FCB on Linux

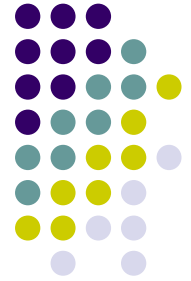


A Typical File Control Block

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

Inode on Linux






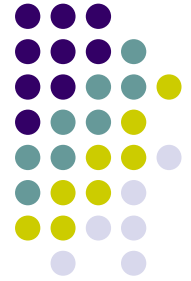
Question

- Which of the following is incorrect about file control block (FCB)?
 - A. a data structure containing information of a file
 - B. OS needs a FCB to access a file
 - C. FCB of a file is updated when a file is accessed
 - D. FCBs reside in memory

Layered File System

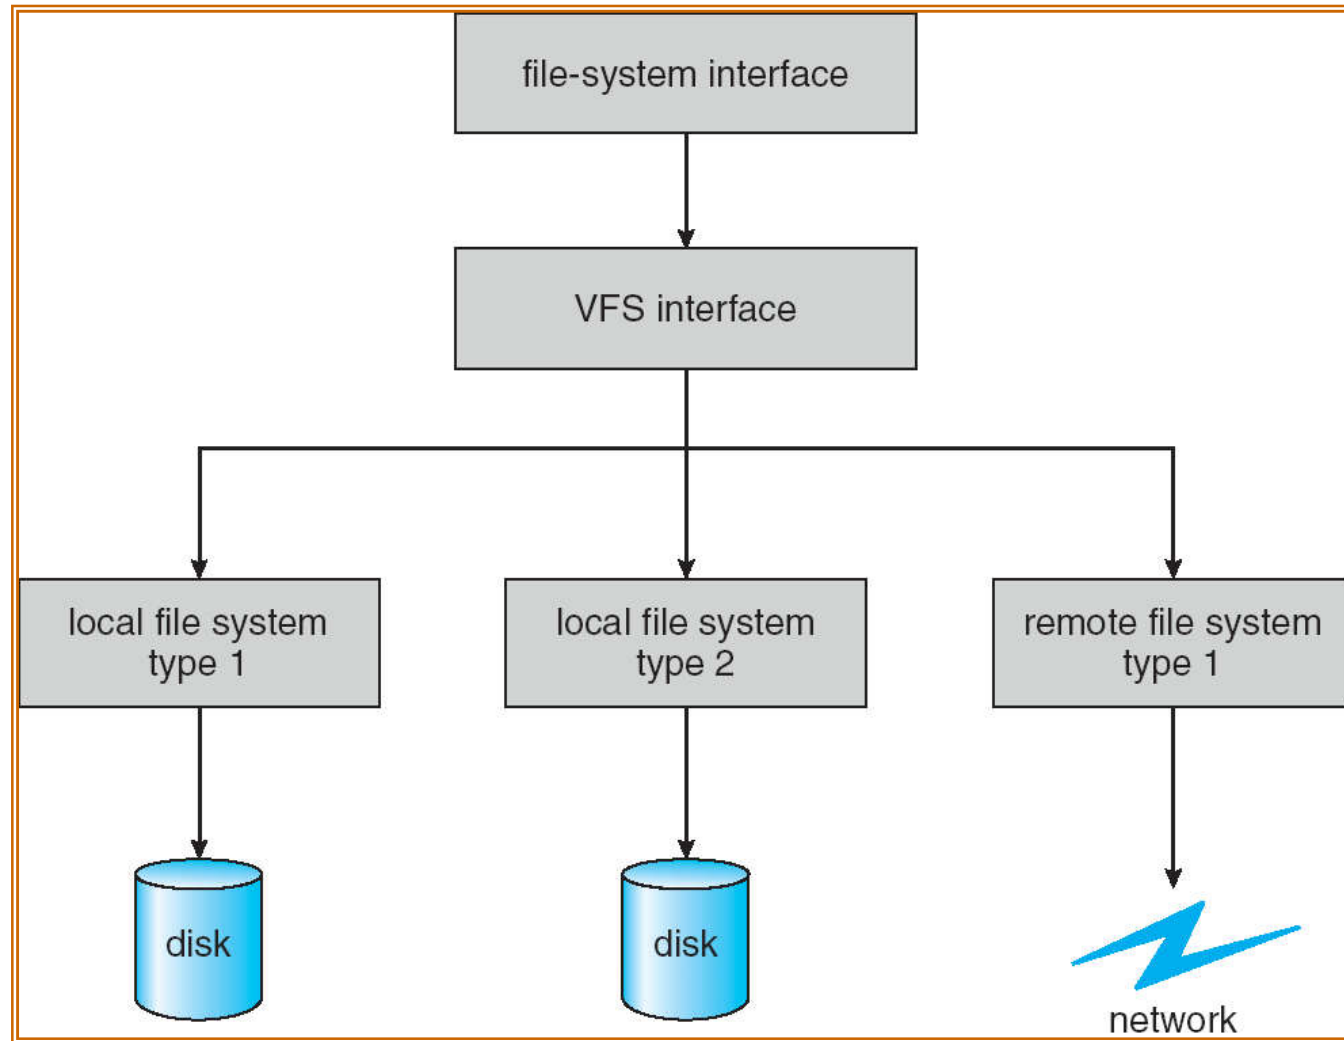


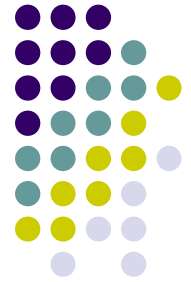
Virtual File Systems



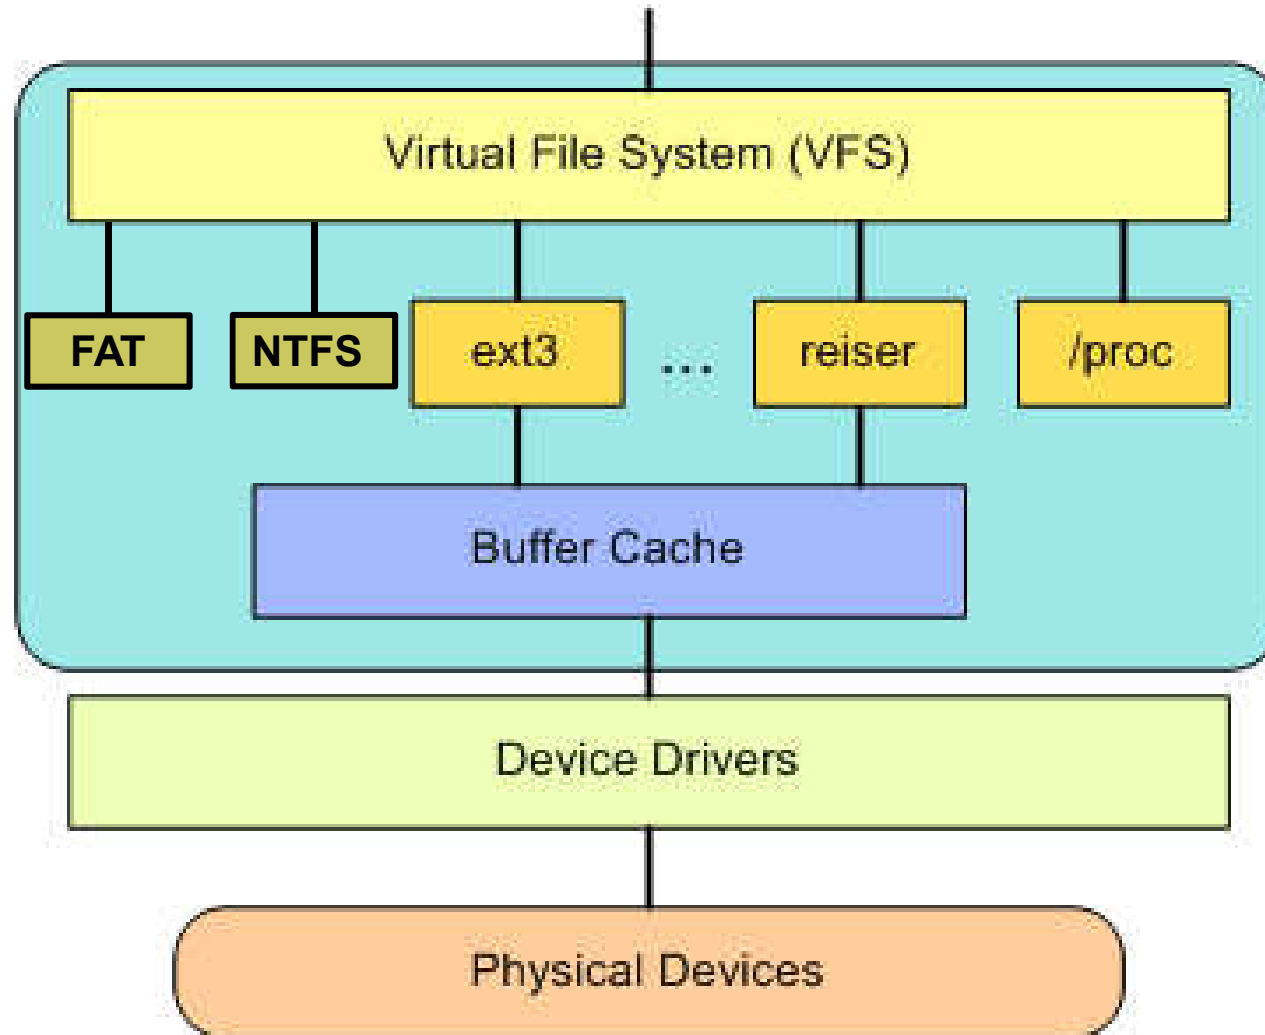
- Virtual File Systems (VFS)
 - provide an object-oriented way of implementing file systems
 - allow the same system call interface (the API) to be used for different types of file systems
 - the API is to the VFS interface, rather than any specific type of file system.

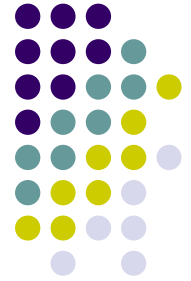
Schematic View of Virtual File System






Linux VFS





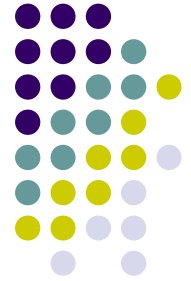
Question

- Which of the following is incorrect about VFS?
 - A. VFS allows an OS to support many different file systems
 - B. VFS provides the same API for all file systems
 - C. VFS is available in all OSes
 - D. VFS hides the detailed implementation of each file system from programmers



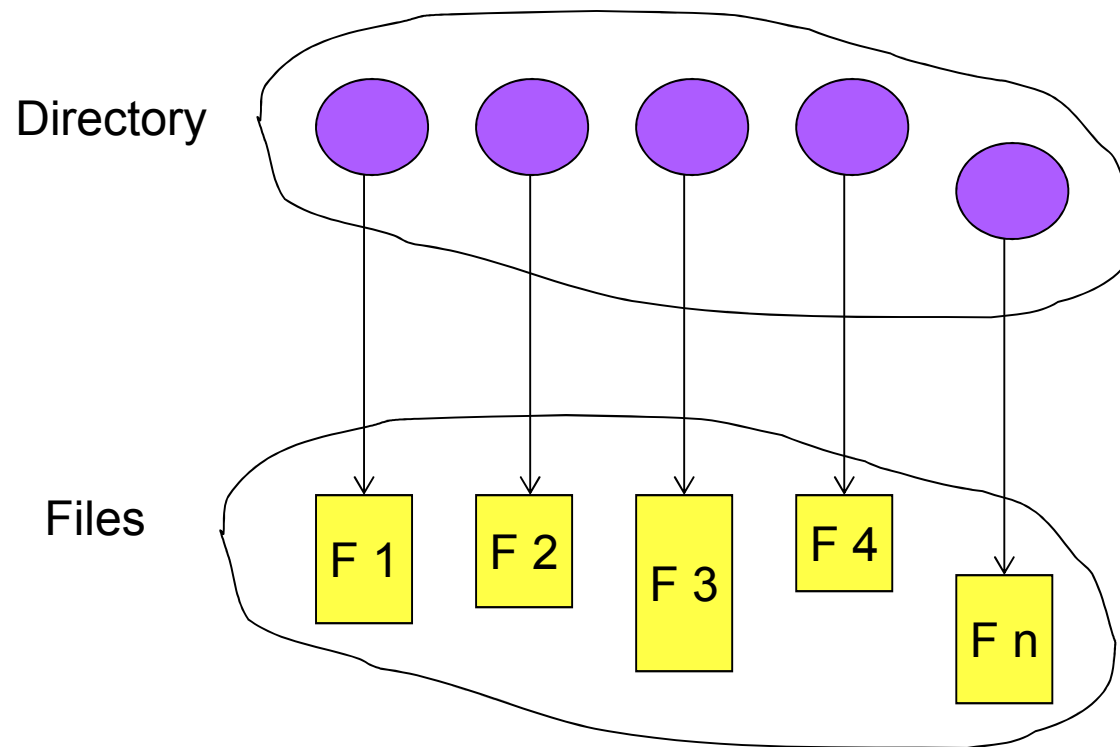
Question

- Which of the following is a correct description of a directory?
 - A. a directory is a disk partition to contain files
 - B. a directory is actually a file containing partial information about its files
 - C. a directory is a container of files' data
 - D. a directory contains the FCB and data of files



Directory Structure

- A collection of nodes containing information about all files



Directory Implementation



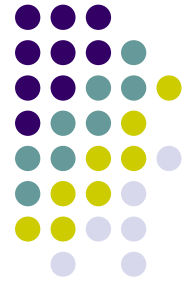
- **Linear list**

- list of file names with pointer to the data blocks
- simple to program
- time-consuming to execute
- FAT http://en.wikipedia.org/wiki/File_Allocation_Table
- Linux FS (Ext3)

- **Hash Table**

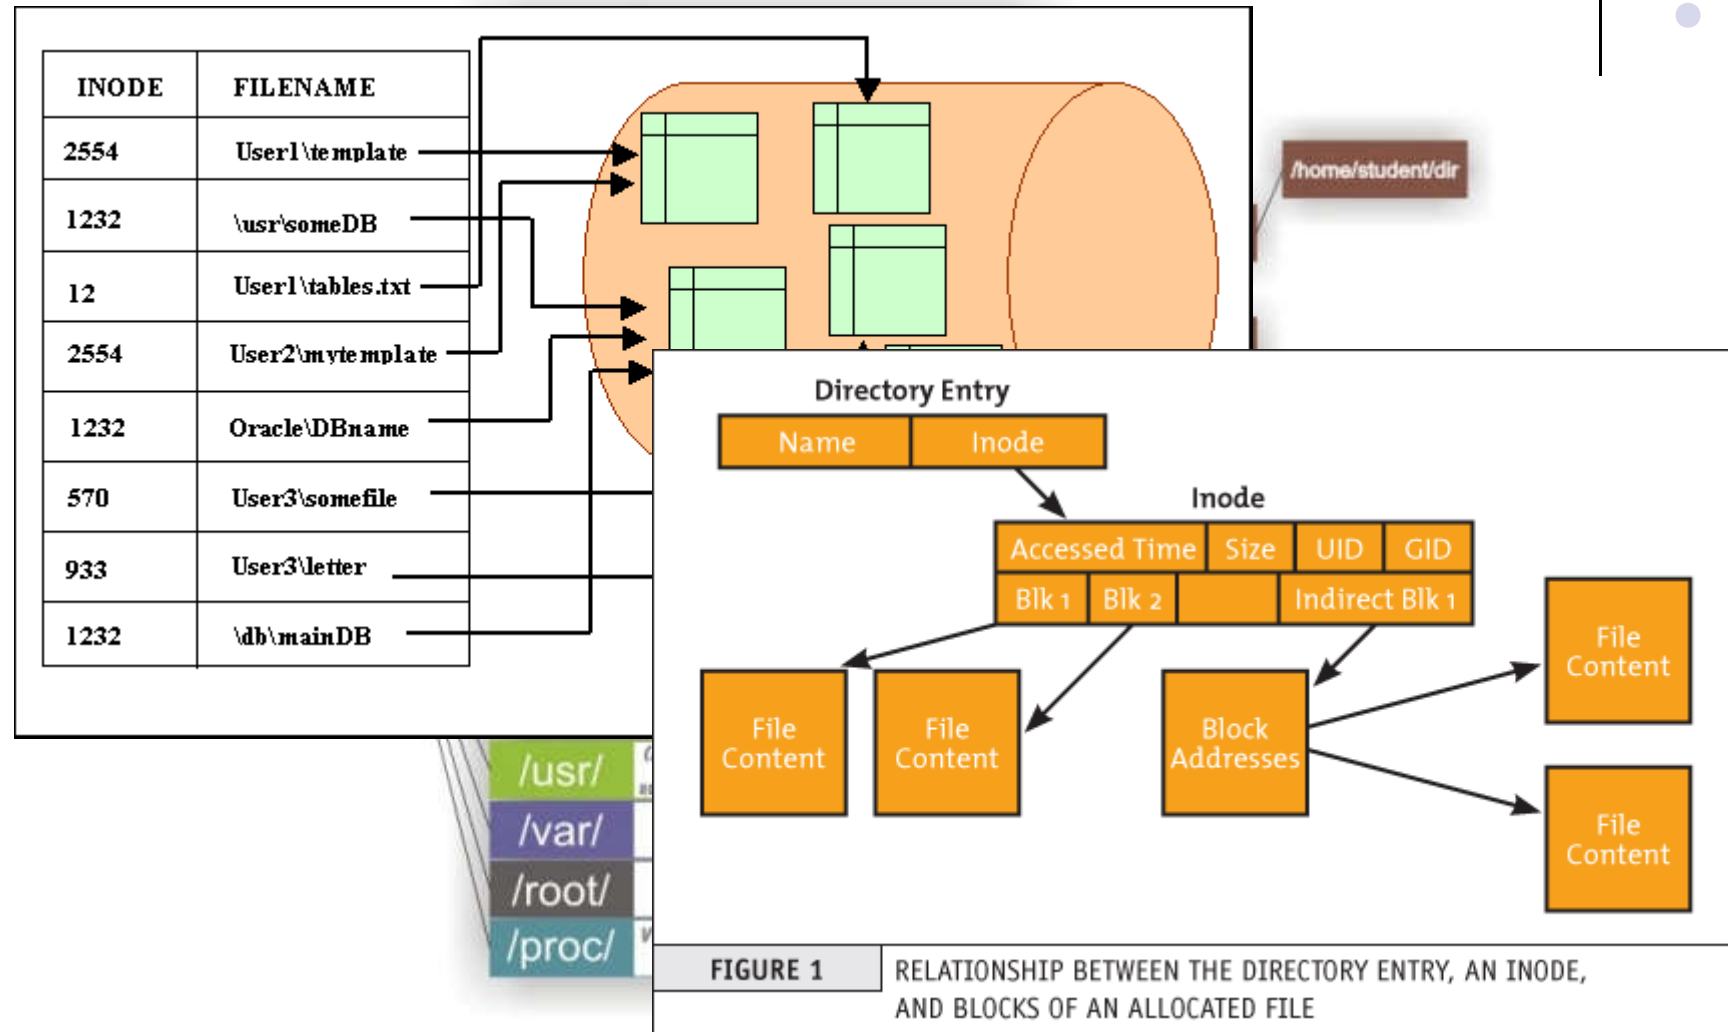
- linear list with hash data structure
- decreases directory search time
- **collisions** resolution
- fixed size

Directory Implementation



- **Balanced binary tree**
 - RAISERFS <http://en.wikipedia.org/wiki/ReiserFS>


Directory in Linux





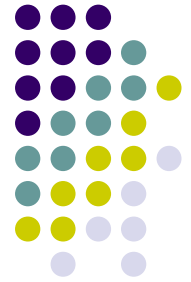
File management API

File Operations




- Create
- Write
- Read
- Reposition within file
- Delete
- Truncate
- $Open(F_i)$
 - search the directory structure on disk for entry F_i , and move the content of entry to memory
- $Close(F_i)$
 - move the content of entry F_i in memory to directory structure on disk


Operations Performed on Directory



- Search for a file
- Create a directory
- Delete a directory
- List a directory
- Rename a directory
- Traverse the file system




File/directory protection



Protection

- File owner/creator should be able to control:
 - what can be done
 - by whom
- Types of access
 - Read
 - Write
 - Execute
 - Append
 - Delete
 - List

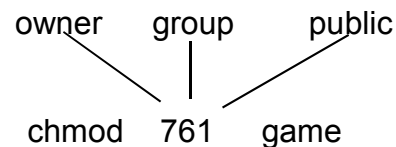
Access Lists and Groups



- Mode of access: read, write, execute
- Three classes of users

			RWX
a) owner access	7	⇒	1 1 1
			RWX
b) group access	6	⇒	1 1 0
			RWX
c) public access	1	⇒	0 0 1

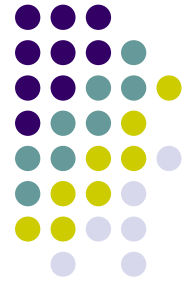
- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say *game*) or subdirectory, define an appropriate access.



Attach a group to a file

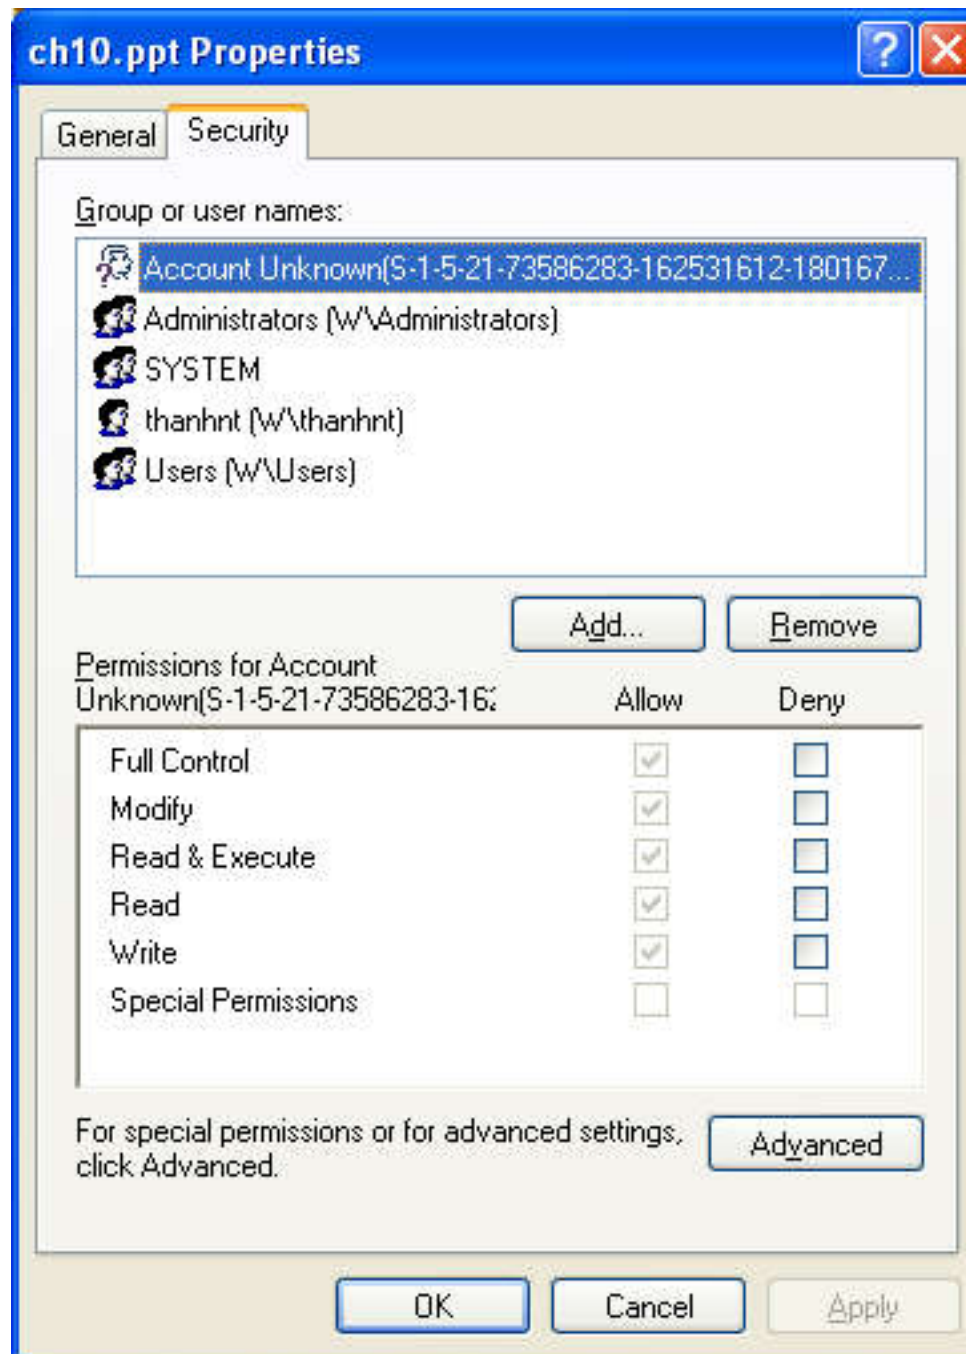
chgrp G game

Question



- Suppose a file has access mode 664. Which is correct ?
 - A. Any user can execute the file
 - B. Users of the owner group can execute the file
 - C. Any user can read the file
 - D. The owner cannot write the file

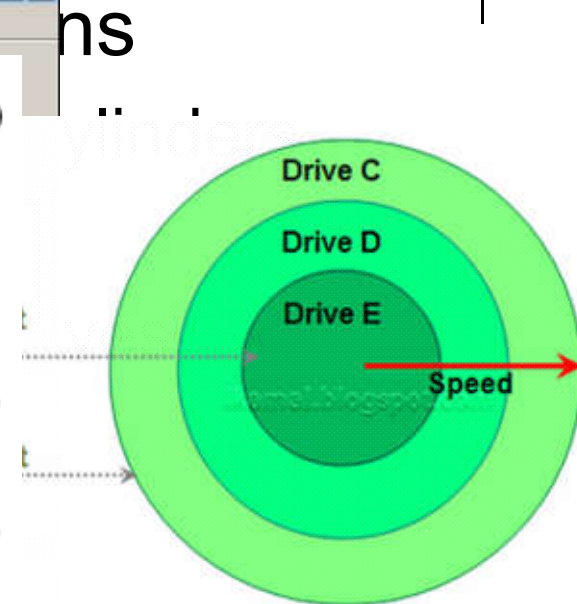
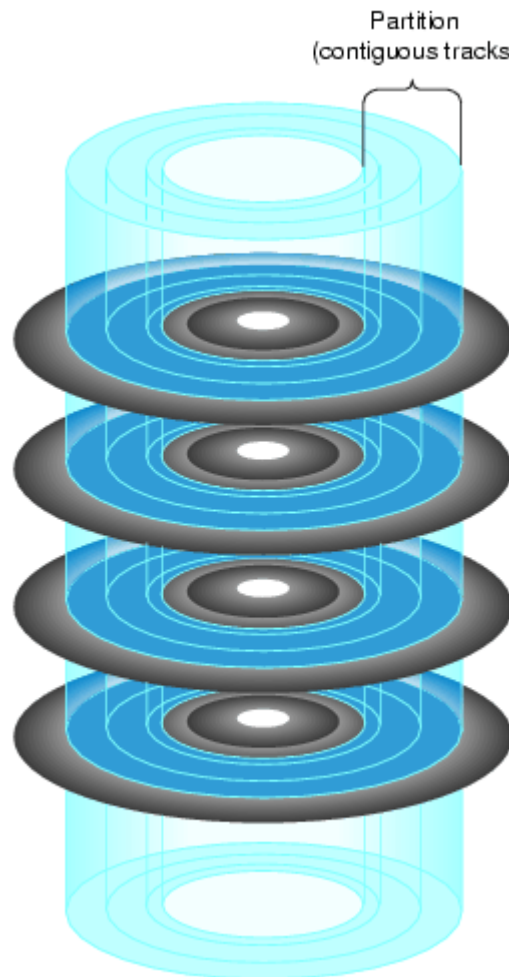
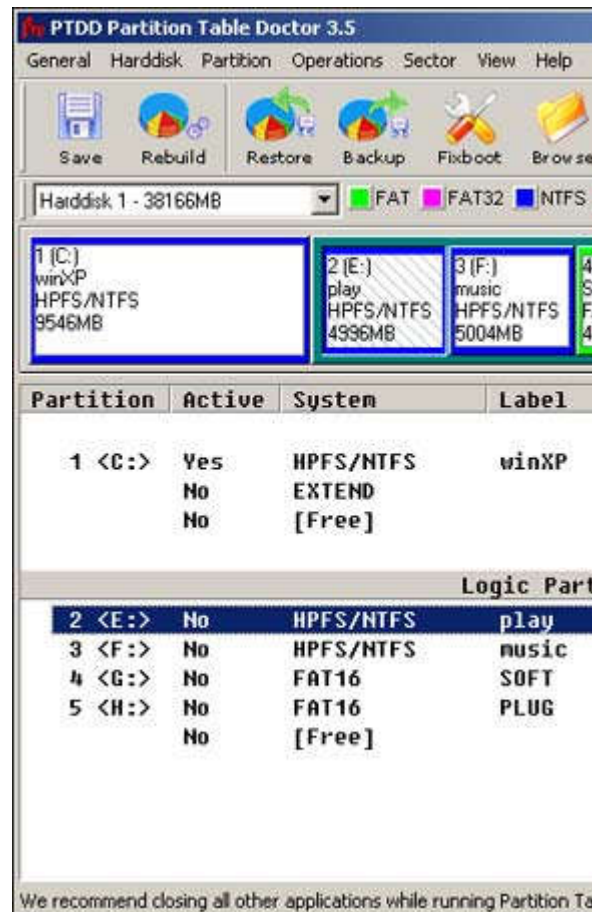

Windows XP Access-control List Management





Storage Allocation Methods


Disk partitions



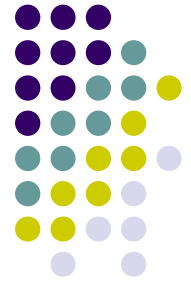


Partition organization

- Organization is specific to each OS
 - region for storing data is divided into equal **blocks**



Allocation Methods

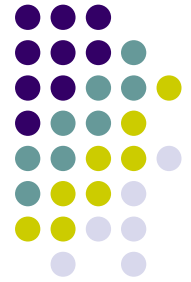


- An allocation method refers to how disk blocks are allocated for files
 - Contiguous allocation
 - Linked allocation
 - Indexed allocation
- Each method needs an appropriate way to access file

Contiguous Allocation



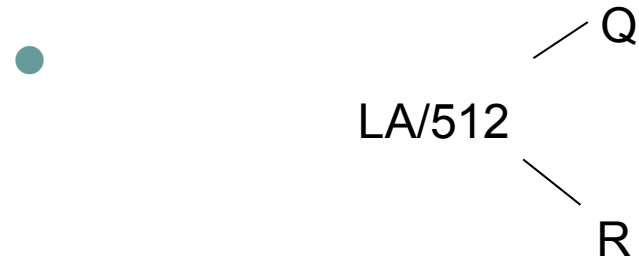
- Each file occupies a set of contiguous blocks on the disk
- Simple
 - only starting location (block #) and length (number of blocks) are required
- Random access
- Wasteful of space
 - dynamic storage-allocation problem
- Files cannot grow



Contiguous Allocation

- Mapping from logical to physical

- LA is position to access

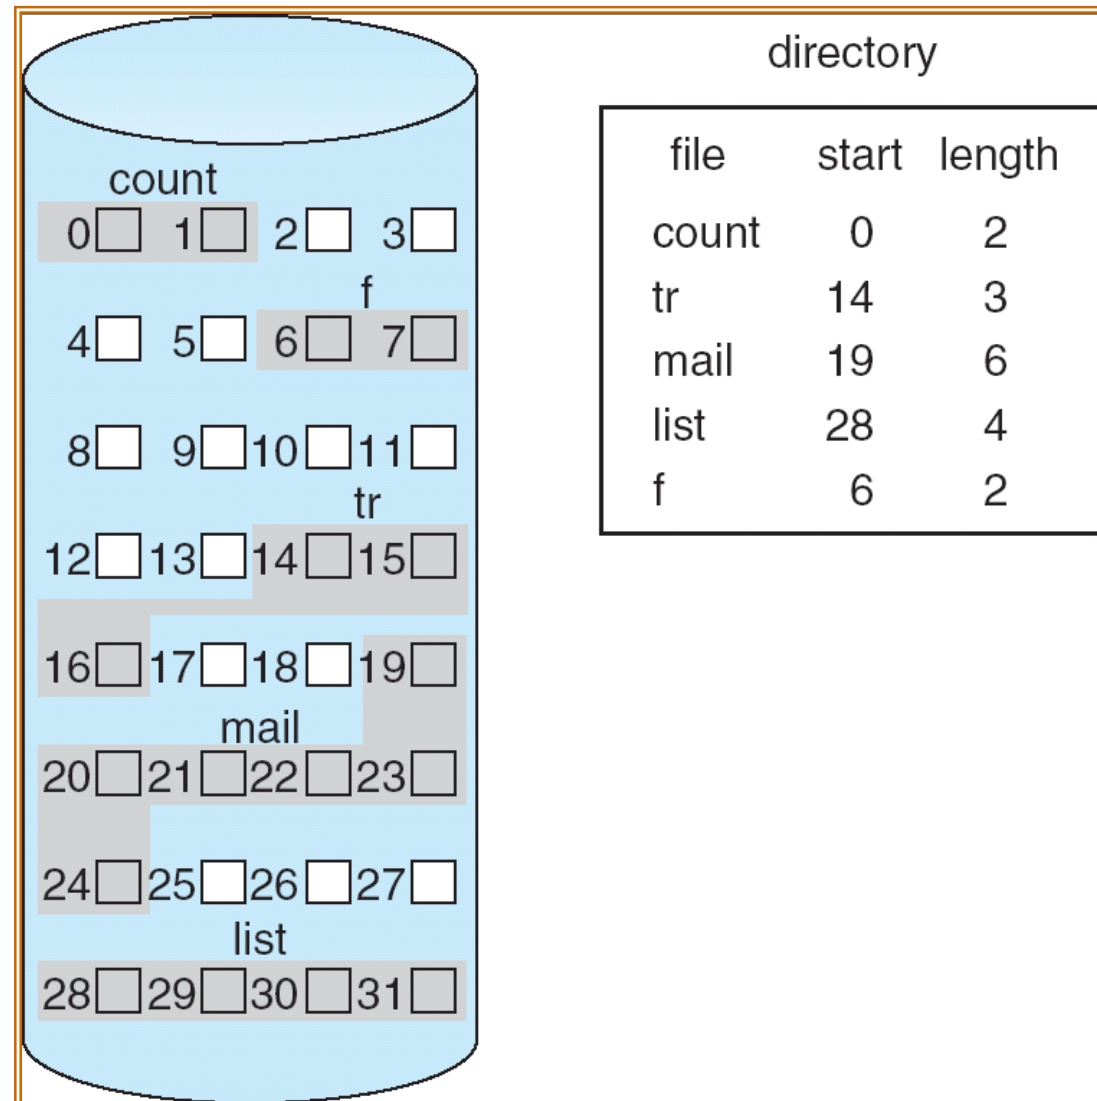



Suppose block size is 512KB

Block to be accessed = $Q + \text{starting address}$

Displacement/offset into block = R


Contiguous Allocation of Disk Space





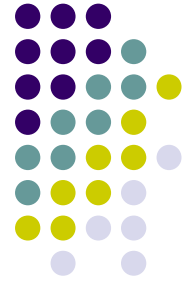
Question

- A system using contiguous allocation
 - block size is 2KB
 - a file has the length of 12.3MB
- Which of the following is the correct location of file at position 50.5KB
 - A. block 25, offset 512
 - B. block 24, offset 1024
 - C. block 25, offset 510
 - D. block 24, offset 2512



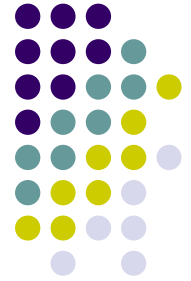
Extent-Based Systems

- Many newer file systems (I.e. Veritas File System)
 - use a modified contiguous allocation scheme
- Extent-based file systems
 - allocate disk blocks in **extents**
- An **extent** is a contiguous block of disks
 - Extents are allocated for file allocation
 - A file consists of one or more extents
 - extents of a file are not necessarily contiguous



Question

- An extent-based file system
 - an extent has *100* blocks
 - a block has size *2KB*
 - a file has size *25.3MB*
- Which of the following is the correct extent number (started by 0) of file at position 15MB?
 - A. 74
 - B. 75
 - C. 76
 - D. 77



Question

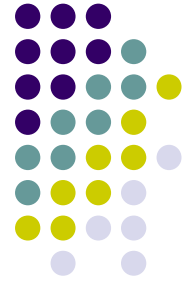
- An extent-based file system
 - an extent has *100* blocks
 - a block has size *2KB*
 - a file has size *25.3MB*
- Which of the following is the correct block (started by 0) number in the extent containing the data of file at position *15MB*?

A. 77

B. 78

C. 79

D. 80



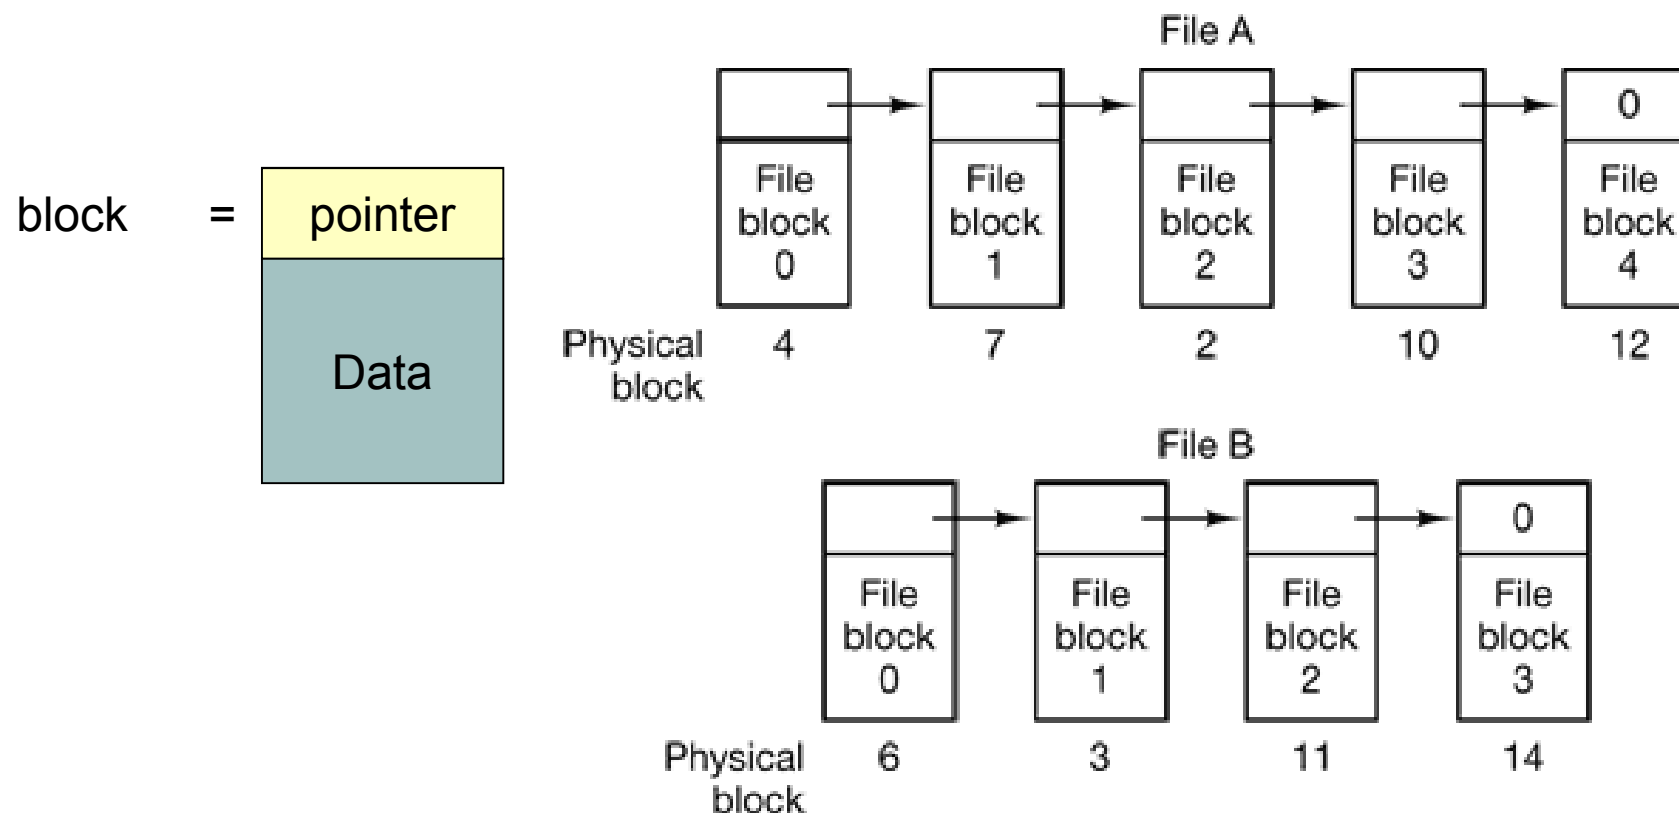
Question

- An extent-based file system
 - an extent has *100* blocks
 - a block has size *2KB*
 - a file has size *25.3MB*
- Which of the following is the correct offset of the of file at position *15MB* in the block containing data?
 - A. 0
 - B. 2047
 - C. 2048
 - D. 512




Linked Allocation

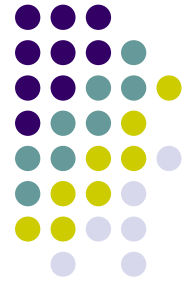
- Each file is a linked list of disk blocks:
 - blocks may be scattered anywhere on the disk.



Linked Allocation

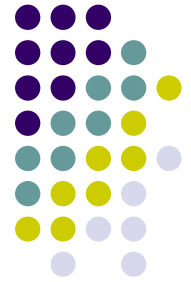


Question



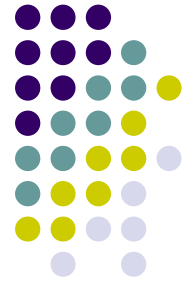
- A system uses linked list allocation
 - partition size 500GB (1GB=1024MB, 1MB=1024KB,...)
 - block size 1KB
- Which of the following is the correct size of the pointer of each block?
 - A. short (2 bytes)
 - B. int (4 bytes)
 - C. float (4 bytes)
 - D. long (8 bytes)

Question



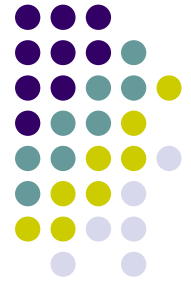
- A system uses linked list allocation
- Which of the following is the correct reason for **no direct access**?
 - A. because we don't know the location of block n directly
 - B. because data blocks of a file may be scattered
 - C. because of security reason
 - D. because the information of data block location is hidden

Question



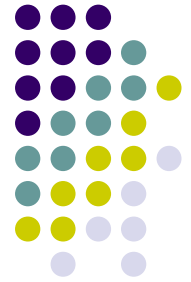
- A system uses linked list allocation
- Which of the following is the correct way to access block n of a file?
 - A. read the first block to find the location of block n
 - B. look up the location of block n from a table
 - C. recursively read block $n-1$ to find out the location of block n
 - D. there is no way to find the location of block n

Linked Allocation (Cont.)



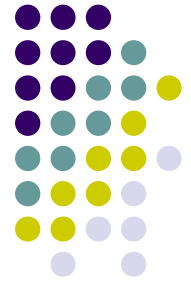
- Simple
 - need only starting address
- Free-space management system
 - no waste of space
- No random access
 - File-allocation table (FAT) – disk-space allocation used by MS-DOS and OS/2.

Question



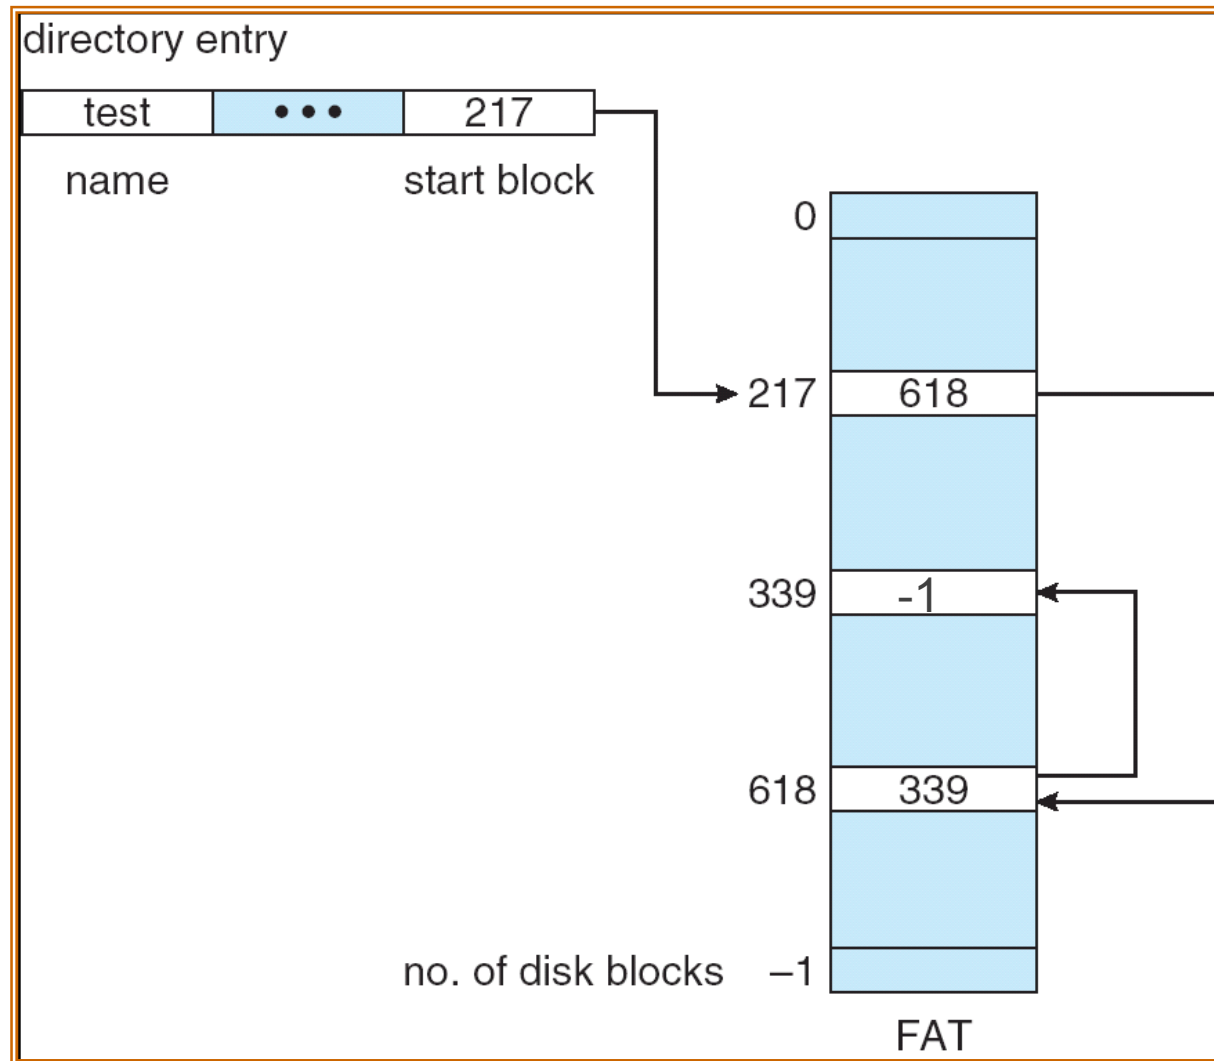
- A system uses linked list allocation
 - block size 2KB
 - pointer size 4 bytes
 - file size 15.4MB
- Which of the following is the correct block of file at position 15.25KB?
 - A. 7
 - B. 8
 - C. 9
 - D. 10

Question



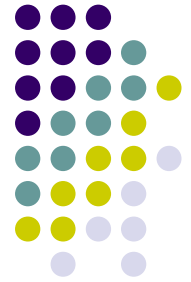
- A system uses linked list allocation
 - block size 2KB
 - pointer size 4 bytes
 - file size 15.4MB
- Which of the following is the correct offset in the block of file at position 15.25KB?
 - A. 1311
 - B. 1312
 - C. 1313
 - D. 1314

File-Allocation Table (DOS)



- Separate the pointers from data
- pointers are stored in File Allocation Table (FAT)
- The system stores two identical copies of FAT

Question

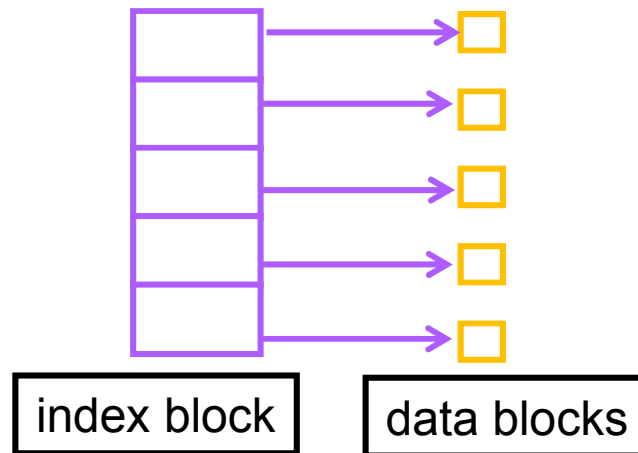


- Which of the following is incorrect about FAT?
 - A. it is fast to access the block n of a file
 - B. it is slow to access the block n of a file
 - C. if FAT is corrupted the whole partition is corrupted
 - D. the system keeps two copies of FAT in order to reduce the risk of FAT corruption

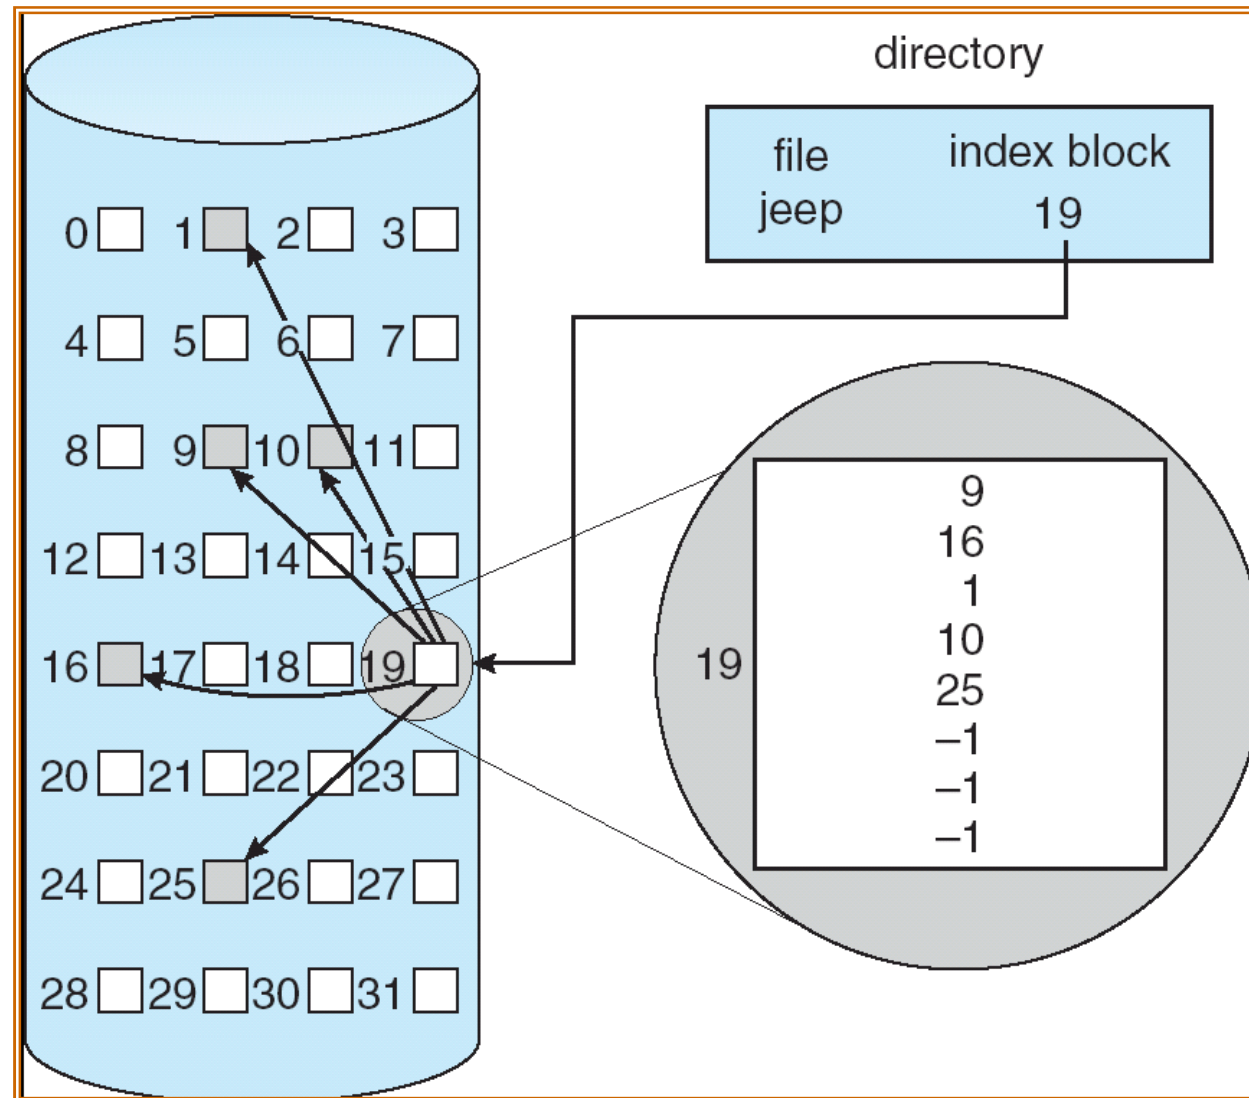
Indexed Allocation



- Brings all pointers together into the *index block*
- Logical view

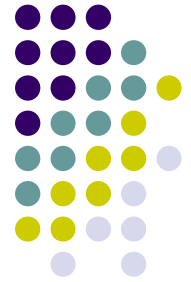


Example of Indexed Allocation

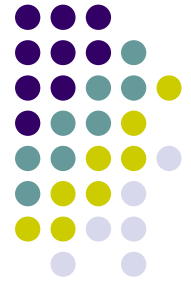


Indexed Allocation (Cont.)


- Need index block
- Random access
- Dynamic allocation without external fragmentation
 - have overhead of index block



Question



- System uses indexed allocation
 - block size 4KB
 - pointer size 4 bytes
- Which of the following is the correctly maximum file size?
 - A. 4MB
 - B. 8MB
 - C. 16MB
 - D. 32MB




Question

- System uses indexed allocation
 - block size 4KB
 - pointer size 4 bytes
 - file size 3MB
- Which of the following is the correct block (starting from 0) and offset at file position 35KB?
 - A. $(block, offset) = (9, 3071)$
 - B. $(block, offset) = (9, 3070)$
 - C. $(block, offset) = (8, 3072)$
 - D. $(block, offset) = (8, 3070)$



Indexed Allocation (cont'd)

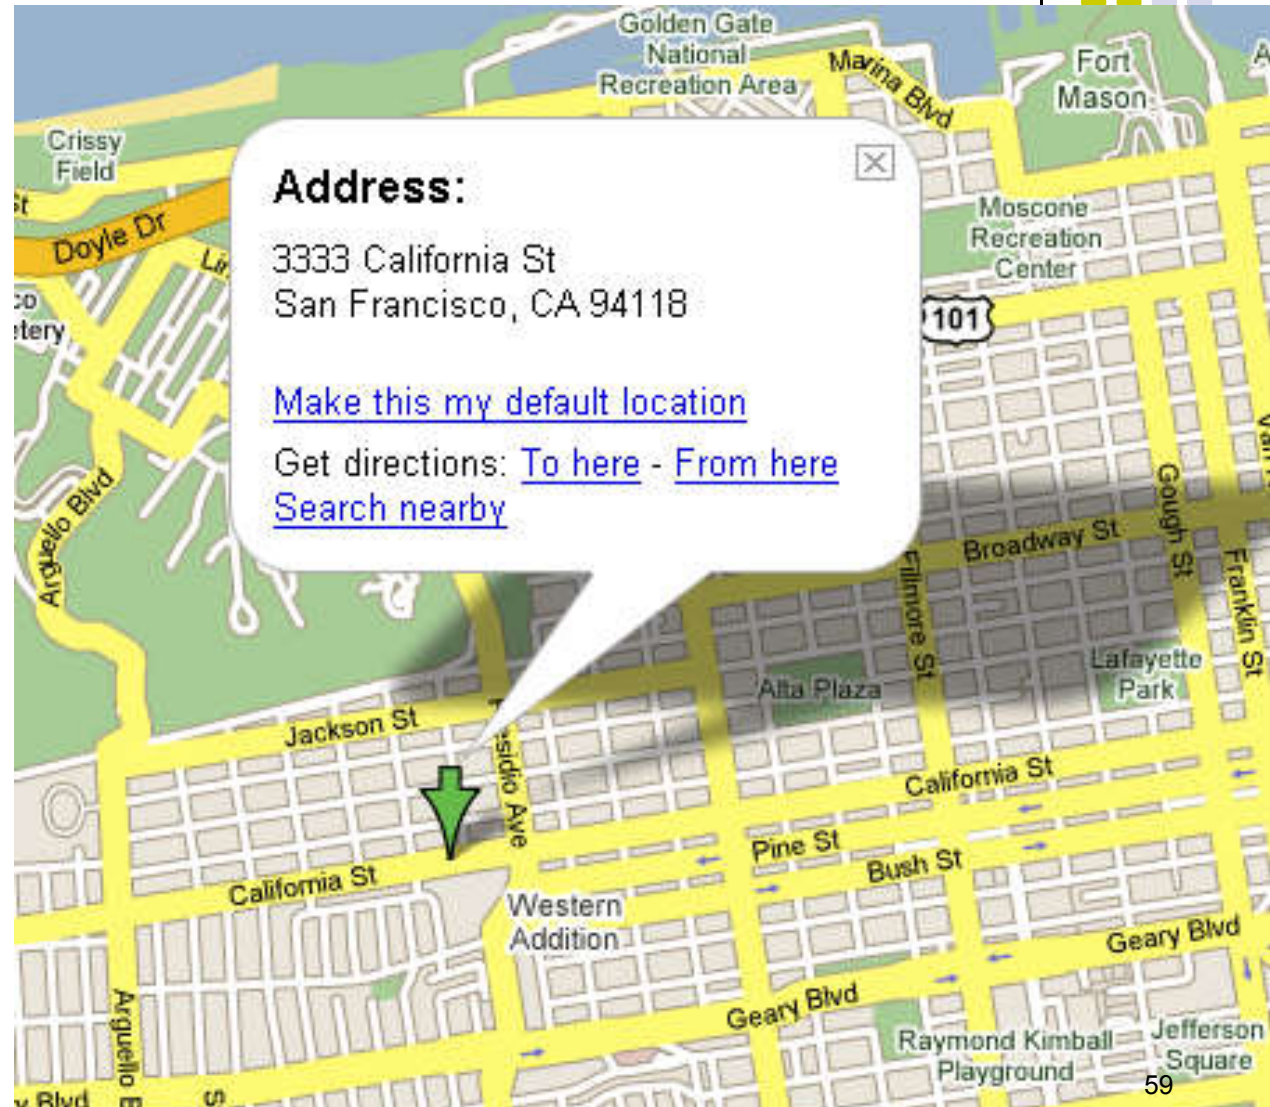
- Two-level index
 - Two level of index block
 - Method to increase file size



Address locating



Go to **USA** →
San Francisco
→
California St
→ **3333**



Indexed Allocation (cont'd)



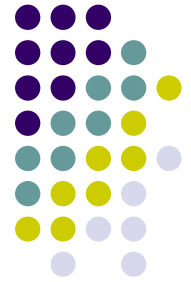
- A system uses two-level index
 - block size 4KB
 - pointer size 4 Bytes
- Which of the following is the correct maximum file size?
 - A. 4GB
 - B. 1GB
 - C. 8GB
 - D. 2GB

Indexed Allocation (cont'd)



- A system uses two-level index
- Which of the following is the correct steps to locate the data at file position n ?
 - A. Identify block number \rightarrow block number in block table \rightarrow offset
 - B. Identify block number in outer index block \rightarrow block number \rightarrow offset
 - C. Identify offset \rightarrow block number
 - D. none of the above

Indexed Allocation (cont'd)




- A system uses two-level index
 - block size 4KB
 - pointer size 4 Bytes
 - File size 20MB
- Which is the correct address of file at position 15MB?
 - A. (4,3072,0)
 - B. (3,768,1023)
 - C. (3,768,0)
 - D. (3,3072,0)

Indexed Allocation (Cont.)



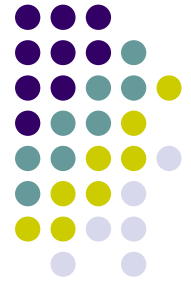
- Linked scheme
 - no limit on file size
- Combine linked list with index block
 - index blocks are linked
 - last pointer of a index block is the address of the next one



index table

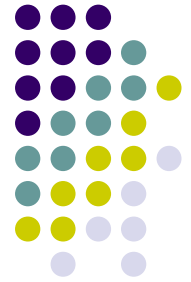
data blocks

Question

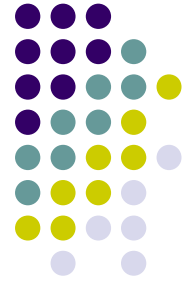


- Linked index block allocation
 - block size 2KB
 - pointer size 4 bytes
 - File size 20MB
- Which of the following is the correct **steps** to read file at position 15.5MB
 - A. identify index block → offset → block number → read
 - B. identify offset → block number → index block → read
 - C. identify offset → index block → block number → read
 - D. identify index block → block number → offset → read

Question

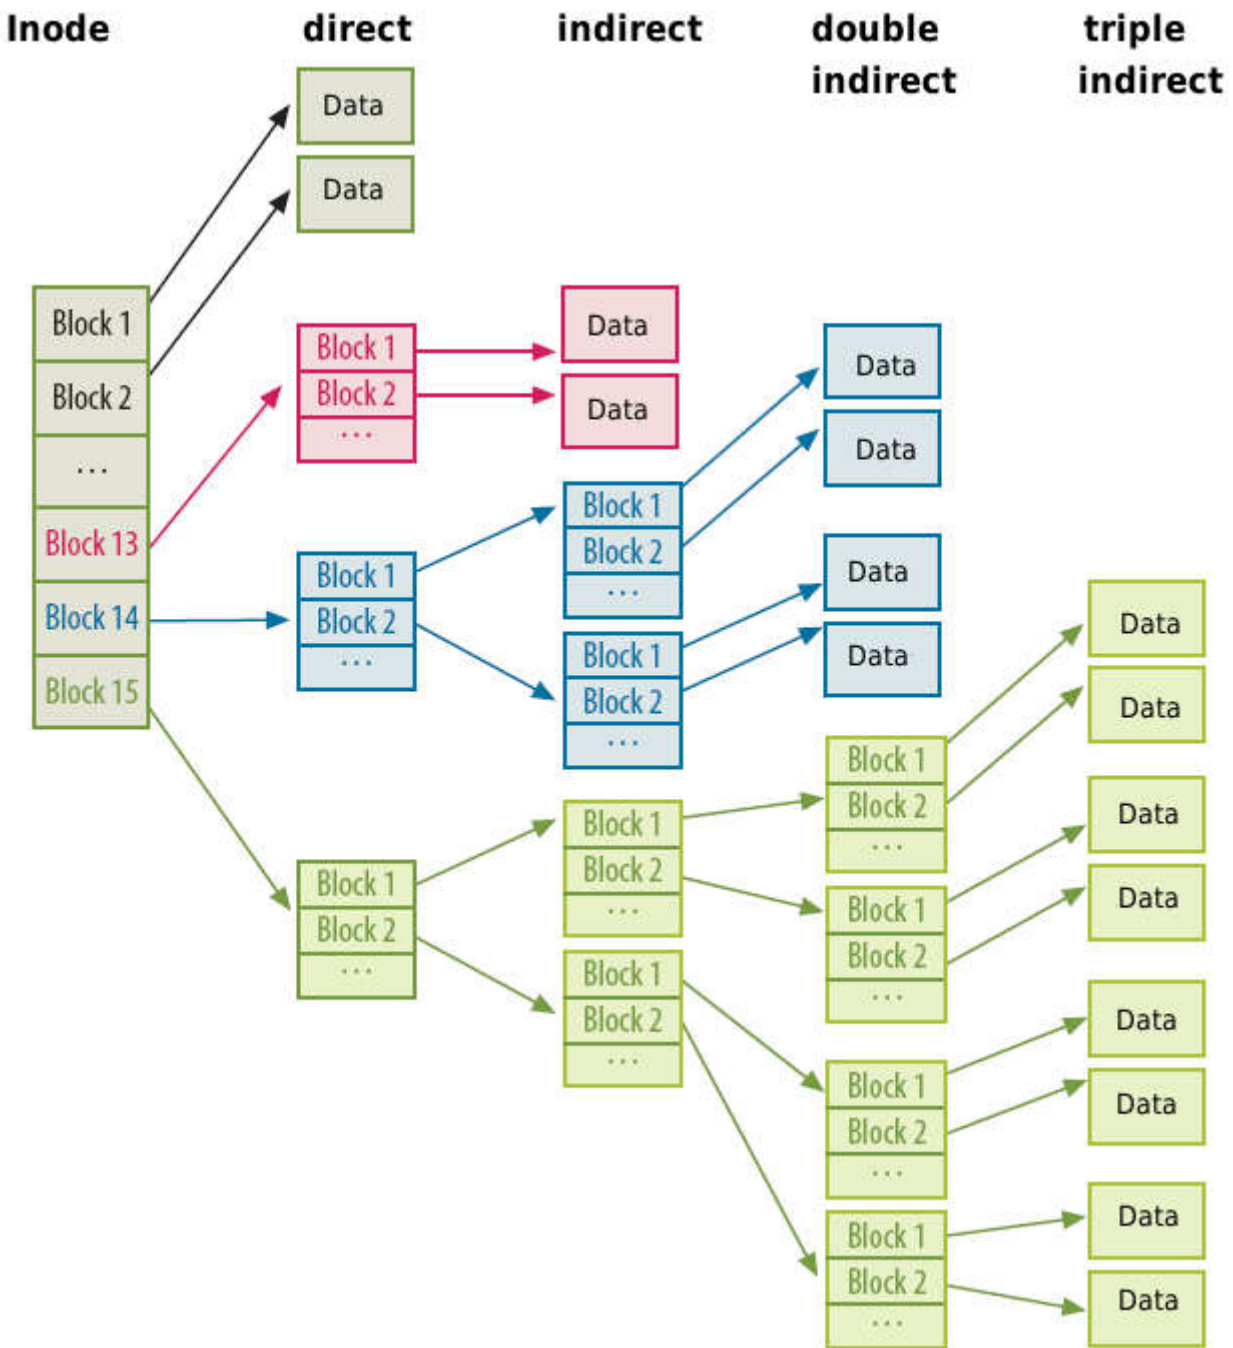


- Linked index block allocation
 - block size 2KB
 - pointer size 4 bytes
 - File size 20MB
- Which of the following is the correct **index block** number at file position 15.5MB (start from 0)
 - A. 13
 - B. 14
 - C. 15
 - D. 16

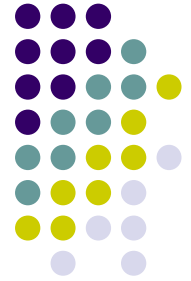


Question

- Linked index block allocation
 - block size 2KB
 - pointer size 4 bytes
 - File size 20MB
- Which of the following is the correct **block** number and **offset** at file position 15.5MB
 - A. $(block, offset) = (271, 2047)$
 - B. $(block, offset) = (271, 0)$
 - C. $(block, offset) = (270, 2047)$
 - D. $(block, offset) = (270, 0)$

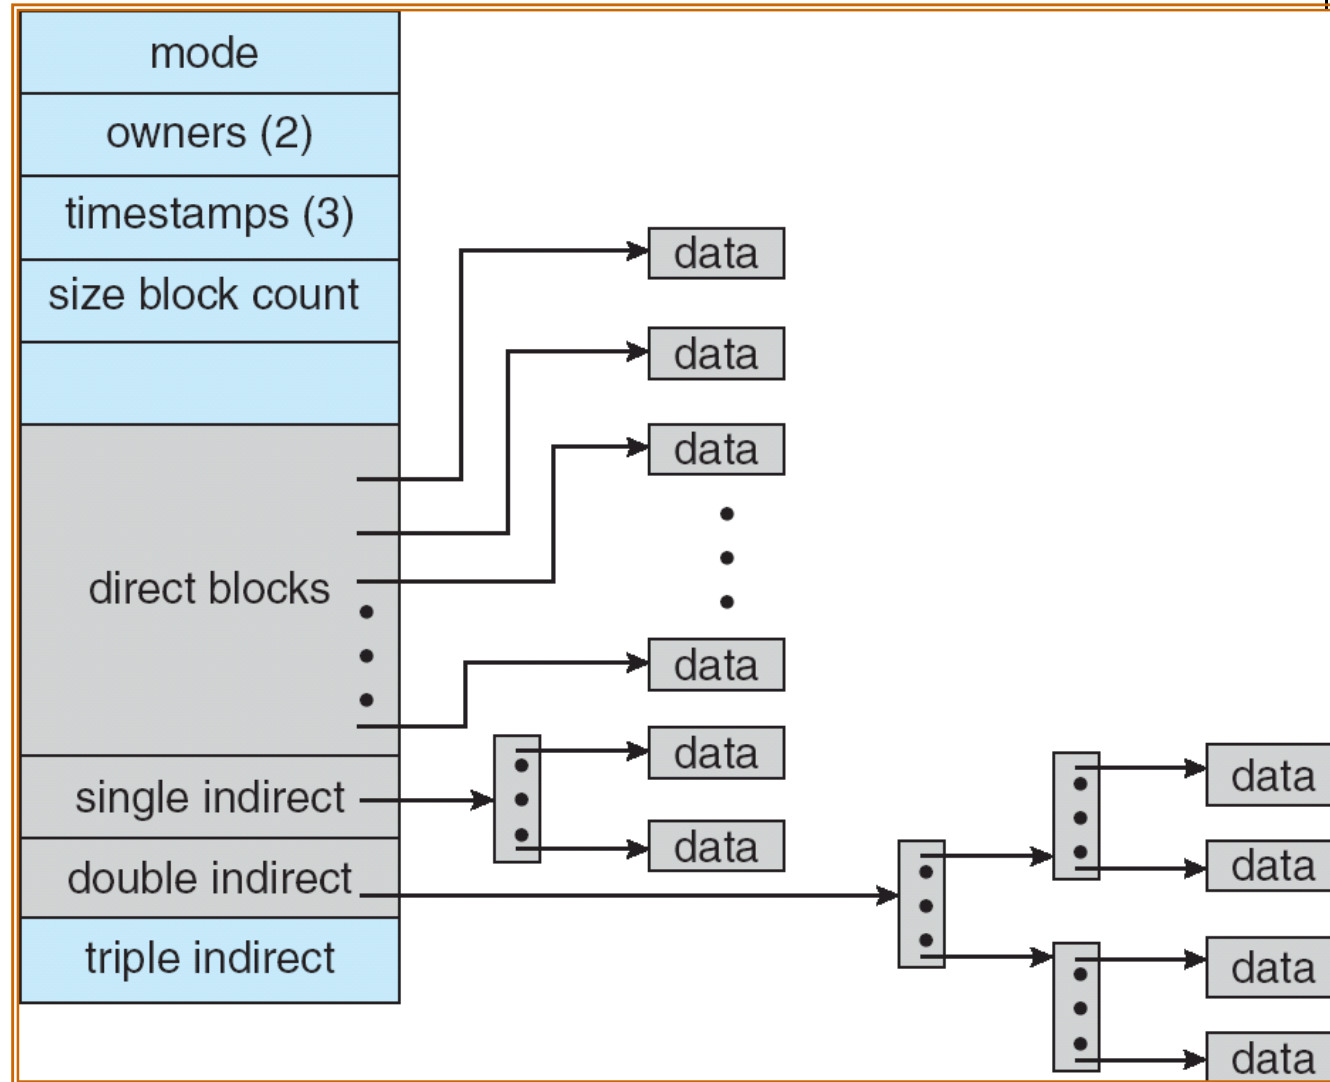


Combined Scheme: UNIX (4K bytes per block)

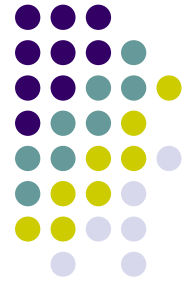


LINUX allocation

(10 direct pointers)

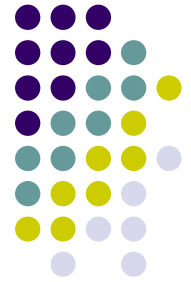


Question



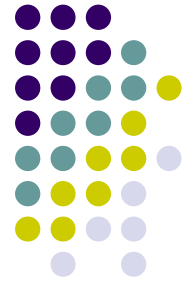
- A UNIX system
 - pointer size 4 bytes
 - block size 4 KB
 - 12 direct pointers, 1 single indirect, 1 double indirect, 1 triple indirect pointers
- Which of the following is the correct **maximum file size**?
 - A. $(12+2^{10}+2^{20}+2^{30})\text{KB}$
 - B. $4*(2^{10}+2^{20}+2^{30})\text{KB}$
 - C. 2^{32}KB
 - D. $4*(12+2^{10}+2^{20}+2^{30})\text{KB}$

Question



- A UNIX system
 - pointer size 4 bytes
 - block size 4 KB
 - 12 direct pointer, 1 single indirect, 1 double indirect, 1 triple indirect pointers
- Which of the following is the correct **maximum number of indexed blocks**?
 - A. $(1+2^{10}+2^{20})$
 - B. $(2+2^{10}+2^{20})$
 - C. $(3+2^{11}+2^{20})$
 - D. 2^{20}

Question

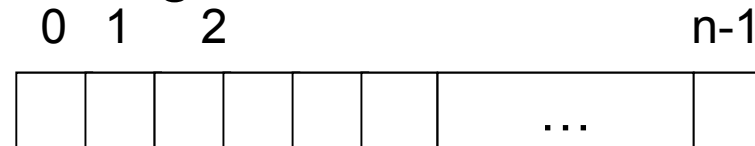


- A UNIX system
 - pointer size 4 bytes
 - block size 4 KB
 - 12 direct pointer, 1 single indirect, 1 double indirect, 1 triple indirect pointers
 - File size 78MB
- Which of the following is the correct **location** at file position 95KB?
 - A. triple indirect block, $(block, offset)=(12,3071)$
 - B. double indirect block, $(block, offset)=(12,3071)$
 - C. single indirect block, $(block, offset)=(11,3071)$
 - D. single indirect block, $(block, offset)=(11,3072)$

Free-Space Management



- Bit vector (n blocks), e.g. Linux (ext3)
 - Easy to get contiguous blocks



$$\text{bit}[i] = \begin{cases} 0 \Rightarrow \text{block}[i] \text{ free} \\ 1 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

- Bit map requires extra space

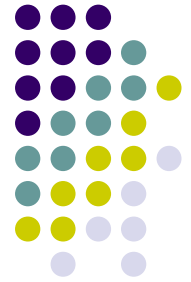
Example:

block size = 2^{12} bytes

disk size = 2^{30} bytes (1 gigabyte)

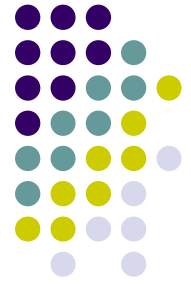
$n = 2^{30}/2^{12} = 2^{18}$ bits (or 32K bytes)

Free-Space Management (Cont.)



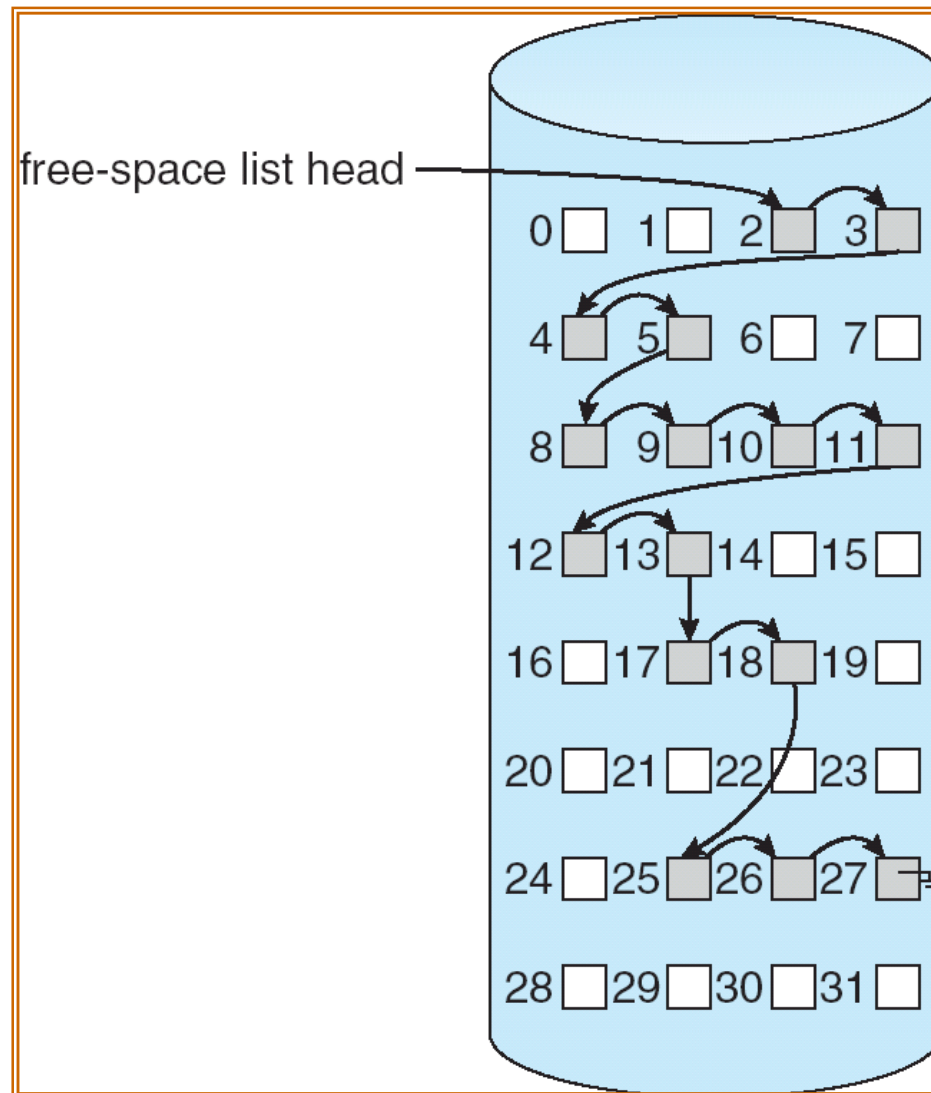
- Linked list (free list)
 - Cannot get contiguous space easily
 - No waste of space
- Grouping
 - Use linked blocks to store pointers to free blocks
 - Last pointer in a block is the address of the next one
- Counting
 - several contiguous blocks are freed/allocated for a file
 - each entry has
 - address of the first free block
 - number of contiguously free blocks

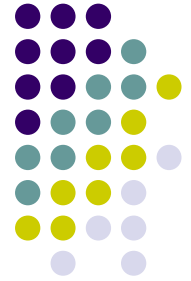
Free-Space Management (Cont.)



- Need to protect:
 - Pointer to free list
 - Bit map
 - Must be kept on disk
 - Copy in memory and disk may differ
 - Cannot allow for block[i] to have a situation where $\text{bit}[i] = 1$ in memory and $\text{bit}[i] = 0$ on disk
 - Solution:
 - Set $\text{bit}[i] = 1$ in disk
 - Allocate block[i]
 - Set $\text{bit}[i] = 1$ in memory

Linked Free Space List on Disk



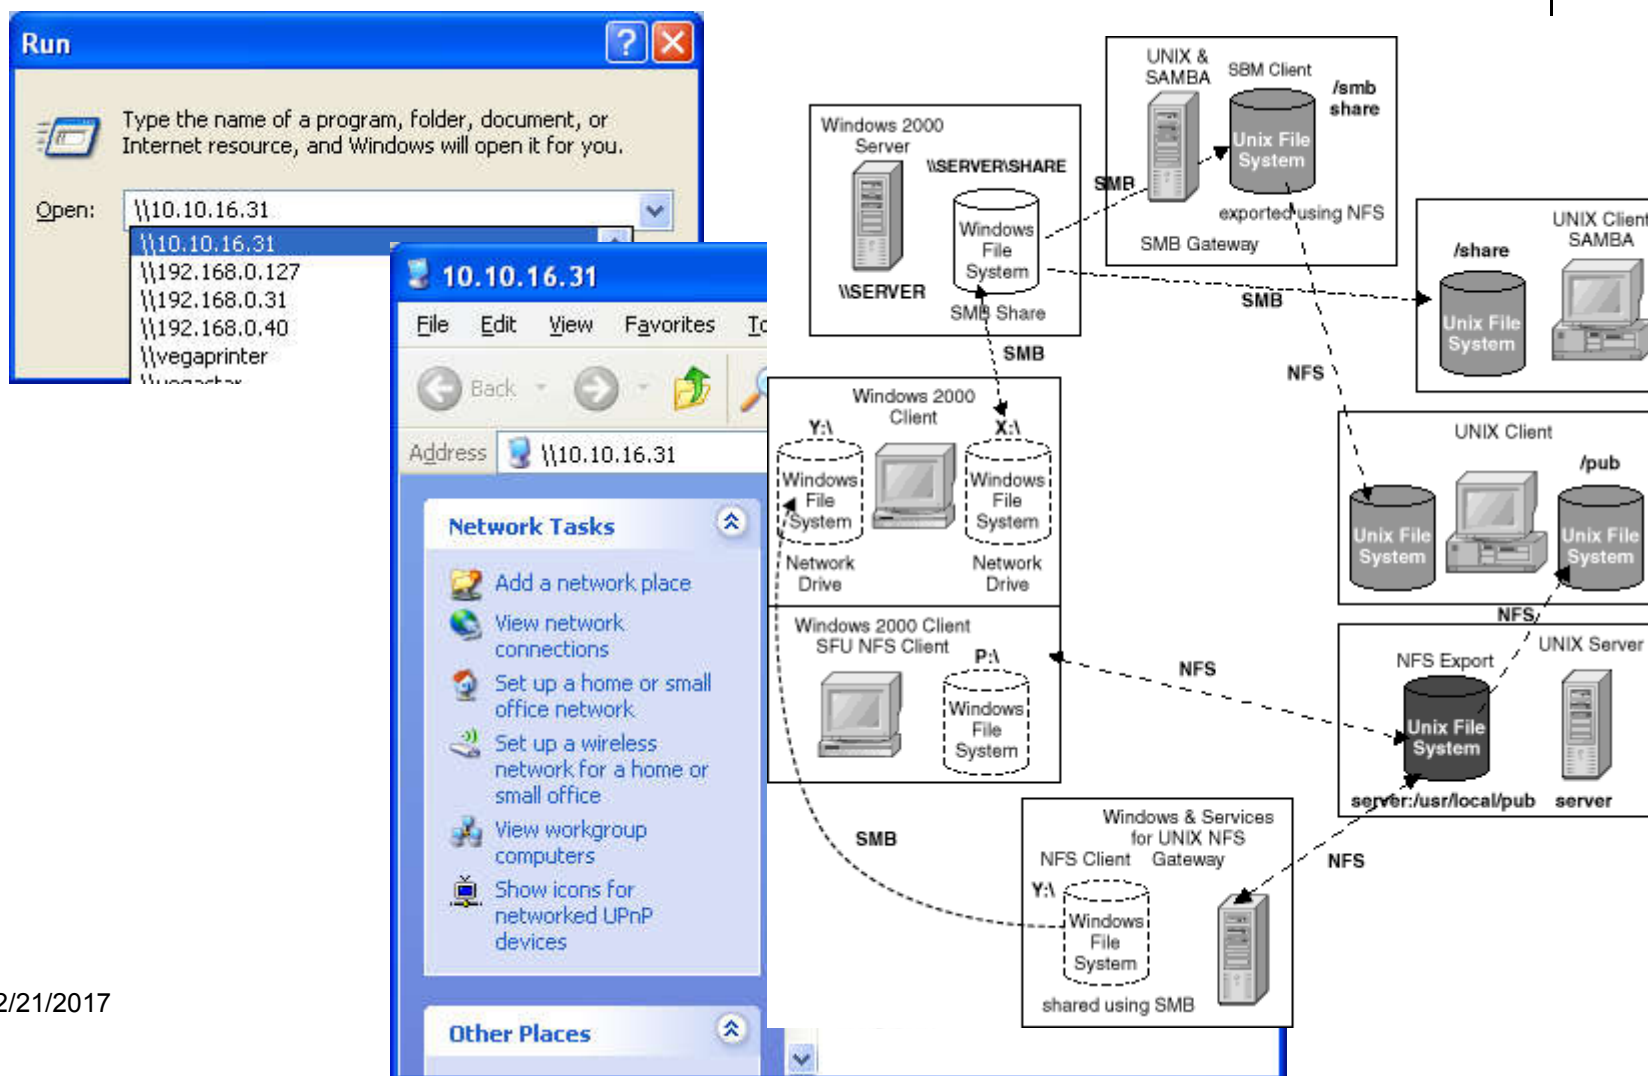


Shared file systems

- A file system is shared to other machines
 - the file system may be large and powerful
 - file sharing is needed in many applications
 - available in many systems, e.g., Windows, Linux (NFS)



Shared file systems



Network File System (NFS) Architecture

